

Exploring Excess Return Generation through European Fixed Income Arbitrages

Alessio Ottaviani

EDHEC Business School

Supervisor: Prof. Riccardo Rebonato

Abstract:

Contents

1	Introduction	8
2	Data, Strategy Construction, and Performance	12
2.1	Data and Strategy Construction	13
2.2	Index Construction and Weighting Schemes	17
2.3	Performance Analysis	18
3	Benchmark Factor Models	21
3.1	Benchmark Specifications	21
3.2	Results and Discussion	24
4	Factor Selection and Spanning Regressions	28
4.1	Candidate Factor Universe	28
4.2	Principal Component Analysis	29
4.3	Adaptive Elastic Net	31
4.4	Conditional Alpha	35
4.5	Discussion, Method Comparison, and Robustness	36
5	Cross-Strategy Interdependencies	37
5.1	Mispricing Construction and Stress Regimes	38
5.2	P1: Cross-Market Co-Movement	39
5.3	P2: Stress Amplification	40
5.4	P3: Lead–Lag Transmission	41
5.5	P4: Persistence and Funding-Dependence	43
5.6	P5: The Common Factor	44
5.7	Discussion: Why These Arbitrages Persist	45

6	Conclusion	48
A	Appendix	56
A.1	BTP Italia: Indexation Coefficient and Deflation Floors	56
A.2	Trade Construction Details	57
A.2.1	Entry/Exit Rules and Thresholds	57
A.2.2	Trade-Level Statistics	57
A.2.3	Transaction Costs	57
A.3	Signal-Weighted vs Equal-Weighted Comparison	58
A.4	Volatility-Managed Portfolios	60
A.5	Candidate Factor Universe	61
A.5.1	Complete Factor List with Sources	61
A.6	PCA Robustness	66
A.6.1	Robustness across the Number of Components	66
A.6.2	Subperiod and Regime Analysis	66
A.6.3	Conditional Alpha: Ferson–Schadt and Threshold Robustness (PCA) . . .	70
A.6.4	Variance Diagnostics	70
A.7	AEN Robustness	71
A.7.1	Sensitivity to Hyperparameters	71
A.7.2	Factor Selection Across Methods	71
A.7.3	Conditional Alpha: Detailed Results (AEN)	72
A.7.4	Rolling Alpha	74
A.8	Subperiod and Rolling Window Analysis	75
A.8.1	Duarte et al. (2007)	76
A.8.2	Brooks et al. (2020) — Active Fixed Income	77
A.8.3	Fung & Hsieh (2004)	79
A.9	Cross-Strategy Interdependencies: Additional Tests	81
A.9.1	Purged Correlations	81

A.9.2	DCC-GARCH	82
A.9.3	Granger Causality: Lag Robustness	83
A.10	VIF Diagnostics	84
Tables		86

List of Tables

1	Entry and Exit Thresholds by Strategy	57
2	Trade-Level Statistics by Strategy	58
3	Transaction Fee Schedule (bps per unit DV01 or duration)	58
4	Equal-Weighted Strategy Returns	59
5	Manipulation-Proof Performance Measure by Risk Aversion (Equal-Weighted) . . .	59
6	Moreira-Muir Volatility-Managed Portfolio Performance (Equal-Weighted)	60
7	Moreira-Muir Volatility-Managed Portfolio Performance	60
8	Candidate Risk Factors	61
9	PCA Spanning Robustness across the Number of Components	66
10	PCA Subperiod Analysis (Contemporaneous)	66
11	PCA Alpha by Regime	68
12	Conditional Alpha (PCA, Contemporaneous): Christopherson, Ferson, and Glass- man (1998) Model	70
13	Conditional Alpha (PCA, Contemporaneous): Stress-Regime Dummies (common β)	70
14	Selected Factors by Method and Strategy	71
15	Conditional Alpha: Christopherson, Ferson, and Glassman (1998) Model	73
16	Conditional Alpha by Stress Regime	74
17	Sub-Period and Regime Analysis: Duarte et al. (2007)	76
18	Sub-Period and Regime Analysis: Active FI (Brooks et al. (2020))	78
19	Sub-Period and Regime Analysis: Fung & Hsieh (2004)	80
20	Correlations After Purging Common Factors	82
21	Granger Causality: Robustness to Lag Length	84
22	Variance Inflation Factors — Duarte et al. (2007), EUR Factors	84
23	Variance Inflation Factors — Brooks et al. (2020), EUR Factors	85
24	Variance Inflation Factors — Fung & Hsieh (2004), EUR Factors	85

25	Summary Statistics for Fixed Income Arbitrage Strategies	87
26	Manipulation-Proof Performance Measure by Risk Aversion	88
27	Duarte Factor Model Regressions	89
28	Fung & Hsieh Factor Model Regressions	90
29	Active Fixed Income Factor Model Regressions	91
30	Alpha Estimates Across Benchmark Factor Models	93
31	PCA Spanning Regressions	94
32	Top Factor Loadings by Principal Component	94
33	Out-of-Sample PCA Spanning Alpha	94
34	Adaptive Elastic Net: Post-Selection OLS on Stable Factors	95
35	Alpha by Stress Regime: PCA and Best-Subset	96
36	Factor Selection Method Comparison	97
37	Unconditional and Regime-Dependent Correlations	98
38	Cross-Strategy Spanning Regressions (Δm)	99
39	Stress Amplification: Interaction Regressions	100
40	AR(1) Persistence of Mispricing Levels	101
41	Common Factor (PC1) and Funding Stress	102

List of Figures

1	Rolling Alpha with Regime Shading (PCA)	69
2	Cumulative Variance Explained by Rolling PCA	71
3	Rolling Alpha with Regime Shading (AEN)	75
4	Rolling 36-Month Alpha with Regime Shading — Duarte et al. (2007)	77
5	Rolling 36-Month Alpha with Regime Shading — Brooks et al. (2020)	79
6	Rolling 36-Month Alpha with Regime Shading — Fung & Hsieh (2004)	81
7	Dynamic Conditional Correlations (DCC-GARCH)	83
8	The TIPS–Treasury Basis, 2005–2025	103
9	Market-Level Basis Time Series	104
10	Cumulative Return Indices and Drawdowns	105
11	Rolling Alpha Across Benchmark Factor Models	106
12	PCA Scree Plot	107
13	Pairwise Correlations by Stress Regime	108
14	Generalized Impulse Response Functions	109

1 Introduction

Duarte et al. [2007] pose the question that motivates our analysis: do fixed-income arbitrage strategies earn genuine alpha, or are their returns simply compensation for systematic risk—“picking up nickels in front of a steamroller”? Using a ten-factor spanning regression, they find that the more complex strategies produce significant alpha that survives equity and bond market risk controls. We take this question to a new setting—European fixed-income markets—and to three strategies that have not previously been studied jointly as systematic arbitrages, namely the BTP Italia inflation-linked basis, the European corporate CDS–bond basis, and the iTraxx index skew.

A second set of questions then arises. Does this alpha withstand a far more demanding test of risk—and which risks does each strategy ultimately bear? Is the alpha that survives roughly constant, or time- and state-dependent? And, perhaps above all, why does it persist? The presidential address of Duffie [2010] provides a framework: when intermediary balance-sheet capacity is the binding constraint on arbitrage, mispricings across different markets should co-move, the co-movement should intensify in stress, and the transmission should follow a liquidity pecking order from more liquid to less liquid markets. Brunnermeier and Pedersen [2009] provide the mechanism—shocks to funding liquidity and market liquidity reinforce each other in a spiral that forces simultaneous deleveraging across positions—and Mitchell and Pulvino [2012] document the empirical counterpart, showing that arbitrage crashes propagate across strategies as capital withdraws.

Each of the three arbitrage strategies we study has an established literature, but none has been analysed in the way we propose. The inflation-linked bond basis was documented by Fleckenstein et al. [2014] in the context of US Inflation-Linked Treasury Bonds (TIPSs), who found an average mispricing of 54.5 basis points between TIPSs and nominal Treasuries over the 2004–2009 period, with peaks exceeding 200 basis points during the 2008–2009 financial crisis (Figure 8). This basis has since compressed substantially, oscillating in a 0–30 bps band post-2009, reflecting broader institutional participation in the U.S. inflation-linked market—and, as Fleckenstein et al. [2014]

show, the basis narrows as hedge fund capital flows into the trade, providing direct evidence for the slow-moving capital channel. In their Section VII, they further document that changes in the TIPS–Treasury mispricing are correlated with changes in the corporate CDS–Bond basis and the CDX index–component basis, establishing the cross-strategy co-movement that we extend to the European setting in Section 5. In contrast, the analogous basis on Italian BTP Italia bonds persists to date. To the best of our knowledge, the BTP Italia has not been the subject of any prior academic study, and its inflation-linked basis has not previously been analysed as a systematic arbitrage strategy. Its persistence is particularly striking because the BTP Italia eliminates the ballooning risk that characterises TIPSs and other conventional inflation-linked sovereigns,¹ and because its retail-dominated investor base creates a structural segmentation that insulates the basis from the institutional capital flows that closed the U.S. mispricing.

Moving to the second strategy we analyse, the CDS–Bond basis is the difference between the CDS spread and the bond credit spread for the same reference entity, a relationship that should be zero by no arbitrage in a frictionless market. It was first documented by Hull et al. [2004] and Longstaff et al. [2005] in its pre-crisis positive form. Mitchell and Pulvino [2012] show that during 2007–2009, tightening financial constraints drove the basis deeply negative, and Bai and Collin-Dufresne [2019] provide a comprehensive analysis of its cross-sectional determinants and its connection to limits of arbitrage. However, this literature focuses predominantly on U.S. corporate bond markets, and, to our knowledge, very little attention has been paid to the euro-area corporate counterpart, which is the second focus of this study.

This brings us to the third strategy of interest, the CDS index skew, which is defined as the deviation of the traded index spread from its intrinsic, no-arbitrage value implied by the single-name constituents. In a U.S. context, this strategy has been studied by Junge and Trolle [2015], who use it as a market-wide liquidity measure, and by Boyarchenko et al. [2016], who document that

¹Conventional inflation-linked bonds adjust the principal upward with realised inflation; in deflationary scenarios or when inflation expectations spike, the redemption value can grow to many multiples of par (the so-called *ballooning* risk), introducing path-dependent valuation risk. The BTP Italia instead pays inflation revaluation as part of each semiannual coupon and redeems at par regardless of the inflation path, eliminating this source of uncertainty. The persistence of a positive arbitrage in the BTP Italia basis is therefore especially surprising: the instrument has *less* risk than a conventional inflation-linked bond, not more.

post-crisis regulatory capital requirements have made basis trades between the index and single-name CDS “significantly less economical,” leading to persistently wider skews; both papers focus on the U.S. CDX index. To the best of our knowledge, neither the negative basis trade on European investment-grade corporates in the iTraxx Main universe nor the index skew across the four iTraxx sub-indices has been studied as a systematic arbitrage strategy in the academic literature.

Beyond being the first study to analyse these three strategies jointly—and, for the BTP Italia, the first to analyse it at all—our paper extends the Duarte et al. [2007] framework in three respects:

1. Each of our three strategies has a deterministic payoff at maturity: if the position can be held to maturity, the profit is certain regardless of the path of market variables in the interim. The strategies in Duarte et al. [2007], by contrast—swap spreads, yield curve trades, mortgage hedging, volatility selling, capital structure trades—do not share this property: their payoffs depend on models, mean-reversion assumptions, or stochastic processes (prepayment, realised volatility) that leave the outcome uncertain even at maturity. This yields a cleaner test of the alpha-versus-risk question: when the terminal payoff is certain, a positive average return cannot be compensation for payoff risk, and any surviving alpha points to limits to arbitrage rather than to a missing risk premium.
2. We subject the alpha estimates to a progressive battery of factor controls: three classical benchmark specifications from the fixed-income hedge fund literature [Duarte et al., 2007, Fung and Hsieh, 2004, Brooks et al., 2020], a full-sample principal component analysis following Connor and Korajczyk [1986, 1988] on the candidate factor universe, and the Adaptive Elastic Net of Zou and Zhang [2009] with stability selection [Meinshausen and Bühlmann, 2010], which identifies a sparse set of risk drivers for each strategy individually. The alpha that survives this sequence of increasingly demanding tests is not compensation for any of the identified systematic risk exposures. Because these benchmark loadings are estimated in sample, we re-estimate each benchmark out of sample, forming the factor hedge from past-only loadings; the alpha survives, confirming that it is an abnormal return implementable in real time rather than an arte-

fact of in-sample factor combination. Using the conditional framework of Christopherson et al. [1998] we show that it rises systematically during periods of intermediary funding stress.

3. We test the cross-strategy implications of slow-moving capital systematically, drawing on Duffie [2010], Brunnermeier and Pedersen [2009], Mitchell and Pulvino [2012], and Fleckenstein et al. [2014]: whereas the existing literature has studied each of these arbitrages in isolation, we examine all three jointly in the European market and ask whether their mispricings co-move, whether the co-movement intensifies in stress, and whether a common factor tied to intermediary capital—rather than to macroeconomic fundamentals—explains it. The European setting provides a discriminating test of the mechanism: all three arbitrages are run by the same type of institutional arbitrageur, but the underlying instruments differ in who holds them. The CDS–Bond and iTraxx legs are held and arbitrated by leveraged institutions; the BTP Italia basis, by contrast, rests on an inflation-linked bond held predominantly by unlevered retail households, who face no margin calls and do not unwind their positions under funding stress. We relate the resulting evidence to the structural limits to arbitrage that keep the mispricings open, including the post-crisis regulatory capital costs documented by Boyarchenko et al. [2016].

We find that all three strategies generate positive and economically significant excess returns net of transaction costs, with annualised Sharpe ratios from 0.71 (BTP Italia) to 1.71 (CDS–Bond Basis), positive skewness, and contained drawdowns. The alpha survives controls based on three widely used benchmark frameworks (estimated both in and out of sample), on rolling principal components, and on the Adaptive Elastic Net for the BTP Italia and CDS–Bond strategies; the iTraxx skew is absorbed by its selected factors unconditionally but reappears as a positive residual during high-stress periods and serves as the leading mispricing in the cross-strategy hierarchy, Granger-causing both the CDS–Bond ($p = 0.011$) and the BTP Italia ($p = 0.007$) mispricings. The residual alpha rises during periods of intermediary funding stress, in a manner consistent with the slow-moving capital mechanism of Duffie [2010]. The co-movement of the three mispricings is governed by a common factor that loads on intermediary balance-sheet variables—the He et al. [2017] capital ratio, the Libor–OIS spread, and the Amihud and Mendelson [2015] illiquid-

ity measure—and not on macroeconomic fundamentals (a placebo regression on macro indicators yields uniformly insignificant coefficients). The persistence of the mispricing levels partially mirrors the funding-dependence of each arbitrage trade: the CDS–Bond basis, which requires continuous repo financing of the cash bond, has the longest half-life at approximately three months, roughly one-third longer than the BTP Italia and the iTraxx skew, which mean-revert within two months and are essentially identical to each other. In monthly mispricing changes, the institutional pair (CDS–Bond and iTraxx) co-moves ever more tightly as stress builds, with correlations rising monotonically from -0.21 in low-stress to $+0.32$ in high-stress regimes. The retail-held BTP Italia basis moves the other way: its correlation with the CDS–Bond basis is significantly negative on the full sample at -0.30 , deepening to -0.47 in high-stress months. This opposing behaviour provides independent evidence of investor-base segmentation: a single intermediary-capital shock pushes the two institutional mispricings wider and the unlevered, retail-held BTP Italia tighter.

The paper is organized as follows. Section 2 describes the three strategies, their construction, and their performance. Section 3 evaluates the alpha against three classical benchmark frameworks. Section 4 extends the analysis to a candidate universe of 85 factors using PCA and the Adaptive Elastic Net. Section 5 tests the cross-strategy predictions of the slow-moving capital literature. Section 6 concludes.

2 Data, Strategy Construction, and Performance

This section introduces the three European fixed-income arbitrage strategies studied in this paper: the inflation-linked basis on Italian government bonds (BTP Italia), the negative basis between corporate bonds and their credit default swaps (CDS–Bond Basis), and the CDS index skew—the difference between the traded iTraxx spread and its intrinsic, no-arbitrage spread implied by the single-name constituents (iTraxx Combined). Each strategy exploits a distinct relative-value mispricing in European markets, and each raises the question of whether the mispricing generates risk-adjusted returns that survive transaction costs and rigorous performance diagnostics.

We describe in Section 2.1 how each strategy is constructed, paying particular attention to the institutional features of the BTP Italia that distinguish it from conventional inflation-linked instruments; in Section 2.2 how individual trade returns are aggregated into investable indices, comparing equal-weighted and signal-weighted schemes and incorporating regime-dependent transaction costs; and in Section 2.3 the performance of the resulting strategy indices, verified against manipulation-proof and volatility-managed robustness checks.

2.1 Data and Strategy Construction

Data are sourced from the MOT market operated by Borsa Italiana (BTP Italia and nominal BTP prices), Bloomberg (inflation swap rates and corporate bond spreads), CMA / CME Group (single-name CDS spreads), and J.P. Morgan DataQuery (iTraxx index and constituent data, cross-validated against Barclays Live). The iTraxx skew sample begins in June 2004 with the launch of the first on-the-run series of the iTraxx Europe Main 5-year index, and the CDS–Bond sample begins in September 2008. The BTP Italia sample begins with the first issuance in March 2012. All three samples extend to May 2025. Following Duarte et al. [2007], each individual trade is modelled as a separate fund that opens when the basis crosses an entry threshold and closes when the basis reverts through an exit threshold or the underlying instrument matures. Entry and exit rules are summarised in Table 1; trade-level statistics are reported in Table 2. Over the full sample, the strategies generate 260 trades for BTP Italia, 931 for CDS–Bond Basis, and 206 for iTraxx Combined. Figure 9 plots the market-level basis time series for each strategy.

BTP Italia: The Inflation-Linked Basis. The BTP Italia is an Italian government bond indexed to domestic inflation, measured by the ISTAT FOI consumer price index (excluding tobacco). It differs in two economically meaningful ways from conventional inflation-linked sovereigns such as the U.S. TIPS studied by Fleckenstein et al. [2014].

The first is the absence of *ballooning risk*. A conventional inflation-linked bond accumulates the inflation adjustment in the principal over the life of the bond, so that the notional outstanding

at time t equals $100 \times I_t$, where I_t is the cumulative inflation factor since issuance. In a credit event, the recovery rate (RR) is applied to the original face value of 100, not to the inflation-adjusted notional: an investor holding a TIPS with $I_t = 1.25$ would recover $RR \times 100$ and lose the entire accumulated revaluation of 25 percentage points. The BTP Italia eliminates this exposure entirely. The capital revaluation is paid out semiannually alongside the coupon, so the inflation factor resets to unity at every coupon date and the notional exposed to credit risk remains at par throughout the life of the bond. The two legs of the arbitrage therefore carry identical sovereign credit exposure, which isolates a pure relative-value component in the basis and removes the differential-default contamination that affects the TIPS–Treasury basis in stress periods. The technical specification of the indexation coefficient and the deflation floors on both coupon and principal is reported in Appendix A.1.

The second feature is the investor base. The Italian Treasury introduced the BTP Italia in 2012 explicitly as a retail instrument, structuring its primary placement in two sequential phases: an initial phase reserved exclusively for retail investors (which typically absorbs the bulk of each issuance), followed by a shorter institutional phase. Investors who purchase at issuance and hold to maturity receive an additional loyalty bonus (the *premio fedeltà*) that further incentivises buy-and-hold behaviour. The BTP Italia is therefore held predominantly by households with long holding periods and no regulatory balance-sheet constraints. This segmentation matters because, as we will document in Section 5, it is what makes the BTP Italia basis behave differently from the two institutional credit strategies during stress episodes: retail holders do not face margin calls and do not unwind their positions when intermediary capital withdraws.

We construct the BTP Italia basis following the replication logic of Fleckenstein et al. [2014], adapted to the Italian instrument. The inflation-linked cash flows of the BTP Italia are converted into fixed-rate equivalents using the term structure of zero-coupon inflation swaps, yielding a fixed-rate equivalent yield y_t^{IL} . The basis is defined as the difference between this yield and the yield to

maturity of the maturity-matched nominal BTP y_t^{NOM} :

$$\text{Basis}_t^{\text{BTP}} = y_t^{\text{IL}} - y_t^{\text{NOM}}. \quad (1)$$

A positive (negative) basis indicates that the BTP Italia, once converted to fixed, yields more (less) than the nominal bond, and the trade is executed accordingly through an asset swap (ASW) on the BTP Italia against an offsetting position in the maturity-matched nominal BTP. Both the BTP Italia and the nominal BTP are accepted as general collateral in the euro repo market, so the two legs carry identical funding costs, which rules out funding-cost differentials as an explanation for the observed basis.

CDS–Bond Basis: The Negative Basis Trade. The CDS–Bond basis for issuer i at date t is defined as the difference between the CDS spread and the z-spread of a reference corporate bond:

$$\text{Basis}_{i,t} = \text{CDS spread}_{i,t} - \text{z-spread}_{i,t}, \quad (2)$$

where the z-spread is computed as the constant spread that, when added to the bootstrapped euro zero-coupon swap curve, equates the discounted bond cash flows to the observed market price. When the basis is negative, credit protection via CDS is cheaper than the credit spread embedded in the bond price. The negative basis trade exploits this gap: the arbitrageur purchases the bond (earning the z-spread) and simultaneously buys CDS protection (paying the lower CDS spread), locking in a positive carry that persists until convergence or bond maturity. We restrict the strategy to the negative basis trade; the reverse trade (short bond, sell protection) requires sourcing the corporate bond through a reverse repo, which for corporate names is subject to recall risk—particularly during periods of stress, when the securities lender may recall the borrowed bond, forcing premature unwinding of the position [Bai and Collin-Dufresne, 2019]. In a default event, the bondholder recovers $\text{RR} \times 100$ on par while the CDS contract pays $(1 - \text{RR}) \times 100$, so the combined position returns par regardless of realised recovery—provided the CDS maturity is no shorter than

the bond's, which we ensure by matching hedge tenors to bond residual life (Appendix A.2). The tradable universe is restricted to single-name CDS whose reference entity belongs to the current on-the-run iTraxx Europe Main 5-year index, which ensures liquidity and provides a transparent filter for investment-grade credit quality. When a name exits the on-the-run index at a roll date, no new trades are initiated on that name; existing positions are carried to convergence or maturity. All CDS are centrally cleared through LCH or ICE, eliminating direct counterparty credit risk on the protection leg.

iTraxx Combined: The CDS Index Skew. The CDS index skew is the difference between the traded spread of an iTraxx index and its intrinsic spread—the no-arbitrage level implied by the constituents' CDS curves under standard market conventions:

$$\text{Skew}_t = S_t^{\text{index}} - S_t^{\text{intrinsic}} \quad (3)$$

A positive (negative) skew indicates that the index trades wide (tight) relative to the sum of its parts.² Three forces sustain a non-zero equilibrium skew. First, the index is far more liquid than its single-name constituents, especially in stress [Junge and Trolle, 2015]. Second, the two legs face different sources of demand: institutional investors use the index as the natural macro credit hedge, while dealer banks trade single names for more granular hedging needs, including the hedging of counterparty CVA exposure—forces that can push the two legs in non-synchronised directions and make the skew bidirectional. Third, executing the replicating basket of single-name positions carries operational costs that the index trade does not. Our strategy trades the skew across the four iTraxx sub-indices—Main, Senior Financials, Subordinated Financials, and Crossover—restricting new positions to the on-the-run series at the time of entry. Entry thresholds are calibrated to the historical volatility of each sub-index skew (Table 1). Unlike the CDS–Bond basis, the iTraxx skew is a derivative-on-derivative trade with no cash leg financed in repo. The two legs are traded

²Each constituent's survival curve is bootstrapped from its CDS quotes, and the index's intrinsic spread is the single running spread that reprices the resulting portfolio under the standard quoting convention (flat hazard rate, 40% recovery; see O'Kane, 2008).

under a single netting set—either centrally cleared at the same CCP or bilaterally with the same dealer under an ISDA Master Agreement and CSA—so that long and short positions offset for daily variation margin and balance-sheet purposes.

2.2 Index Construction and Weighting Schemes

Following Duarte et al. [2007], individual trades are treated as separate fictional hedge funds, each with its own entry date, exit date, and daily return stream. The strategy index on any given day is the weighted average of the returns of the funds that are active on that day. The equal-weighted (EW) scheme of Duarte et al. [2007] assigns identical weight $1/K_t$ to every active trade:

$$r_t^{\text{EW}} = \frac{1}{K_t} \sum_{k=1}^{K_t} r_{k,t}. \quad (4)$$

The signal-weighted (SW) scheme, inspired by the curvature signal of Rebonato and Ronzani [2021], weights each trade proportionally to the absolute distance between the basis at entry and its historical equilibrium level,

$$r_t^{\text{SW}} = \sum_{k=1}^{K_t} \frac{|b_{k,0} - \theta_{t_k}|}{\sum_{j=1}^{K_t} |b_{j,0} - \theta_{t_j}|} r_{k,t}, \quad (5)$$

where $b_{k,0}$ is the basis at the entry date t_k of trade k , θ_{t_k} is the expanding-window average of the cross-sectional basis from the start of the sample up to t_k , and $r_{k,t}$ is the signed return of trade k (positive when the position profits, with the direction—long or short—fixed at entry). The absolute value ensures that the weight reflects the magnitude of the dislocation rather than its sign, which is already carried by $r_{k,t}$. The rationale follows Rebonato and Ronzani [2021]: a basis that deviates further from its historical equilibrium represents a richer opportunity and warrants a larger allocation conditional on convergence. Signal-weighted indices dominate equal-weighted indices on both Sharpe and manipulation-proof performance for all three strategies (Appendix A.3, Table 4), and we adopt the SW scheme as the baseline throughout the paper; EW results in the appendix confirm that all findings are robust to this choice.

Market conditions are classified into three regimes based on the level of the iTraxx Europe Main 5-year spread: LOW (below 60 bps), MID (60–100 bps), and HIGH (above 100 bps). The same classification recurs in Section 3, where we decompose alpha by regime. Transaction costs are applied at the individual trade level as entry and exit fees proportional to DV01, and fee levels are regime-dependent: bid-offer spreads and fees rise monotonically from LOW to HIGH, reflecting the empirical widening of dealer quotes during market dislocation. This tiered structure ensures that costs are conservative precisely in the stress periods when arbitrage returns are largest. The full fee schedule is reported in Table 3; from this point forward, all strategy returns are net of these regime-dependent costs.

2.3 Performance Analysis

Table 25 reports summary statistics computed from daily data. All strategy returns are excess returns by construction: the BTP Italia trade is self-financing (long and short legs are both repoed at the same general collateral rate), the CDS–Bond trade earns the z-spread net of the CDS premium (where the z-spread is itself measured over the swap curve, which proxies for the funding cost), and the iTraxx skew trade requires no capital beyond initial margin. Panel A describes the basis levels: all three series are highly persistent, with first-order autocorrelations between 0.87 and 0.98, consistent with the slow-moving capital literature [Duffie, 2010, Brunnermeier and Pedersen, 2009, Mitchell et al., 2007] that we examine formally in Section 5. Panel B describes the strategy returns. The CDS–Bond basis produces the highest annualised return (5.37%) and Sharpe ratio (1.71), with an annualised volatility of 3.15%. iTraxx Combined returns 1.96% p.a. with a Sharpe ratio of 0.76 over the longest sample of the three. BTP Italia earns 1.67% p.a. with a Sharpe ratio of 0.71 over a sample that begins with the first issuance in April 2012. All three strategies exhibit positive return skewness—0.60 for BTP Italia, 2.81 for the CDS–Bond basis, and 0.43 for iTraxx Combined—which cuts against the popular characterisation of fixed-income arbitrage as “picking up nickels in front of a steamroller” [Duarte et al., 2007]. Maximum drawdowns are moderate and short-lived, as Figure 10 illustrates: the deepest drawdown is 4.4% for the CDS–Bond basis during the 2008–

2009 credit crisis, 3.0% for BTP Italia, and 2.5% for iTraxx Combined, and all three are recovered within months.

Manipulation-Proof Performance. Sharpe ratios are vulnerable to manipulation by strategies whose dynamic trading rules alter the shape of the return distribution, which is precisely the structure that defines the indices we construct. Table 26 therefore reports the Manipulation-Proof Performance Measure (MPPM) of Goetzmann et al. [2007] for risk-aversion parameters $\rho \in \{2, 3, 4\}$. The MPPM is the unique performance metric that satisfies four properties: it produces a single-valued score, is independent of dollar scale, cannot be enhanced by information-free trading or leverage, and is consistent with standard financial market equilibrium. These properties make it the appropriate benchmark for strategies, like ours, that use dynamic entry and exit rules.

The key finding is the remarkable stability of the MPPM across risk-aversion levels. For BTP Italia, the measure is 1.56, 1.55, and 1.54 percent per year at $\rho = 2, 3, 4$ respectively; for CDS–Bond Basis, 5.06, 5.01, and 4.97; for iTraxx Combined, 1.83, 1.82, and 1.80. The near-invariance of the MPPM to ρ is a direct consequence of the low return variance and the limited drawdowns documented above. The MPPM penalises strategies that depend on extreme returns through the concave transformation $(1 + r_t)^{1-\rho}$, so a strategy that earns its Sharpe ratio by occasionally suffering large losses sees its MPPM deteriorate sharply as ρ increases. Conversely, when returns are moderate and tails are thin, the transformation is approximately linear and the measure varies little with ρ . Read together with the positive return skewness documented above, the stability we observe indicates that the strategies’ risk-adjusted performance does not rely on tail events: an investor with $\rho = 4$ evaluates these strategies almost identically to an investor with $\rho = 2$. This is strong evidence that the documented Sharpe ratios do not reflect compensation for realised tail risk that would dominate the assessment of a risk-averse investor.

Volatility-Managed Portfolios. A second concern that the literature has raised about fixed-income arbitrage is that unconditional Sharpe ratios may average over a time-varying risk-return trade-off, with periods of attractive compensation alternating with periods in which risk is poorly

priced. Moreira and Muir [2017] propose a simple test of this hypothesis. The managed portfolio scales exposure inversely to the previous month’s realised variance,

$$f_{t+1}^{\sigma} = \frac{c}{\hat{\sigma}_t^2} f_{t+1}, \quad (6)$$

where $\hat{\sigma}_t^2$ is the realised variance over month t and the scaling constant c is chosen so that the managed and unmanaged portfolios have the same unconditional standard deviation. A statistically significant managed-portfolio alpha relative to the buy-and-hold strategy is diagnostic of a time-varying ratio μ_t/σ_t^2 : an investor can earn additional risk-adjusted returns by taking less risk precisely when volatility is high, which can only be rationalised if the risk-return trade-off is itself time-varying.

Table 7 reports results for the three strategies. The CDS–Bond basis produces a large and statistically significant managed-portfolio alpha of approximately 1.68% per year for the SW index, consistent with a time-varying funding-liquidity mechanism: when volatility is low, intermediary balance sheets are healthy and arbitrage capital is available to exploit the negative basis; when volatility spikes, capital withdraws and the trade-off deteriorates. For BTP Italia and iTraxx Combined, the managed-portfolio alphas are slightly negative (−0.22% and −0.16% p.a. respectively, both insignificant at any conventional level), indicating that the realised-variance signal does not identify exploitable variation in the risk-return trade-off for these two strategies. The contrast between strategies is itself informative and foreshadows the regime-dependent alpha analysis of Section 3, where we decompose alpha across stress regimes and show that the funding liquidity pattern documented here for the CDS–Bond basis generalises into a uniform regime effect across all three strategies once we control for benchmark factor exposures.

Having established that all three strategies generate positive and robust risk-adjusted returns net of transaction costs, we turn next to a formal assessment of these alpha estimates against the canonical benchmark factor models for fixed-income arbitrage.

3 Benchmark Factor Models

This section estimates three canonical benchmark spanning regressions from the fixed-income arbitrage literature—Duarte et al. [2007], Fung and Hsieh [2004], and Brooks et al. [2020]—and tests whether the alpha of the three strategies survives once their systematic factor exposures are controlled for. A central question, motivated by the slow-moving capital framework of Duffie [2010], is whether the residual alpha varies systematically with intermediary funding stress. We describe in Section 3.1 the three specifications and the construction of their euro-denominated counterparts, and present in Section 3.2 the regression results, the cross-model alpha synthesis, the regime decomposition, the conditional alpha model of Christopherson et al. [1998], and an out-of-sample re-estimation that forms the benchmark hedge from past-only loadings.

3.1 Benchmark Specifications

Three benchmark specifications are widely used in the empirical literature on fixed-income arbitrage and active fixed-income management, and they provide the natural reference points against which our strategies should be evaluated.

Duarte et al. [2007], in the paper that motivates our research design, test whether fixed-income arbitrage excess returns are compensation for exposures to equity, banking-sector, and term-structure risk:

$$\begin{aligned} r_{i,t} = & \alpha_i + \beta_1 Mkt_t + \beta_2 SMB_t + \beta_3 HML_t + \beta_4 UMD_t + \beta_5 RS_t + \beta_6 RI_t \\ & + \beta_7 RB_t + \beta_8 R2_t + \beta_9 R5_t + \beta_{10} R10_t + \varepsilon_{i,t}, \end{aligned} \quad (7)$$

where Mkt , SMB , HML , and UMD are the Fama–French–Carhart equity factors (market, size, value, and momentum), RS is the S&P bank stock index excess return, RI and RB are A/BAA-rated industrial and bank bond index excess returns, and $R2$, $R5$, $R10$ are the CRSP Fama bond portfolios at two-, five-, and ten-year maturities.

Fung and Hsieh [2004] propose a seven-factor *asset-based style* model that tests whether hedge fund returns are explained by equity, fixed-income, and trend-following exposures:

$$r_{i,t} = \alpha_i + \beta_1 S\&P_t + \beta_2 SC-LC_t + \beta_3 BdOpt_t + \beta_4 FXOpt_t + \beta_5 ComOpt_t + \beta_6 \Delta 10Y_t + \beta_7 \Delta CredSpr_t + \varepsilon_{i,t}, \quad (8)$$

where *S&P* is the S&P 500 excess return, *SC-LC* is the small-minus-large-cap total return, *BdOpt*, *FXOpt*, and *ComOpt* are the three trend-following factors constructed as lookback straddle returns on bond, currency, and commodity futures, $\Delta 10Y$ is the change in the 10-year Treasury yield, and $\Delta CredSpr$ is the change in the Baa–Treasury credit spread.

Brooks et al. [2020] take a different angle, decomposing the returns of active fixed-income managers into eight *traditional risk premia* organised along three channels of risk-taking—duration, credit, and volatility selling:

$$r_{i,t} = \alpha_i + \beta_1 Term_t + \beta_2 GlTerm_t + \beta_3 GlAgg_t + \beta_4 InflLink_t + \beta_5 CorpCred_t + \beta_6 EmDebt_t + \beta_7 EmCurr_t + \beta_8 USTVol_t + \varepsilon_{i,t}, \quad (9)$$

where *Term* and *GlTerm* are the total returns on the US Treasury and the currency-hedged global Treasury indices, both in excess of cash (3-month Treasury bills), *GlAgg* is the total return on the currency-hedged global aggregate bond index in excess of cash, *InflLink* is the total return on the currency-hedged global inflation-linked Treasury index in excess of cash, *CorpCred* is a 50/50 blend of a US high-yield corporate bond index (in excess of duration-matched Treasuries) and a leveraged loan index (in excess of 3-month LIBOR), *EmDebt* is the duration-adjusted excess return on an emerging-market debt index, *EmCurr* is the return on an equal-weighted basket of emerging-market currencies, and *USTVol* is the return on delta-hedged straddles on 10-year Treasury futures.

Each benchmark is estimated in two parallel panels. The first uses the original U.S.-denominated factors as specified by the respective authors, which is the natural way to replicate each framework

as it was first published. The second uses euro equivalents, since all three strategies trade European instruments and may therefore be better explained by European factor exposures.

For the euro panel, the Fama–French factors—used directly in Duarte et al. [2007] and indirectly in Fung and Hsieh [2004]—are sourced from the Kenneth French data library; because these are reported in USD, we convert them to EUR following Glück et al. [2021]. For long factors (market excess return):

$$F_t^{\text{EUR}} = \frac{1 + F_t^{\text{USD}} + r_{f,t}^{\text{USD}}}{1 + r_t^{\text{USD/EUR}}} - 1 - r_{f,t}^{\text{EUR}}, \quad (10)$$

and for long-short factors (SMB, HML, UMD):

$$LS_t^{\text{EUR}} = \frac{LS_t^{\text{USD}}}{1 + r_t^{\text{USD/EUR}}}, \quad (11)$$

where F_t^{USD} and LS_t^{USD} are the original factors in US dollars, $r_{f,t}^{\text{USD}}$ and $r_{f,t}^{\text{EUR}}$ are the US and euro risk-free rates (Fama–French T-bill rate and one-month Euribor, respectively), and $r_t^{\text{USD/EUR}}$ is the period return on the dollar-to-euro exchange rate. All remaining factors enter the euro panel with native European counterparts: European bank stocks, euro-area corporate bond indices, and German government bonds in the Duarte specification; euro-area term, credit, and inflation-linked premia in Brooks et al.; and the 10-year Bund yield and Baa–Bund credit spread in Fung and Hsieh. The Fung and Hsieh trend-following lookback straddle factors are currency-independent by construction (D. Hsieh, private communication) and enter both panels identically. The complete factor list is reported in Appendix A.5.1.

All specifications are estimated by ordinary least squares with Newey–West HAC standard errors [Newey and West, 1987], with the truncation lag set as $\lfloor T^{1/4} \rfloor$ following Newey and West [1994]. Variance inflation factors for each specification are reported in Appendix A.10. All factors exhibit VIF below 10 except for the Fama bond portfolios (R2, R5, R10) in the Duarte specification (VIF $\approx$ 15–55), which reflects the well-known correlation structure of maturity-sorted Treasury returns.

3.2 Results and Discussion

Factor Model Regressions. Tables 27, 28, and 29 report the full regression results. Across all three frameworks, all three strategies, and both currency denominations, the estimated intercept is positive and statistically significant at the one-percent level. The annualised alphas are economically meaningful: approximately +5% p.a. for the CDS–Bond basis across all specifications, +1.6–3.5% for iTraxx Combined, and +1.6–2% for BTP Italia.

The adjusted R^2 is low across all specifications, ranging from one to eighteen percent, indicating that the standard factor sets leave most of the return variation unexplained—a question we take up in Section 4. The significant factor loadings that do emerge differ across strategies and are economically intuitive. The CDS–Bond basis loads on credit and rates factors: the A-rated industrial corporate bond return is positive and significant in the Duarte specification in both currency denominations, and the change in the 10-year Treasury yield enters negatively in the Fung–Hsieh specification, again in both denominations—the natural exposures of a long-credit, cash-financed position. BTP Italia and iTraxx Combined load instead on equity factors: the market, size, and value factors in Duarte and the small-minus-large-cap spread in Fung–Hsieh, with the broad equity market significant for BTP Italia but not for iTraxx Combined. These loadings carry negative signs, consistent with the modest indirect equity exposure embedded in the two basis trades. The Brooks et al. specification, which uses broader risk premia rather than individual asset returns, produces fewer significant loadings but delivers comparable alpha estimates, reinforcing that the intercepts are not an artefact of any single factor set.

Cross-Model Comparison and Regime Analysis. Table 30 consolidates the alpha estimates from the three frameworks and decomposes them by market regime. Panel A reports the unconditional alphas side by side: for each strategy, all three frameworks deliver positive and significant alpha, and the estimates are consistent across specifications despite their different factor structures.

Panel B decomposes alpha into two sub-samples defined by the level of the iTraxx Europe Main 5-year spread, using the same regime classification introduced in Section 2.3: HIGH-stress

months are those in which the spread exceeds 100 basis points, and NORMAL the remaining observations. In every cell of the table, alpha during HIGH-stress months is substantially larger than alpha during NORMAL months, in most specifications by a factor of two to four. The pattern holds across all three benchmark specifications and across all three strategies, and it does not depend on the precise threshold used to define the HIGH regime: in Appendix A.8 we report the same exercise at threshold levels of 80, 100, and 120 basis points, with similar results.

This regime dependence is consistent with the theoretical mechanism that Duffie [2010] formalises in his presidential address. When intermediary capital is abundant, arbitrage forces are strong and mispricings are quickly closed, compressing the residual premium that remains after compensating for the systematic factor loadings. When intermediary capital is scarce—as it is during episodes of widening credit spreads—the arbitrageurs who ordinarily close the gap retreat from the trade or are forced to unwind existing positions, the mispricings widen, and the residual premium for those who can still deploy capital rises sharply. This is precisely the regime dependence documented in Panel B, and it is also the empirical pattern that motivates the cross-strategy analysis in Section 5, where we test whether the same intermediary balance-sheet shocks drive co-movement of mispricings across the three strategies.

Conditional Alpha. Panel C of Table 30 formalises the regime effect using the conditional performance evaluation framework of Christopherson et al. [1998], which lets the alpha intercept vary continuously with a publicly observable state variable rather than across a binary regime split. The specification is

$$r_{i,t} = \alpha_0 + \alpha_1 z_t + \beta' \mathbf{X}_t + \varepsilon_t, \quad (12)$$

where $r_{i,t}$ is the strategy excess return, $\mathbf{X}_t$ collects the factor returns of the chosen benchmark, and z_t is the standardised level of the iTraxx Europe Main 5-year spread: $z_t = (s_t - \bar{s})/\hat{\sigma}_s$. The state variable z_t is constructed to be mean-zero and unit-variance over the sample, so the intercept α_0 retains its interpretation as the unconditional alpha and the slope α_1 measures the additional alpha earned per one-standard-deviation increase in funding stress. Standard errors are again computed

with the Newey and West [1994] HAC correction. The Christopherson et al. [1998] formulation improves on the binary split in Panel B by using the full continuous variation in the stress proxy rather than collapsing it into a dummy, which is more efficient when the underlying state varies meaningfully over time.

The estimates of α_1 in Panel C are positive for all three strategies and all three benchmark specifications, and the unconditional intercept α_0 remains significant in every case, confirming that a baseline alpha persists even after conditioning on the stress state. For CDS–Bond Basis and iTraxx Combined, α_1 is significant at the one-percent level across all three frameworks: a one-standard-deviation increase in the iTraxx Main spread raises CDS–Bond alpha by approximately three percentage points per year and iTraxx alpha by one to three percentage points. For BTP Italia, α_1 is positive in all specifications but less precisely estimated (significant only at the 5–10% level), consistent with the shorter sample available for this strategy. For all three strategies, the implied alpha at $z = +1\sigma$ is approximately three to nine times larger than at $z = -1\sigma$, indicating that the bulk of the unconditional alpha is earned during periods of elevated funding stress. Reading Panel B and Panel C together, the unconditional alpha masks a pronounced time variation consistent with the slow-moving capital mechanism of Duffie [2010].

Out-of-Sample Alpha. A residual concern, central to the literature on out-of-sample performance evaluation [Moreira and Muir, 2017, Cederburg et al., 2020], is that the benchmark loadings in Tables 27–29 are themselves estimated in sample, so the surviving alpha could in principle reflect an in-sample-optimal combination of the factors rather than an abnormal return that an investor could have earned in real time. Panel D of Table 30 addresses this directly. For each strategy and each benchmark, we re-estimate the factor loadings on an expanding past-only window in the recursive out-of-sample design of Welch and Goyal [2008], form the factor hedge from the lagged loadings, and define the out-of-sample alpha as the mean of the resulting hedged return, tested with the Newey and West [1994] HAC correction. The expanding window is initialised with a 60-month minimum, so that the first benchmark fit has at least five observations per parameter under the

ten-factor Duarte et al. [2007] specification. The out-of-sample alpha remains positive and statistically significant for all three strategies and all three frameworks. The point estimates are broadly comparable to, and in most cases modestly below, their in-sample counterparts in Panel A—the attenuation expected once the hedge can no longer be fitted to the same returns it is applied to—and the ranking across strategies is preserved, with the CDS–Bond basis the largest and BTP Italia, estimated over the shortest sample, the least precisely identified. The rolling-window counterpart, which additionally relaxes the assumption of constant loadings, is reported in Appendix A.8 and delivers the same conclusion.

Rolling Alpha. A natural concern with regime-based or conditional analyses is that the estimated relationship between alpha and stress may be driven by a single sub-period—perhaps the global financial crisis of 2008–2009, or the European sovereign debt crisis of 2011–2012, or the COVID-19 shock of March 2020—rather than by a stable structural feature of the data. Figure 11 addresses this concern directly by plotting rolling 36-month alpha estimates for each strategy under each of the three benchmark specifications. The rolling alpha is positive across essentially the whole sample for all three strategies, so the estimates are not the product of a single favourable sub-period but persist throughout, and the elevated-alpha intervals are not confined to one crisis: they recur across the distinct stress episodes spanned by the sample. The three benchmark estimates move broadly together within each panel, reinforcing the cross-model robustness documented in Panel A of Table 30. Detailed sub-period splits and threshold robustness for each framework individually are reported in Appendix A.8.

Taken together, the five sets of results in this section establish two findings that are central to the rest of the paper. First, the alpha of the three arbitrage strategies is robust: it survives controls based on three widely used benchmark frameworks for fixed-income arbitrage, in both currency denominations; it is not an artefact of in-sample factor combination, surviving when each benchmark hedge is re-estimated out of sample; and it does so for every cross-section and every sub-period that we examine. Second, the alpha is not constant. It rises systematically when intermediary cap-

ital is scarce, consistent with the slow-moving capital theory, and the regime effect is large: in several specifications the alpha earned in high-stress months is a multiple of the alpha earned in normal months. The first finding indicates that the observed excess returns are not explained by the systematic exposures that the prior literature has identified. The second finding points to a deeper economic mechanism, but it does not yet identify which specific risk factors are responsible for the residual variation. We turn to that question in Section 4, where we use penalised regression methods on a much larger candidate factor universe to identify the factors that these three benchmarks omit.

4 Factor Selection and Spanning Regressions

The benchmark factor models of Section 3 are designed to capture known systematic exposures, but they leave open the question of whether a broader set of risk factors can better explain strategy excess returns and absorb the residual alpha. Starting from a candidate universe of 85 risk factors assembled in Section 4.1, we proceed in two steps. We begin in Section 4.2 with a principal component analysis following Connor and Korajczyk [1986, 1988], which extracts a small number of latent factors from this candidate set and tests whether these diffuse combinations span strategy excess returns. We then turn in Section 4.3 to the Adaptive Elastic Net of Zou and Zhang [2009], a penalised regression that selects individual observable factors from the same candidate set, and we compare its output with alternative machine learning methods in Section 4.5. As in Section 3, we examine whether the alpha that survives factor controls varies with intermediary funding stress.

4.1 Candidate Factor Universe

We assemble a candidate set of 85 risk factors from a comprehensive review of the empirical asset pricing, intermediary, and macro-finance literatures. Each factor is selected on the basis of a published theoretical or empirical link to fixed-income arbitrage returns, intermediary balance-sheet constraints, or broad market risk conditions. The factors are organised into six economic categories: credit risk (Panel A), liquidity and funding (Panel B), equity (Panel C), volatility (Panel D),

macro-financial (Panel E), and interest rates and term structure (Panel F). The complete list with definitions, data sources, and references is reported in Table 8 (Appendix A.5.1). This candidate universe serves as the common input to both the PCA of Section 4.2 and the penalised regressions of Section 4.3.

The candidate set combines tradable returns (equity and bond portfolios, currency baskets, trend-following straddles) with non-tradable state-variable proxies for intermediary capital, funding liquidity, and macroeconomic conditions (e.g., HKM_IC, Libor-OIS, Δ ILLIQ, Δ UR, Δ UF). Under the mimicking-portfolio theorem of Huberman et al. [1987], extended to multi-factor settings by Breeden et al. [1989], any non-tradable factor admits a mimicking portfolio constructed as the projection of the factor onto the span of the tradable asset returns, and the resulting pricing relation is identical to that obtained from the original factors. The intercept α_i in the spanning regressions of this section therefore retains its interpretation as the component of strategy returns that is not compensation for any of the identified risk exposures, regardless of whether those exposures are entered as tradable or non-tradable proxies. This follows the empirical tradition of Fung and Hsieh [2004], who similarly mixes tradable returns with non-tradable state variables in a spanning framework.

4.2 Principal Component Analysis

When the number of candidate factors is large relative to the sample size, regressing strategy returns on the full set is unreliable: the strong collinearity of the candidate factors inflates the variance of the coefficient estimates and overfits the sample. The standard response in the asset-pricing literature is to summarise the common variation in a small number of principal components and test whether these diffuse combinations span strategy returns. Estimating the latent common factors from a large panel by principal components and using them to measure the abnormal return of a managed position is the performance-measurement framework of Connor and Korajczyk [1986, 1988]; by the mimicking-portfolio argument of Section 4.1, the resulting intercept α_i coincides

with that obtained from a fully tradable factor set, so the test retains its interpretation as the component of strategy returns that is not compensation for the identified common risks.

We estimate the components from the full estimation sample rather than from a trailing window. For an in-sample spanning test—as distinct from a return *forecast*—using the entire sample to estimate the loadings is both standard and desirable: it delivers the most precise estimate of the common-factor space, and with the cross-sectional dimension N and the time dimension T both large it yields consistent and asymptotically normal second-stage estimates of α_i and β_i [Giglio and Xiu, 2021]. The concern that full-sample loadings embed information unavailable in real time is addressed directly below: the Out-of-Sample Alpha results of Table 33 re-estimate both the components and the betas recursively, using only data available before each evaluation date.

We standardise each factor to zero mean and unit variance over the estimation sample and retain the first $K = 8$ components, the number selected by the Bai and Ng [2002] IC_{p2} criterion (the less parsimonious IC_{p1} indicates thirteen). The eight components capture 64.7% of the variance in the balanced factor panel (Figure 12), and robustness of the alpha to $K \in \{5, 8, 11, 13\}$ is reported in Appendix A.6.1.

The spanning regression takes the form

$$r_{i,t} = \alpha_i + \beta_i' \mathbf{PC}_t + \varepsilon_{i,t},$$

where $r_{i,t}$ is the monthly excess return of strategy i and $\mathbf{PC}_t$ is the $K \times 1$ vector of principal component scores at time t , estimated contemporaneously ($\mathbf{PC}_t \rightarrow r_{i,t}$). Inference uses Newey–West HAC standard errors [Newey and West, 1987]. Because the regressors are estimated principal components, the second-stage standard errors are in principle subject to a generated-regressor adjustment; under $N, T \rightarrow \infty$ with $\sqrt{T}/N \rightarrow 0$ this adjustment is asymptotically negligible [Bai and Ng, 2006], and we verify that the point estimates and their significance are unchanged under a stationary block bootstrap that re-estimates the components within each replication [Politis and Romano, 1994].

Table 31 reports the estimates. The intercept remains positive and significant for [strategie]: +XX% p.a. for BTP Italia, +XX% for CDS–Bond, and +XX% for iTraxx Combined. The adjusted $\bar{R}^2$ is low throughout, reaching at most XX%, indicating that the diffuse principal components do not span strategy excess returns. Table 32 reports, for each retained component, the candidate factors with the largest absolute loading. The components admit a clear economic reading: PC1 is dominated by credit and equity-market factors, PC2 by global rates and bond returns, and PC3 by swap spreads and dealer CDS.

Out-of-Sample Alpha. A residual concern, central to the literature on out-of-sample performance evaluation [Moreira and Muir, 2017], is that the full-sample loadings could reflect an in-sample-optimal combination of the factors rather than a return an investor could have earned in real time. Table 33 addresses this directly. At each evaluation date we re-estimate the principal components *and* the spanning betas on a past-only window—expanding (initialised at 60 months) as the baseline and a 60-month rolling window as a robustness check—and define the out-of-sample alpha as the mean of the hedged return, tested with the Newey and West [1994] HAC correction. Because the components are themselves re-estimated within each window, the design is invariant to the sign and rotation indeterminacy of principal components. The out-of-sample alpha remains [positivo/significativo]: +XX% p.a. (BTP Italia), +XX% (CDS–Bond), +XX% (iTraxx Combined) under the expanding window, broadly comparable to—and modestly below—the in-sample estimates of Table 31, confirming that the alpha is not an artefact of full-sample factor estimation.

Because PCA is an unsupervised method—it extracts factors that maximise variance in the candidate set, not factors that explain variation in strategy returns—it cannot identify strategy-specific risk drivers. This motivates the supervised penalised regression of the next subsection.

4.3 Adaptive Elastic Net

Preprocessing. The 85 candidate factors described in Section 4.1 are preprocessed for penalised regression. Stationarity is verified via augmented Dickey–Fuller tests, and non-stationary series are

first-differenced. Following the normalisation convention of Zou and Zhang [2009], each factor is standardised to unit L_2 -norm rather than unit variance, which ensures that the penalisation treats all factors symmetrically regardless of scale. Near-duplicate pairs with absolute correlation exceeding 0.95 are identified and the member of each pair with the lower marginal correlation to the strategy excess return is dropped, reducing the effective dimension to approximately 82 factors per strategy. Even after this reduction, the space of possible factor subsets— $2^{82} \approx 4.8 \times 10^{24}$ —rules out exhaustive search and motivates the use of penalised estimation methods with formal variable selection guarantees.

Estimation. The Adaptive Elastic Net of Zou and Zhang [2009] proceeds in two stages. In Stage 1, a standard Elastic Net is estimated by solving

$$\hat{\beta}^{\text{EN}} = \arg \min_{\beta} \left\{ \frac{1}{T} \sum_{t=1}^T (r_t - \alpha - \beta' \mathbf{f}_t)^2 + \lambda_2 \|\beta\|_2^2 + \lambda_1 \|\beta\|_1 \right\}, \quad (13)$$

which combines an ℓ_2 ridge penalty that handles collinearity with an ℓ_1 LASSO penalty that induces sparsity. The Elastic Net produces a preliminary estimate $\hat{\beta}^{\text{EN}}$ that is used to construct adaptive weights $w_j = |\hat{\beta}_j^{\text{EN}}|^{-\gamma}$ with $\gamma = 1$. Factors with large Stage 1 coefficients receive small weights (weak penalisation), while factors with near-zero Stage 1 coefficients receive large weights (strong penalisation), so the adaptive weighting concentrates the penalty on the irrelevant factors and relaxes it on the relevant ones. In Stage 2, the Adaptive Elastic Net re-solves with factor-specific ℓ_1 penalties:

$$\hat{\beta}^{\text{AEN}} = \arg \min_{\beta} \left\{ \frac{1}{T} \sum_{t=1}^T (r_t - \alpha - \beta' \mathbf{f}_t)^2 + \lambda_2 \|\beta\|_2^2 + \lambda_1 \sum_{j=1}^p w_j |\beta_j| \right\}. \quad (14)$$

Both stages use the same λ_2 ; the pair (λ_1, λ_2) is selected by minimising the Hannan–Quinn information criterion (HQC), and the solution path is computed via coordinate descent [Friedman et al., 2010].

The AEN is preferred to the standard LASSO for two reasons that are directly relevant to our setting. First, the LASSO does not possess the oracle property: Zou and Zhang [2009] show that under conditions on the design matrix that are routinely violated when candidate factors are correlated—precisely the situation in our data—the LASSO can be inconsistent for variable selection. The adaptive weighting restores consistency and delivers the oracle property, meaning that the AEN performs asymptotically as if the true set of relevant factors were known in advance. Second, the ℓ_2 ridge component handles the grouping effect: when several candidate factors are highly correlated (as are, for example, the various credit spread measures in our universe), the LASSO arbitrarily selects one and drops the rest, while the Elastic Net tends to select or drop them together, producing more stable and economically interpretable factor sets.

Stability Selection. Even with the oracle property, the factors selected by a single run of the AEN on a finite sample may be sensitive to small perturbations in the data. To address this selection uncertainty, we apply the stability selection framework of Meinshausen and Bühlmann [2010], adapted for time-series dependence. The procedure draws $B = 100$ stationary bootstrap resamples with an expected block length of 9 months, re-estimates the full AEN pipeline on each resample, and records which factors are selected. A factor is retained in the final “stable” set only if its selection frequency across the 100 bootstrap replicates exceeds $\hat{\pi}_j \geq 80\%$. Factors that are selected occasionally but not consistently are discarded as unreliable.

The key theoretical result of Meinshausen and Bühlmann [2010] is that stability selection controls the expected number of false selections: for a threshold $\pi^* \in (1/2, 1)$ and a maximum model size q per bootstrap replicate, the expected number of falsely selected variables is bounded by $q^2/[p(2\pi^* - 1)]$, where p is the total number of candidates. With $p \approx 82$ factors and a threshold of 80%, this bound is small, providing a formal guarantee against overfitting that no single-run LASSO or Elastic Net can offer. Because correlated factors may split bootstrap votes among themselves (each individually falling below the threshold even though the cluster as a whole is consistently selected), we additionally report cluster-level selection frequencies in Appendix A.7.2.

Post-Selection OLS. Once the stable factor set $\hat{S}_i$ has been identified for each strategy, we discard the penalised coefficient estimates and re-estimate the spanning regression by unrestricted OLS on the original (un-standardised) factors:

$$r_{i,t} = \alpha_i + \beta_i' \mathbf{f}_{\hat{S}_i,t} + \varepsilon_{i,t}, \quad (15)$$

with Newey–West HAC standard errors. This two-step procedure—penalised selection followed by unpenalised inference—avoids the well-known downward bias of penalised coefficient estimates while retaining the sparsity and interpretability of the selected model. The post-selection OLS regression is the basis for all alpha estimates, factor loadings, and $\bar{R}^2$ values reported in Table 34.

Results. Table 34 reports the post-selection OLS estimates. Alpha is statistically significant at the one-percent level for BTP Italia (+1.39% p.a.) and CDS–Bond Basis (+3.94% p.a.); for iTraxx Combined, the alpha is positive but not statistically significant (+0.47%, $t = 1.04$), suggesting that the skew strategy’s excess returns are largely explained by the selected factors. The models are parsimonious—six, eight, and five factors respectively—yet deliver adjusted $\bar{R}^2$ of 21.7%, 26.7%, and 23.6%, substantially higher than either the benchmark frameworks of Section 3 or the PCA spanning regressions above.

The selected factors are economically interpretable and differ across strategies in ways that reflect their distinct market microstructures. For BTP Italia, the stable set includes the Amihud and Mendelson [2015] illiquidity measure, the Libor–OIS spread, the 10-year swap spread, the CDX Investment Grade index, and two equity style factors (SMB, momentum)—a combination that points to liquidity, funding conditions, and interest rate exposures as the primary systematic drivers of the inflation-linked basis trade. For CDS–Bond, the selected factors include financial uncertainty (ΔUF , Ludvigson et al., 2021), European investment-grade bond returns (RI_EU), emerging-market currency risk, the 1-year European prime broker CDS spreads (PB_EU_CDS_1Y), the 5-year swap spread, U.S. credit spread changes, the option-implied equity premium (EP_SVIX, Martin, 2017), and the currency trend-following factor—consistent with a trade whose profitability

depends on dealer balance-sheet capacity, funding conditions, and cross-market risk appetite. For iTraxx Combined, the dominant factors are the option-implied equity premium (EP_SVIX), the Libor–OIS spread, European equity volatility ($\Delta V2X$), real uncertainty (ΔUR , Ludvigson et al., 2021), and U.S. Treasury settlement fails—reflecting the index skew’s sensitivity to the aggregate state of intermediary capital and market-wide uncertainty rather than to idiosyncratic credit events. Each strategy thus loads on a distinct combination of risk drivers, consistent with the view that the alpha documented in Section 3 and in this section is not compensation for a common systematic factor: each strategy is exposed to a different mix of identified risks, and the residual premium that survives these controls is strategy-specific in origin. To verify that the results do not depend on the choice of penalised estimator, we re-estimate the spanning regression using four alternative methods—LASSO, standard Elastic Net, Adaptive LASSO, and the ALASSO-Chen variant of Chen et al. [2025]—and report the comparison in Section 4.5. The full list of factors selected by each method is reported in Appendix A.7.2.

4.4 Conditional Alpha

Table 35 decomposes alpha by stress regime under both the PCA and AEN factor specifications. The CDS–Bond basis exhibits the sharpest contrast: under AEN factors, alpha rises from +4.15% (NORMAL) to +8.17% (HIGH), nearly doubling while remaining significant at the one-percent level in both regimes; the PCA specification delivers a similar pattern (+4.84% to +9.13%). For BTP Italia, the HIGH-regime alpha is substantially larger under both PCA (+5.37%, significant at the five-percent level) and AEN (+2.71%), but the AEN estimate does not reach statistical significance ($t = 1.43$), reflecting the smaller number of HIGH-stress months in the shorter BTP Italia sample. For iTraxx Combined, the alpha is significant in both regimes under PCA factors (+3.38% in HIGH, +1.82% in NORMAL), while under AEN factors it is significant only during HIGH stress (+1.51%) and turns negative and insignificant during NORMAL months—indicating that the AEN factors absorb the skew’s unconditional return but not the stress-related component. This regime dependence is consistent with the slow-moving capital framework of Duffie [2010]: when

intermediary balance sheets are under pressure, mispricings widen and the residual excess return rises. Detailed Christopherson et al. [1998] conditional alpha estimates and threshold robustness checks (at 80, 100, and 120 bps) are reported in Appendix A.6.3 and A.7.3. Sub-period stability of the alpha estimates under both PCA and AEN is confirmed in Appendix A.6.2, and rolling 36-month alpha estimates with regime shading are plotted in Appendix A.7.4.

4.5 Discussion, Method Comparison, and Robustness

Table 36 compares alpha and model dimension across five machine learning methods: LASSO, Elastic Net, Adaptive LASSO, AEN (our baseline), and ALASSO-Chen [Chen et al., 2025]. The first four methods use HQC for hyperparameter selection; ALASSO-Chen uses AIC, which imposes a lighter penalty and therefore selects substantially more factors. For BTP Italia, the LASSO, Elastic Net, and Adaptive LASSO converge to the same seven-factor model under HQC tuning; the AEN stability selection drops one factor, yielding a six-factor specification with nearly identical alpha (+1.39% p.a., significant at the one-percent level). For CDS–Bond Basis, alpha ranges from +3.91% to +4.97% across methods, always significant at the one-percent level. For iTraxx Combined, the alpha is insignificant across all methods ($t = 1.04$ – 1.68), confirming that the skew strategy’s excess returns are absorbed by the selected risk factors—a finding that is itself informative, as it indicates that the iTraxx skew does not offer a genuine risk-adjusted premium beyond its systematic exposures. The ALASSO-Chen specification selects 51 factors for BTP Italia and raises $\bar{R}^2$ to 48%, but the alpha loses statistical significance ($t = 1.62$), suggesting overfitting: the additional factors absorb noise rather than genuine risk exposures. AEN is preferred because it combines the oracle property of Zou and Zhang [2009] with the stability selection filter of Meinshausen and Bühlmann [2010], producing the sparsest model that retains only factors with robust bootstrap selection frequencies. Sensitivity of the factor selection to the correlation pre-cleaning threshold (ρ) and the adaptive weight exponent (γ) is verified in Appendix A.7.1: in all configurations, the AEN selects the same core factors with identical post-selection OLS results.

Having established that alpha survives both diffuse (PCA) and targeted (AEN) factor controls for the BTP Italia and CDS–Bond strategies, that the iTraxx skew is absorbed by its selected risk factors, and that the residual alpha rises during stress in a manner consistent with Duffie [2010], we turn next to the question of whether the three mispricings are linked to one another and, if so, what economic mechanism connects them.

5 Cross-Strategy Interdependencies

The preceding sections have established that all three strategies generate alpha that survives both classical benchmark and data-driven factor controls, and that this alpha rises during periods of intermediary funding stress. A natural question follows: are the three mispricings linked to one another, and if so, what economic mechanism ties them together? The slow-moving capital literature introduced in Section 1—Duffie [2010], Brunnermeier and Pedersen [2009], Mitchell and Pulvino [2012], Boyarchenko et al. [2016], and Fleckenstein et al. [2014]—yields five testable predictions about how arbitrage mispricings behave when they draw on a common pool of slow-moving intermediary capital: (P1) mispricings serviced by a common pool of intermediary capital should co-move [Duffie, 2010]; (P2) the co-movement should intensify when that capital is scarce [Duffie, 2010, Brunnermeier and Pedersen, 2009]; (P3) mispricing changes in one strategy should contain information for predicting those of others, with effects propagating asymmetrically along the liquidity gradient [Mitchell and Pulvino, 2012, Fleckenstein et al., 2014]; (P4) the persistence of mispricing levels should reflect the funding-dependence of each arbitrage trade, with strategies that require continuous balance-sheet funding displaying the longest half-lives [Boyarchenko et al., 2016]; and (P5) a common factor extracted from the co-movement should be explained by intermediary balance-sheet variables rather than by macroeconomic fundamentals [Duffie, 2010, Brunnermeier and Pedersen, 2009].

A feature of our three-strategy setting gives these tests their identifying power. All three arbitrages are run by the same type of leveraged institutional arbitrageur, but the instruments differ in who holds them: the CDS–Bond and iTraxx legs trade among institutions, whereas the BTP Italia

inflation-linked bond is held predominantly by retail investors. This heterogeneity in the holder base is the source of identification we exploit throughout—it lets us separate the slow-moving-capital mechanism, under which co-movement depends on sharing a capital pool, from a common macroeconomic driver, under which all three mispricings would move together regardless of who holds the instruments.

We devote one subsection to each prediction, stating in each case the prediction, the test, and the verdict where the evidence is presented. Section 5.1 first defines the mispricing series and the stress regimes. Sections 5.2–5.6 then test P1 through P5 in turn. Section 5.7 interprets the evidence as a whole, explains why these arbitrage opportunities persist in post-crisis European credit markets, and collects the verdicts in a summary mapping.

5.1 Mispricing Construction and Stress Regimes

Following Duffie [2010], we work with mispricing magnitudes $m_{i,t}$ rather than strategy returns, because it is the level of the mispricing—the distance of the basis from its no-arbitrage value—that measures the degree to which capital is failing to close the gap. For each strategy, $m_{i,t}$ is defined as the absolute value of the basis, sign-adjusted so that an increase $\Delta m > 0$ denotes a widening of the mispricing for all three strategies. The BTP Italia mispricing is the cross-sectional mean of the absolute basis across all outstanding issues with at least six months to maturity; the CDS–Bond mispricing is the cross-sectional mean of the absolute value of all negative single-name bases in the iTraxx Main universe on each day; and the iTraxx Combined mispricing is the cross-sectional mean of the absolute skew across the four on-the-run sub-indices. All three series are aggregated to monthly frequency. Market conditions are classified into three regimes—LOW, MEDIUM, and HIGH—based on the iTraxx Europe Main 5-year spread, with thresholds at 60 and 100 basis points. These are the same thresholds that define the stress-dependent transaction-cost tiers applied to every return series in the paper (Appendix A.2.3) and that demarcate the HIGH stress regime in the benchmark and conditional-alpha analyses of Sections 3 and 4.

5.2 P1: Cross-Market Co-Movement

The first prediction is that mispricings serviced by a common pool of intermediary capital should co-move: when the capital that closes one arbitrage is withdrawn, the same withdrawal leaves the others unhedged, so their mispricings widen together. We test this with pairwise Pearson correlations of mispricing changes Δm , computed unconditionally and within each stress regime (Table 37, Panel B).

The prediction holds for the pair whose underlying instruments are all institutionally held and traded—the CDS–Bond basis and the iTraxx skew, hereafter the “institutional pair.” Their mispricing changes correlate at +0.19 unconditionally, significant at the ten percent level, and the correlation rises monotonically across regimes from -0.21 in the LOW regime to $+0.32$ in the HIGH regime, with the LOW-to-HIGH movement significant at the one percent level under a Fisher z -test. The Forbes and Rigobon [2002] correction for the heteroskedastic increase in mispricing volatility under stress attenuates the shift but leaves it significant at the ten percent level, so the intensification reflects a genuine increase in the cross-market link rather than the mechanical inflation of correlations that volatility differences alone would produce. When dealer balance sheets are healthy the two markets price almost independently; when balance sheets are depleted, common funding shocks force the simultaneous unwinding of positions in both, and the link turns strongly positive. P1 is confirmed for the institutional pair.

The BTP Italia pairs move in the opposite direction, and it is worth being explicit that this does not contradict P1. The CDS–Bond versus BTP Italia correlation is -0.30 unconditionally, significant at the one percent level, and deepens from -0.13 in LOW to -0.47 in HIGH, the latter significantly different from zero at the five percent level. The BTP Italia versus iTraxx correlation is negative but indistinguishable from zero on the full sample (-0.10), and reverses sign across regimes, from $+0.15$ in LOW to -0.14 in HIGH. The dynamic conditional correlations estimated by a bivariate GARCH–DCC model (Appendix A.9.2, Figure 7) show that the negative sign of both BTP Italia pairs is a persistent feature of the data—the conditional correlation between CDS–Bond and BTP Italia remains below zero in every month of the sample, with mean -0.29 , and

that between BTP Italia and iTraxx is negative throughout, with mean -0.10 —not an artefact of averaging over offsetting positive and negative episodes.

P1 does not predict that all mispricings co-move—only that those drawing on a common capital pool do. The BTP Italia inflation-linked bond is retail-held and does not draw on the institutional pool, so P1 predicts no positive co-movement, and there is none. This is exactly the holder-base heterogeneity that identifies the mechanism: the sign of the co-movement turns on whether the capital pool is shared—positive within the institutional pool, negative against the retail-held basis—rather than being common to all three, as a macroeconomic driver would require. We take up the economics of the negative sign in Section 5.7.

5.3 P2: Stress Amplification

The second prediction sharpens the first: the co-movement should not merely be present but should intensify as intermediary capital becomes scarce. The regime correlations of the previous subsection already provide one reading—co-movement within the institutional pair rises from -0.21 in LOW to $+0.32$ in HIGH—but they treat stress as a discrete state and cannot say whether transmission rises continuously with it, nor separate the incremental effect of stress from the average level of transmission. We therefore test the prediction parametrically with an interaction regression:

$$\Delta m_{i,t} = \alpha + \beta_0 \Delta m_{j,t} + \beta_1 (\Delta m_{j,t} \times z_t) + \gamma z_t + \varepsilon_t, \quad (16)$$

where z_t is the standardised iTraxx Main 5-year spread, a proxy for intermediary funding stress. The coefficient β_1 captures the incremental transmission during stress: $\beta_1 > 0$ means that a given shock to Δm_j has a larger effect on Δm_i when balance sheets are strained. The theory is explicit on the mechanism: Duffie [2010] notes that a depleted intermediary is “less able or predisposed to absorb supply and demand shocks,” and Brunnermeier and Pedersen [2009] show that the destabilising effects of funding constraints are “largest when traders are near their constraint.”

Table 39 reports the results. Within the institutional pair, the interaction coefficient on iTraxx mispricing carries the predicted positive sign, and the direct effect of stress on the CDS–Bond basis is significant in both interaction regressions, at the one percent level ($\gamma = +1.60$) when BTP Italia mispricing is the explanatory variable and at the five percent level ($\gamma = +0.95$) when iTraxx mispricing is. Three pieces of evidence converge to support P2 within the institutional pair. First, the regime correlation rises monotonically across stress states (Section 5.2). Second, the Forbes–Rigobon correction confirms that this rise reflects a genuine increase in the cross-market link rather than a heteroskedasticity artefact. Third, the within-pair interaction term β_1 is correctly signed and, in the CDS–Bond $\rightarrow$ iTraxx direction, significant at the ten percent level ($\beta_1 = +0.08$). All three readings point the same way, and P2 is confirmed for the institutional pair. For the pairs involving BTP Italia the interaction coefficient is significantly negative ($\beta_1 = -0.64$ for CDS–Bond $\rightarrow$ BTP Italia, significant at the one percent level): stress widens, rather than amplifies, the link between the institutional pair and the retail-held basis—the mirror image of P2 that the holder-base heterogeneity of Section 5.7 explains.

5.4 P3: Lead–Lag Transmission

The third prediction concerns timing: if capital moves slowly across markets, a mispricing change in one strategy should help predict future changes in another, and the transmission should run from more liquid to less liquid markets as capital is reallocated across the dealer network. We test predictability with spanning regressions in the design of Fleckenstein et al. [2014],

$$\Delta m_{i,t} = \alpha + \sum_{j \neq i} \sum_{l=0}^L \beta_{j,l} \Delta m_{j,t-l} + \varepsilon_t, \quad (17)$$

with $L = 2$ lags and Newey–West HAC standard errors. The objects of interest are the lagged coefficients: a significant loading on $\Delta m_{j,t-1}$ means that last month’s mispricing in strategy j predicts this month’s in strategy i , a genuine lead–lag relation rather than contemporaneous co-movement.

Table 38 reports the results. The significant coefficients trace a liquidity pecking order: one-month-lagged iTraxx changes predict the widening of the CDS–Bond mispricing with coefficient $+0.71$ at the ten percent level, while CDS–Bond changes enter the BTP Italia equation contemporaneously with coefficient -0.42 at the five percent level. The two institutional mispricings thus lead the BTP Italia basis with a delay of zero to one month, the timing dictated by the speed at which dealers can rotate inventory across markets.

The spanning regressions establish predictability but not direction. To fix the direction we embed the three Δm series in a first-order VAR and test for Granger causality. The hierarchy is sharp. The iTraxx index Granger-causes both the CDS–Bond basis ($p = 0.011$) and the BTP Italia basis ($p = 0.007$), and the link survives at two and three lags. Neither direction reverses: BTP Italia does not Granger-cause either institutional mispricing at conventional levels, nor does CDS–Bond Granger-cause iTraxx. The directional link from iTraxx to the CDS–Bond basis strengthens to the one-percent level on the longer bivariate sample available for those two series alone. Shocks thus originate in the most liquid, standardised index market, where dealer inventory adjusts fastest, and reach both the less liquid cash basis and the retail-held inflation-linked basis with a delay. The generalised impulse responses of Pesaran and Shin [1998], invariant to the ordering of variables in the VAR, are consistent with the same picture (Figure 14). A one-standard-deviation shock to the iTraxx index produces a positive response in the CDS–Bond basis that peaks at one month, and a negative response in the BTP Italia basis that peaks at one month; both decay within two to three months, with confidence bands that narrow rapidly toward zero. Shocks in the reverse direction—originating in either of the less liquid bases and propagating back to the iTraxx skew—are indistinguishable from zero across the full horizon, with bootstrap bands that contain zero throughout. The asymmetric pattern matches the forced-deleveraging mechanism of Mitchell and Pulvino [2012], by which intermediaries adjust their most liquid positions first when funding contracts and their less liquid positions only with a delay. P3 is confirmed: predictability, direction, and dynamic response all run from the liquid index to the less liquid bases. Robustness to lag length ($p = 1, 2, 3$) appears in Appendix A.9.3.

5.5 P4: Persistence and Funding-Dependence

The fourth prediction links the persistence of each mispricing to how funding-dependent its trade is: the more a strategy relies on continuously financing a cash position on a dealer balance sheet, the longer a shock to its mispricing should take to decay, because the capital that closes it arrives only slowly. Duffie [2010] makes this quantitative—the half-life of the reversal that follows a shock is decreasing in the mean rate at which arbitrage capital reaches the trade—so the cross-strategy ordering of half-lives should track the ordering of funding frictions. We estimate first-order autocorrelations of the mispricing levels,

$$m_{i,t} = c_i + \phi_i m_{i,t-1} + u_{i,t}, \quad (18)$$

with Newey–West standard errors (six monthly lags), and read off the implied half-life $h_i^* = \ln(0.5)/\ln(\phi_i)$.

Table 40 reports the estimates. The CDS–Bond basis is the most persistent ($\phi = 0.79$, half-life ≈ 3.0 months), with the BTP Italia and the iTraxx skew essentially identical at $\phi \approx 0.73$ and half-life ≈ 2.2 months, all three coefficients significant at the one percent level. The CDS–Bond basis requires the arbitrageur to hold the cash bond on a repo book and finance it continuously, so a widening persists until fresh balance-sheet capacity arrives—the longest half-life of the three, roughly one-third longer than its counterparts. The two remaining mispricings do not separate: the iTraxx skew is a pure derivative position, but the on-balance-sheet exposure of trading the index against its constituents carries non-trivial repo and counterparty costs that the half-life of the basis itself reflects; the BTP Italia is refreshed by primary auctions reserved to households and stabilised by the loyalty bonus, both of which slow its mean reversion below what a pure-derivative basis would otherwise deliver. P4 is therefore confirmed at the upper end of the funding-intensity spectrum—the cash-bond basis is the most persistent—but not in the finer ordering below it: the differential between the two less funding-intensive strategies is too small to register in monthly

persistence. The skew would, on funding grounds alone, be expected to revert fastest, yet its half-life matches the BTP Italia's.

The impulse responses corroborate this reading from a second angle. The own-responses on the diagonal of Figure 14 decay within the first month: innovations to Δm dissipate quickly precisely because the level m mean-reverts over a multi-month horizon, at a speed governed by funding-dependence. This is the “sharp reaction ... and a subsequent and more extended reversal” that Duffie [2010] identifies as the signature of slow-moving capital: the impact is absorbed within a month, while the level reverts over the horizon set by the rate of capital arrival.

5.6 P5: The Common Factor

The fifth prediction concerns the source of the co-movement. If slow-moving intermediary capital is the binding constraint, the common factor behind the three mispricings should load on intermediary balance-sheet variables and not on macroeconomic fundamentals; Duffie [2010] identifies “the level of capital commonly available to bear losses” as the link tying mispricings across markets, and Brunnermeier and Pedersen [2009] tie the commonality of liquidity to the speculators’ shadow cost of capital. We extract the first principal component (PC1) of the three Δm series and regress it on two sets of variables (Table 41): intermediary stress proxies, and a placebo set of macroeconomic indicators that should not explain PC1 if the slow-moving capital interpretation is correct.

The contrast is sharp. In Panel A, one proxy per channel of the intermediary literature—the He et al. [2017] capital ratio (−4.29, significant at the five percent level), the Libor–OIS spread (+3.34, significant at the one percent level), the European prime-broker CDS spread (−0.018, significant at the five percent level), and the Amihud and Mendelson [2015] illiquidity measure (+1.73, significant at the one percent level)—all carry the theoretically predicted sign, so that a fall in intermediary capital, a rise in funding costs, and an increase in market illiquidity all predict wider mispricings; the model explains 23.1% of the variation in the common factor, jointly significant at the one percent level, and Panel B confirms the same picture with an extended proxy

set ($\bar{R}^2 = 23.5\%$, jointly significant at the one percent level). Panel C runs the placebo: regressed on macroeconomic and real-activity uncertainty, long-term inflation expectations, and economic policy uncertainty, the common factor yields four uniformly insignificant coefficients, a model $\bar{R}^2$ that falls to 11.2%, and a joint F -statistic that fails to reject at conventional levels. A factor loading on macroeconomic fundamentals would have reversed this ranking. P5 is confirmed: the force driving the three mispricings to co-move is the state of intermediary balance sheets, not of the macroeconomy.

5.7 Discussion: Why These Arbitrages Persist

Holder-Base Heterogeneity as the Source of Identification. The evidence of the preceding subsections is unified by the single structural feature introduced at the outset: although all three strategies are run by the same type of institutional arbitrageur, the instruments differ in who holds them, and it is this heterogeneity that identifies the mechanism. The CDS–Bond basis and the iTraxx skew trade in instruments whose holder base is overwhelmingly institutional—hedge funds, dealer proprietary desks, and leveraged accounts—so that a spike in volatility inflates the Value-at-Risk of the combined position, exhausts the desk’s risk budget, and forces simultaneous liquidation across both books [Brunnermeier and Pedersen, 2009], pushing both bases wider at once. The BTP Italia is different: only its inflation-linked leg is structurally retail-dominated. Introduced in 2012 with a primary auction reserved to households and a loyalty bonus (*premio fedeltà*) for buy-and-hold investors, and serving as the simplest inflation hedge available to an Italian household, the bond is held by preferred-habitat investors in the sense of Vayanos and Vila [2021]: they use no leverage, face no margin calls, and do not unwind when dealer capital is scarce. The arbitrageur of the BTP Italia basis is itself institutional and faces the same funding environment as the others—so the strategy is fully part of the experiment, not a bystander to it—but the absence of forced selling on the bondholder side keeps the basis out of the institutional fire sales, so that stress drives it opposite to the institutional pair rather than alongside it. This is what gives the design its identifying power: a common macroeconomic factor would move all three mispricings in the

same direction regardless of who holds the instruments, whereas the sign of the co-movement turns entirely on whether the holder base is shared—the pattern documented under P1 and P2.

Why the Mispricings Persist. Why have these opportunities not been competed away? The answer lies in structural limits to arbitrage that differ across strategies but share a common root in the scarcity of intermediary capital. For the BTP Italia, the binding constraint is the segmentation itself: institutional arbitrageurs can exploit the basis, but the retail holders who form the natural counterparty neither monitor nor arbitrage it [Vayanos and Vila, 2021]. For the CDS–Bond basis, the limits are those of Bai and Collin-Dufresne [2019] and Mitchell and Pulvino [2012]: the trade must fund the corporate bond on a dealer balance sheet, the repo market for European investment-grade corporates is thin and subject to recall risk, and the counterparty risk on the CDS leg consumes regulatory capital—so capacity is procyclical and the mispricing countercyclical, widening precisely when the trade is most attractive but least feasible. For the iTraxx skew, the constraint is post-crisis regulation: Boyarchenko et al. [2016] show that the supplementary leverage ratio and other Basel III requirements raised the capital charge on index-versus-constituent trades—as required leverage ratios rose from roughly 2–3% pre-crisis to 5–6% post-crisis, the return on equity fell by a factor of two to three—rendering them “significantly less economical” and producing “potentially ‘new normal’ levels for both the CDS-bond and the CDX-CDS bases.” The persistence of alpha documented in Sections 3 and 4 is therefore not a puzzle calling for a risk-based explanation; it is the equilibrium of a market in which the cost of intermediary capital has been permanently raised, and slow-moving capital, in the sense of Duffie [2010], keeps mispricings open longer than frictionless models would predict [Mitchell et al., 2007].

Summary: The Five Predictions.

P1. Mispricings drawing on a common capital pool co-move. The institutional pair correlates at +0.19, while the retail-held BTP Italia, outside that pool, correlates significantly negatively with the CDS–Bond basis at –0.30 unconditionally (Table 37). The sign turns on the holder base, exactly as the common-capital mechanism requires. Confirmed.

- P2.** The co-movement intensifies in stress. Institutional co-movement rises monotonically from -0.21 in LOW to $+0.32$ in HIGH, a LOW-to-HIGH increase significant at the one percent level under a Fisher z -test (Tables 37, 39). Confirmed.
- P3.** Mispricing changes transmit along the liquidity gradient. The iTraxx index Granger-causes both other mispricings ($p = 0.007$ and 0.011), with impulse responses tracing the same liquid-to-illiquid path (Table 38, Figure 14). Confirmed.
- P4.** Persistence reflects funding-dependence. The CDS–Bond basis, the only strategy requiring continuous cash-bond financing on a dealer balance sheet, has a half-life about one-third longer than the iTraxx skew and the BTP Italia (Table 40); the two less funding-intensive strategies are essentially identical. Partially confirmed.
- P5.** The common factor is intermediary capital, not macro fundamentals. PC1 loads on balance-sheet proxies with the predicted signs ($\bar{R}^2 = 23.1\%$, jointly significant at the one percent level) but not on a macroeconomic placebo set ($\bar{R}^2 = 11.2\%$, fails to reject at conventional levels) (Table 41). Confirmed.

Four of the five predictions find unambiguous support in the data; only P4—the cross-strategy ordering of persistences—is partially supported, holding at the upper end of the funding-intensity spectrum but not within the institutional pair. Taken together, the evidence locates the geography of slow-moving capital not in asset-class labels but in the network of intermediaries that service these markets—and ultimately in who holds the bonds.

Having documented both the existence of significant alpha across three distinct European fixed-income arbitrage strategies and the economic mechanism—slow-moving intermediary capital, amplified by post-crisis regulation and modulated by investor-base segmentation—that sustains these mispricings, we draw our conclusions in Section 6.

6 Conclusion

We document positive and economically significant alpha in three European fixed-income arbitrage strategies—the BTP Italia inflation-linked basis, the corporate CDS–bond basis, and the iTraxx index skew—and show that this alpha is not compensation for the identified systematic risk exposures, whether the controls are entered as classical benchmark factors or as data-driven selections from a large candidate universe, and it survives when each benchmark hedge is re-estimated out of sample from past-only loadings. The mispricings are linked by a common factor that loads on intermediary balance-sheet variables (and not on macroeconomic fundamentals), and their persistence at the level of each series reflects the funding-dependence of the arbitrage trade: the CDS–Bond basis, the only strategy requiring continuous repo financing of the cash bond, is the most persistent, while the BTP Italia and the iTraxx skew mean-revert more rapidly and are essentially identical to each other. The three correlations exhibit an asymmetric pattern that mirrors the structural divide between institutional and retail holder bases: positive and rising in stress within the institutional pair (CDS–Bond and iTraxx), and negative across holder bases (BTP Italia vs. either institutional strategy). The slow-moving capital literature [Duffie, 2010, Brunnermeier and Pedersen, 2009, Mitchell and Pulvino, 2012], combined with the post-crisis regulatory environment documented by Boyarchenko et al. [2016], provides an economic explanation for why these mispricings persist.

The findings have implications for two strands of the literature. For the limits-to-arbitrage literature, the European setting provides a natural test of the slow-moving capital mechanism in a market with a structurally heterogeneous holder base, and the investor-base-segmentation pattern in cross-strategy correlations offers independent evidence that the geography of arbitrage capital—rather than asset-class labels—determines cross-market co-movement. For the fixed-income asset pricing literature, supervised selection on a large candidate universe recovers strategy-specific risk drivers that the standard benchmark factor sets miss, yet economically significant alpha persists

for the BTP Italia and CDS–Bond strategies—indicating that the residual return is genuinely unspanned rather than compensation for an overlooked systematic factor.

References

- Yakov Amihud and Haim Mendelson. The pricing of illiquidity as a characteristic and as risk. *Multinational Finance Journal*, 19(3):149–168, 2015.
- Clifford S Asness, Tobias J Moskowitz, and Lasse Heje Pedersen. Value and momentum everywhere. *The journal of finance*, 68(3):929–985, 2013.
- Clifford S Asness, Andrea Frazzini, and Lasse Heje Pedersen. Quality minus junk. *Review of Accounting Studies*, 24(1):34–112, 2019.
- Attakrit Asvanunt and Scott Richardson. The credit risk premium. *Journal of Fixed Income*, 26(3):6–24, 2017.
- Jennie Bai and Pierre Collin-Dufresne. The cds-bond basis. *Financial Management*, 48(2):417–439, 2019.
- Jushan Bai and Serena Ng. Determining the number of factors in approximate factor models. *Econometrica*, 70(1):191–221, 2002.
- Jushan Bai and Serena Ng. Confidence intervals for diffusion index forecasts and inference for factor-augmented regressions. *Econometrica*, 74(4):1133–1150, 2006.
- Dion Bongaerts, Frank De Jong, and Joost Driessen. Derivative pricing with liquidity risk: Theory and evidence from the credit default swap market. *Journal of Finance*, 66(1):203–240, 2011.
- Nina Boyarchenko, Pooja Gupta, Nick Steele, and Jacqueline Yen. Trends in credit market arbitrage. Technical report, Staff Report, 2016.
- Douglas T. Breeden, Michael R. Gibbons, and Robert H. Litzenberger. Empirical test of the consumption-oriented CAPM. *Journal of Finance*, 44(2):231–262, 1989.
- Jordan Brooks, Tony Gould, and Scott A Richardson. Active fixed income illusions. *Journal of Fixed Income*, 29(4), 2020.
- Markus K Brunnermeier and Lasse Heje Pedersen. Market liquidity and funding liquidity. *The review of financial studies*, 22(6):2201–2238, 2009.

- Mark M Carhart. On persistence in mutual fund performance. *The Journal of finance*, 52(1):57–82, 1997.
- Peter Carr and Liuren Wu. Variance risk premiums. *Review of Financial Studies*, 22(3):1311–1341, 2009.
- Scott Cederburg, Michael S O’Doherty, Feifei Wang, and Xuemin Sterling Yan. On the performance of volatility-managed portfolios. *Journal of financial Economics*, 138(1):95–117, 2020.
- Yong Chen, Sophia Zhengzi Li, Yushan Tang, and Guofu Zhou. Anomalies as new hedge fund factors. *Journal of Financial and Quantitative Analysis*, 60(8):3660–3693, 2025.
- Jon A Christopherson, Wayne E Ferson, and Debra A Glassman. Conditioning manager alphas on economic information: Another look at the persistence of performance. *The Review of Financial Studies*, 11(1):111–142, 1998.
- Gregory Connor and Robert A Korajczyk. Performance measurement with the arbitrage pricing theory: A new framework for analysis. *Journal of Financial Economics*, 15(3):373–394, 1986.
- Gregory Connor and Robert A Korajczyk. Risk and return in an equilibrium apt: Application of a new test methodology. *Journal of Financial Economics*, 21(2):255–289, 1988.
- Martijn Cremers, Matthias Fleckenstein, and Priyank Gandhi. Treasury yield implied volatility and real activity. *Journal of Financial Economics*, 140(2):412–435, 2021.
- Jefferson Duarte, Francis A Longstaff, and Fan Yu. Risk and return in fixed-income arbitrage: Nickels in front of a steamroller? *The Review of Financial Studies*, 20(3):769–811, 2007.
- Darrell Duffie. Presidential address: Asset price dynamics with slow-moving capital. *The Journal of finance*, 65(4):1237–1267, 2010.
- Bjørn Eraker and Yue Wu. Explaining the negative returns to volatility claims: An equilibrium approach. *Journal of Financial Economics*, 125(1):72–98, 2017.
- Eugene F Fama and Kenneth R French. Common risk factors in the returns on stocks and bonds. *Journal of financial economics*, 33(1):3–56, 1993.

- Eugene F Fama and Kenneth R French. International tests of a five-factor asset pricing model. *Journal of financial Economics*, 123(3):441–463, 2017.
- Matthias Fleckenstein, Francis A Longstaff, and Hanno Lustig. The tips-treasury bond puzzle. *the Journal of Finance*, 69(5):2151–2197, 2014.
- Kristin J Forbes and Roberto Rigobon. No contagion, only interdependence: measuring stock market comovements. *The journal of Finance*, 57(5):2223–2261, 2002.
- Andrea Frazzini and Lasse Heje Pedersen. Betting against beta. *Journal of financial economics*, 111(1):1–25, 2014.
- Jerome H Friedman, Trevor Hastie, and Rob Tibshirani. Regularization paths for generalized linear models via coordinate descent. *Journal of statistical software*, 33:1–22, 2010.
- William Fung and David A Hsieh. The risk in hedge fund strategies: Theory and evidence from trend followers. *The review of financial studies*, 14(2):313–341, 2001.
- William Fung and David A Hsieh. Hedge fund benchmarks: A risk-based approach. *Financial Analysts Journal*, 60(5):65–80, 2004.
- Stefano Giglio and Dacheng Xiu. Asset pricing with omitted factors. *Journal of Political Economy*, 129(7):1947–1990, 2021.
- Maximilian Glück, Benjamin Hübel, and Hendrik Scholz. Currency conversion of fama–french factors: How and why. *Journal of portfolio management*, 47(2):157–175, 2021.
- William Goetzmann, Jonathan Ingersoll, Matthew Spiegel, and Ivo Welch. Portfolio performance manipulation and manipulation-proof performance measures. *The Review of Financial Studies*, 20(5):1503–1546, 2007.
- Bing Han and Yi Zhou. Understanding the term structure of credit default swap spreads. *Journal of Empirical Finance*, 31:18–35, 2015.
- Bing Han, Avanidhar Subrahmanyam, and Yi Zhou. The term structure of credit spreads, firm fundamentals, and expected stock returns. *Journal of Financial Economics*, 124(1):147–171, 2017.

- Joseph Haubrich, George Pennacchi, and Peter Ritchken. Inflation expectations, real rates, and risk premia: Evidence from inflation swaps. *The Review of Financial Studies*, 25(5):1588–1629, 2012.
- Zhiguo He, Bryan Kelly, and Asaf Manela. Intermediary asset pricing: New evidence from many asset classes. *Journal of Financial Economics*, 126(1):1–35, 2017.
- Gur Huberman, Shmuel Kandel, and Robert F. Stambaugh. Mimicking portfolios and exact arbitrage pricing. *Journal of Finance*, 42(1):1–9, 1987.
- John Hull, Mirela Predescu, and Alan White. The relationship between credit default swap spreads, bond yields, and credit rating announcements. *Journal of banking & finance*, 28(11):2789–2811, 2004.
- Antti Ilmanen, Ronen Israel, Rachel Lee, Tobias J. Moskowitz, and Ashwin Thapar. How do factor premia vary over time? A century of evidence. *Journal of Investment Management*, 19(4), 2021.
- Benjamin Junge and Anders B Trolle. Liquidity risk in credit default swap markets. Research Paper 13-65, Swiss Finance Institute, 2015.
- Sven Klingler and Suresh Sundaresan. An explanation of negative swap spreads: Demand for duration from underfunded pension plans. *The Journal of Finance*, 74(2):675–710, 2019. doi: 10.1111/jofi.12750.
- Ralph S. J. Koijen, Tobias J. Moskowitz, Lasse Heje Pedersen, and Evert B. Vrugt. Carry. *Journal of Financial Economics*, 127(2):197–225, 2018.
- Robert B Litterman and José A Scheinkman. Common factors affecting bond returns. *Journal of Fixed Income*, 1(1):54–61, 1991.
- Francis A Longstaff, Sanjay Mithal, and Eric Neis. Corporate yield spreads: Default risk or liquidity? new evidence from the credit default swap market. *The journal of finance*, 60(5):2213–2253, 2005.
- Sydney C Ludvigson, Sai Ma, and Serena Ng. Uncertainty and business cycles: exogenous impulse or endogenous response? *American Economic Journal: Macroeconomics*, 13(4):369–410, 2021.

- Ian Martin. What is the expected return on the market? *The Quarterly Journal of Economics*, 132(1):367–433, 2017.
- Nicolai Meinshausen and Peter Bühlmann. Stability selection. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 72(4):417–473, 2010.
- Mark Mitchell and Todd Pulvino. Arbitrage crashes and the speed of capital. *Journal of Financial Economics*, 104(3):469–490, 2012.
- Mark Mitchell, Lasse Heje Pedersen, and Todd Pulvino. Slow moving capital. *American Economic Review*, 97(2):215–220, 2007.
- Alan Moreira and Tyler Muir. Volatility-managed portfolios. *The Journal of Finance*, 72(4):1611–1644, 2017.
- Whitney K Newey and Kenneth D West. A simple positive semi-definite heteroscedasticity and autocorrelation consistent matrix estimation. *Econometrica*, 55:703–708, 1987.
- Whitney K Newey and Kenneth D West. Automatic lag selection in covariance matrix estimation. *The review of economic studies*, 61(4):631–653, 1994.
- Dominic O’Kane. *Modelling single-name and multi-name credit derivatives*. John Wiley & Sons, 2008.
- L’uboš Pástor and Robert F Stambaugh. Liquidity risk and expected stock returns. *Journal of Political economy*, 111(3):642–685, 2003.
- Loriana Pelizzon, Marti G. Subrahmanyam, Davide Tomio, and Jun Uno. Sovereign credit risk, liquidity, and European Central Bank intervention: Deus ex machina? *Journal of Financial Economics*, 122(1):86–115, 2016. doi: 10.1016/j.jfineco.2016.06.001.
- H Hashem Pesaran and Yongcheol Shin. Generalized impulse response analysis in linear multivariate models. *Economics letters*, 58(1):17–29, 1998.
- Dimitris N Politis and Joseph P Romano. The stationary bootstrap. *Journal of the American Statistical association*, 89(428):1303–1313, 1994.

Riccardo Rebonato and Riccardo Ronzani. Is convexity efficiently priced? evidence from international swap markets. *Journal of Empirical Finance*, 63:392–413, 2021.

Emil N Siriwardane. Limited investment capital and credit spreads. *Journal of Finance*, 74(5):2303–2347, 2019.

Dimitri Vayanos and Jean-Luc Vila. A preferred-habitat model of the term structure of interest rates. *Econometrica*, 89(1):77–112, 2021.

Ivo Welch and Amit Goyal. A comprehensive look at the empirical performance of equity premium prediction. *The Review of Financial Studies*, 21(4):1455–1508, 2008.

Hui Zou and Hao Helen Zhang. On the adaptive elastic-net with a diverging number of parameters. *Annals of statistics*, 37(4):1733, 2009.

A Appendix

A.1 BTP Italia: Indexation Coefficient and Deflation Floors

The BTP Italia indexes both coupons and principal to the Italian consumer price index for blue- and white-collar worker households, excluding tobacco (FOI ex tabacchi) published by Istat. At each coupon date t , the indexation coefficient CI_t is the ratio of the relevant FOI reference index to its value at the previous coupon date $t - 6m$, so that CI_t measures the cumulative inflation over the six-month coupon period. The semiannual coupon and the capital revaluation payment are computed as

$$C_t = \frac{1}{2} c N \tilde{CI}_t, \quad R_t = N (\tilde{CI}_t - 1), \quad (19)$$

where c is the fixed real annual coupon rate set at issuance, N is the subscribed nominal capital, and $\tilde{CI}_t$ is the *modified* indexation coefficient that incorporates a semiannual deflation floor:

$$\tilde{CI}_t = \max(CI_t, 1). \quad (20)$$

If the FOI ex tabacchi falls between consecutive coupon dates, the indexation coefficient is reset to one for that semester, the coupon reverts to the fixed real rate $c/2 \cdot N$, and no capital revaluation is paid; in the following semester, the base index is reset to the higher of the previous semester's level and the highest reading observed in earlier semesters. The principal is redeemed at par at maturity (N), regardless of the cumulative inflation path. Two features of this design are relevant for the arbitrage. First, the semiannual reset of $\tilde{CI}_t$ to unity at every coupon date means that the notional exposed to default risk is always the original face value N , eliminating the path-dependent contamination that affects conventional inflation-linked bonds such as the U.S. TIPS or the euro-area BTP€i. Second, the deflation floor on the coupon guarantees a strictly positive nominal payment in every period, which is absent from instruments with floor only at maturity.

A.2 Trade Construction Details

A.2.1 Entry/Exit Rules and Thresholds

Table 1 summarises entry and exit thresholds for each strategy. All strategies require a minimum of 6 months to maturity at entry and maintain at least one open trade after the first signal. The CDS–Bond universe is restricted to single-name CDS whose reference entity belongs to the current on-the-run iTraxx Main series; names that exit the index are no longer traded. All CDS are centrally cleared. The iTraxx skew strategy trades only on-the-run series across Main, SnrFin, SubFin, and Xover.

Table 1
Entry and Exit Thresholds by Strategy

Strategy	Entry Long	Entry Short	Exit Long	Exit Short
BTP Italia	> 40	< –50	< 10	> –10
CDS–Bond Basis	< –40	–	> 0	–
iTraxx Main	> 5	< –10	< –5	> 10
iTraxx SnrFin	> 8	< –15	< –5	> 5
iTraxx SubFin	> 10	< –20	< –10	> 10
iTraxx Xover	> 15	< –25	< –15	> 10

All thresholds in bps. CDS–Bond Basis is a one-directional negative basis trade (long bond, buy CDS protection). Minimum 6 months to maturity at entry for all strategies. Trades shorter than 3 trading days excluded ex-post.

A.2.2 Trade-Level Statistics

Table 2 reports trade-level statistics for each strategy, including the number of trades, average duration, average profit per trade, and the distribution of exit reasons.

A.2.3 Transaction Costs

Transaction costs are computed at the individual trade level and applied as entry and exit fees proportional to DV01 (bonds) or duration (CDS). Fee levels are regime-dependent: at each trade date, the prevailing level of the iTraxx Main 5Y index determines which of three tiers applies (Table 3). Thresholds are set at 60 and 100 bps, reflecting the empirical widening of bid-ask spreads in

Table 2
Trade-Level Statistics by Strategy

	BTP-Italia	iTraxx <i>combined</i>	CDS-Bond Basis
Total trades	260	206	931
Long	224	146	931
Short	36	60	0
Avg duration (days)	386	225	240
Median duration (days)	212	97	119
Avg P&L per trade (%)	2.079	1.491	2.463
Median P&L per trade (%)	1.960	1.395	1.927
<i>Exit reasons (%)</i>			
Target hit	92.3	97.1	90.7
Maturity	6.9	0.5	1.0
End of sample	0.8	2.4	8.4

Note: This table reports trade-level statistics for each arbitrage strategy. P&L is the cumulative return per trade in percentage points. Duration is in calendar days. Exit reasons may not sum to 100% due to rounding.

stressed markets. For trades held to maturity, the exit fee is waived. All return series in subsequent sections are net of these costs.

Table 3
Transaction Fee Schedule (bps per unit DV01 or duration)

	LOW (Main < 60)	MID (60–100)	HIGH (Main > 100)
BTP Italia	3.0	5.0	8.0
CDS–Bond Basis	4.0	7.0	10.0
iTraxx Main	0.5	0.75	1.0
iTraxx SnrFin	0.5	0.75	1.0
iTraxx SubFin	1.0	2.0	3.0
iTraxx Xover	3.0	5.0	7.0

Fee applied as: entry DV01 $\times$ fee_bps at entry + exit DV01 $\times$ fee_bps at exit. Exit fee waived for trades held to maturity. Regime determined by iTraxx Main 5Y level at trade date: LOW (<60 bps), MID (60–100 bps), HIGH (>100 bps).

A.3 Signal-Weighted vs Equal-Weighted Comparison

Table 4 compares the signal-weighted (SW, Rebonato and Ronzani, 2021) and equal-weighted (EW, Duarte et al., 2007) indices. SW is adopted as the baseline in the main text; EW results con-

firm robustness to the weighting scheme. The equal-weighted counterparts of the manipulation-proof performance measure (Table 26) and of the volatility-managed portfolios (Table 7) are reported in Tables 5 and 6.

Table 4
Equal-Weighted Strategy Returns

This table reports monthly percentage excess returns for the equal-weighted arbitrage strategies, net of transaction costs; it is the equal-weighted counterpart of Panel B of Table 25. The mean and standard deviation are annualized and the Sharpe ratio is their ratio; kurtosis is excess kurtosis; the t -statistic for the mean is corrected for serial correlation (Newey–West); and %Neg is the proportion of negative monthly excess returns.

Strategy	n	Mean (% p.a.)	t -Stat	Std Dev (% p.a.)	Min (%)	Max (%)	Skew	Kurt	%Neg	AC(1)	Sharpe
BTP-Italia	163	1.34	3.01	1.69	-1.25	1.92	0.82	1.88	39.9	0.06	0.79
CDS-Bond Basis	207	4.36	5.86	2.71	-3.36	3.55	0.61	4.86	25.6	-0.02	1.61
iTraxx <i>combined</i>	243	1.37	5.30	1.05	-1.44	1.11	0.26	3.68	39.9	-0.04	1.30

Table 5
Manipulation-Proof Performance Measure by Risk Aversion (Equal-Weighted)

MPPM is the manipulation-proof performance measure of Goetzmann et al. (2007), computed from monthly returns; ρ is the coefficient of relative risk aversion and the risk-free rate is 1-month Euribor. Returns are net of transaction costs.

Strategy	$\rho = 2$ (% p.a.)	$\rho = 3$ (% p.a.)	$\rho = 4$ (% p.a.)
BTP-Italia	1.24	1.22	1.21
CDS-Bond Basis	4.19	4.16	4.13
iTraxx <i>combined</i>	1.35	1.34	1.34

Table 6
Moreira-Muir Volatility-Managed Portfolio Performance (Equal-Weighted)

Following Moreira and Muir (2017). The volatility-managed strategy scales the position inversely to the previous month's realized variance, $f_{t+1}^\sigma = (c/\sigma_t^2) f_{t+1}$, with c chosen to match buy-and-hold volatility. Alpha is the intercept from regressing volatility-managed on buy-and-hold monthly returns, with Newey–West standard errors. *, **, *** denote significance at the 10%, 5%, and 1% levels. Returns are net of transaction costs.

Strategy	Buy-and-Hold			Vol-Managed			Alpha	$t(\alpha)$
	Ret (% p.a.)	Vol (% p.a.)	Sharpe	Ret (% p.a.)	Vol (% p.a.)	Sharpe		
BTP-Italia	1.27	1.67	0.76	0.21	1.67	0.13	-0.33	-0.73
CDS-Bond Basis	4.27	2.63	1.63	2.83	2.63	1.08	+1.17	1.58
iTraxx <i>combined</i>	1.36	1.05	1.29	0.60	1.05	0.57	-0.01	-0.09

A.4 Volatility-Managed Portfolios

Table 7 reports the volatility-managed portfolio results of Moreira and Muir [2017] for the three strategies.

Table 7
Moreira-Muir Volatility-Managed Portfolio Performance

Following Moreira and Muir (2017). The volatility-managed strategy scales the position inversely to the previous month's realized variance, $f_{t+1}^\sigma = (c/\sigma_t^2) f_{t+1}$, with c chosen to match buy-and-hold volatility. Alpha is the intercept from regressing volatility-managed on buy-and-hold monthly returns, with Newey–West standard errors. *, **, *** denote significance at the 10%, 5%, and 1% levels. Returns are net of transaction costs.

Strategy	Buy-and-Hold			Vol-Managed			Alpha	$t(\alpha)$
	Ret (% p.a.)	Vol (% p.a.)	Sharpe	Ret (% p.a.)	Vol (% p.a.)	Sharpe		
BTP-Italia	1.52	1.60	0.95	0.37	1.60	0.23	-0.22	-0.50
CDS-Bond Basis	5.09	3.00	1.70	3.49	3.00	1.16	+1.68**	2.11
iTraxx <i>combined</i>	1.91	1.53	1.25	0.84	1.53	0.55	-0.16	-0.58

A.5 Candidate Factor Universe

A.5.1 Complete Factor List with Sources

Table 8 reports the full set of candidate risk factors used in the penalised regressions of Section 4, together with their data sources, transformations, and economic category.

Table 8
Candidate Risk Factors

Factor ID	Description	Reference
Panel A: Credit Risk Factors		
BTP_BUND	BTP-Bund spread factor (10-year). Duration-matched excess return of Italian over German government bonds, long the Italian and short the German 7-10Y total-return index.	Pelizzon et al. [2016]
BTP_BUND_2Y	BTP-Bund spread factor (2-year). Duration-matched excess return of Italian over German government bonds at the short end, long the Italian and short the German 1-3Y total-return index.	Pelizzon et al. [2016]
CDX_IG	Credit risk factor (U.S. investment grade). Excess return on a fully-collateralised sell-protection position in the 5-year CDX North American Investment Grade index.	Asvanunt and Richardson [2017], Bongaerts et al. [2011]
CREDIT_EU	Bond credit risk factor. Return spread between euro-area BBB and AAA corporate bond total-return indices.	Fama and French [1993]
CREDIT_US	Bond credit risk factor. Return spread between BAA and AAA U.S. corporate bond total-return indices.	Fama and French [1993]
CRED_SPR_EU	Bond credit spread factor. Return spread between euro-area BBB corporate bonds and 10-year German Bunds.	Fung and Hsieh [2001]
CRED_SPR_US	Bond credit spread factor. Return spread between BAA U.S. corporate bonds and 10-year U.S. Treasuries.	Fung and Hsieh [2001]
DEF_EU	Bond default risk factor. Total return spread between long-term European corporate bonds and long-term German Bunds.	Fama and French [1993]
DEF_US	Bond default risk factor. Total return spread between long-term U.S. corporate bonds and long-term U.S. government bonds.	Fama and French [1993]
EMERG_DEBT	Excess return on hard-currency emerging market debt.	Brooks et al. [2020]
EMERG_FX	Return on a basket of emerging market currencies versus the USD (MSCI Emerging Markets Currency Index; currency weights from the MSCI EM Index), in excess of USD cash.	Brooks et al. [2020]
GLOBAL_AGG	Excess return on the Bloomberg Global Aggregate Bond Index.	Brooks et al. [2020]
HY_CORP	Sub-investment-grade corporate credit factor: excess return of a EUR-hedged Pan-European high-yield corporate bond index over cash.	Asvanunt and Richardson [2017]

Note: Factors are organized by category. Each series is reconstructed consistently with the original reference. Total: 75 factors.

Table 8 (continued)

Factor ID	Description	Reference
ITRX_MAIN	Credit risk factor (European investment grade). Excess return on a fully-collateralised sell-protection position in the 5-year iTraxx Europe Main index.	Asvanunt and Richardson [2017], Bongaerts et al. [2011]
PB_CDS_1Y_EU	Intermediary funding factor (Europe, short tenor). Excess return on an equally weighted basket of 1-year sell-protection positions in the senior CDS of five European prime brokers.	Siriwardane [2019], He et al. [2017]
PB_CDS_1Y_US	Intermediary funding factor (United States, short tenor). Excess return on an equally weighted basket of 1-year sell-protection positions in the senior CDS of four U.S. prime brokers.	Siriwardane [2019], He et al. [2017]
PB_CDS_5Y_EU	Intermediary funding factor (Europe). Excess return on an equally weighted basket of 5-year sell-protection positions in the senior CDS of five European prime brokers.	Siriwardane [2019], He et al. [2017]
PB_CDS_5Y_US	Intermediary funding factor (United States). Excess return on an equally weighted basket of 5-year sell-protection positions in the senior CDS of four U.S. prime brokers.	Siriwardane [2019], He et al. [2017]
RI_EU	Excess return on a European A-rated industrial corporate bond index.	Duarte et al. [2007]
SLOPE_3S5S_MAIN	Credit-curve slope factor. Excess return on a DV01-matched position long protection at the 5-year and short protection at the 3-year point of the iTraxx Europe Main curve.	Han et al. [2017], Han and Zhou [2015]
SNRFIN_MAIN	Financial-sector stress factor. Excess return on a DV01-matched position long credit via the iTraxx Senior Financials index and short the iTraxx Europe Main index.	He et al. [2017], Siriwardane [2019]
XOVER_MAIN	Credit-cycle factor. Excess return on a beta-neutral position long credit via the 5-year iTraxx Crossover index and short the 5-year iTraxx Europe Main index, the short leg scaled by a rolling 36-month spread-change beta.	Asvanunt and Richardson [2017]
Panel B: Liquidity Factors		
HKM_IC	Intermediary capital factor (tradable). Value-weighted equity return on the primary-dealer holding companies, in excess of cash and converted to euro; a tradable proxy (0.92 correlation) for the He-Kelly-Manela capital-ratio innovation.	He et al. [2017]
LIQ_V	Value-weighted return on a 10–1 portfolio formed by sorting stocks on historical liquidity betas.	Pástor and Staambaugh [2003]
Panel C: Equity Factors		
BAB_EU	Betting Against Beta (Europe). Long low-beta and short high-beta European equities, each leg rescaled to unit beta; USD factor converted to euro.	Frazzini and Pedersen [2014]
CMA_EU	Investment factor (Conservative Minus Aggressive). Return spread between European conservative and aggressive investment portfolios.	Fama and French [2017]
COM_CARRY	Commodity carry factor. Self-financing long-short return on commodity futures sorted on the slope of the futures curve (carry); converted to euro.	Ilmanen et al. [2021]

Note: Factors are organized by category. Each series is reconstructed consistently with the original reference. Total: 75 factors.

Table 8 (continued)

Factor ID	Description	Reference
COM_MOM	Commodity momentum factor. Self-financing long-short return on commodity futures sorted on their past 12-month return; converted to euro.	Ilmanen et al. [2021]
FX_CARRY	Currency carry factor. Self-financing long-short return, long high-interest-rate and short low-interest-rate developed-market (G10) currencies; converted to euro.	Ilmanen et al. [2021]
FX_MOM	Currency momentum factor. Self-financing long-short return on developed-market (G10) currencies sorted on their past 12-month return; converted to euro.	Ilmanen et al. [2021]
GMOM	Global all-asset momentum factor. Long-short momentum return averaged across the eight Value-and-Momentum-Everywhere strategies (stock selection in the U.S., U.K., Europe and Japan, plus equity-index, government-bond, currency and commodity allocations), converted to euro.	Asness et al. [2013]
GVAL	Global all-asset value factor. Long-short value return averaged across the eight Value-and-Momentum-Everywhere strategies (stock selection in the U.S., U.K., Europe and Japan, plus equity-index, government-bond, currency and commodity allocations), converted to euro.	Asness et al. [2013]
HML_EU	Value factor (High Minus Low). Return spread between European high and low book-to-market portfolios.	Fama and French [2017]
MKT_EU	Market excess return on the European value-weighted equity market portfolio over the 1-month German government bill.	Fama and French [2017]
QMJ_EU	Quality Minus Junk (Europe). Long high-quality and short low-quality (“junk”) European stocks, where quality aggregates profitability, growth, safety, and payout; USD factor converted to euro.	Asness et al. [2019]
RB_EU	Excess return on a European bank bond index.	Duarte et al. [2007]
RMW_EU	Profitability factor (Robust Minus Weak). Return spread between European robust and weak operating profitability portfolios.	Fama and French [2017]
RS_EU	Excess return on the EuroStoxx Banks index.	Duarte et al. [2007]
SMB_EU	Size factor (Small Minus Big). Return spread between diversified European small- and large-cap portfolios.	Fama and French [2017]
UMD_EU	Momentum factor (Up Minus Down). Monthly return on a European zero-investment winners-minus-losers momentum portfolio.	Carhart [1997]
Panel D: Volatility Factors		
ATM_IV_CDX	Excess return on a three-month variance swap on the CDX Investment Grade index spread: realized variance of the on-the-run index spread minus its at-the-money (100% moneyness) implied variance fixed at the start of the window (equivalently, the delta-hedged ATM payer/receiver straddle return).	Carr and Wu [2009]
ATM_IV_ITRX	Excess return on a three-month variance swap on the iTraxx Europe Main index spread: realized variance of the on-the-run index spread minus its at-the-money (100% moneyness) implied variance fixed at the start of the window (equivalently, the delta-hedged ATM payer/receiver straddle return).	Carr and Wu [2009]

Note: Factors are organized by category. Each series is reconstructed consistently with the original reference. Total: 75 factors.

Table 8 (continued)

Factor ID	Description	Reference
IV_BUND	Excess return on a one-month variance swap on the 10-year German Bund future: realized variance of the future minus its at-the-money implied variance fixed at the start of the month (equivalently, the delta-hedged ATM straddle return).	Cremers et al. [2021]
IV_TSY	Excess return on a one-month variance swap on the 10-year U.S. Treasury future: realized variance of the future minus its at-the-money implied variance fixed at the start of the month (equivalently, the delta-hedged ATM straddle return).	Cremers et al. [2021]
BdOpt	Return on the bond trend-following factor, constructed from lookback straddles on government bond futures.	Fung and Hsieh [2004]
ComOpt	Return on the commodity trend-following factor, constructed from lookback straddles on commodity futures.	Fung and Hsieh [2004]
FXOpt	Return on the currency trend-following factor, constructed from lookback straddles on major currency futures.	Fung and Hsieh [2004]
PTFSIR	Return on the short-term interest rate trend-following factor, constructed from lookback straddles on 3-month interest rate futures.	Fung and Hsieh [2001]
PTFSSTK	Return on the equity index trend-following factor, constructed from lookback straddles on stock index futures.	Fung and Hsieh [2004]
SVS_RET_1M	Excess return on a one-month simple variance swap on the S&P 500: realized simple variance minus the option-implied SVIX ² strike fixed at the start of the month.	Martin [2017]
SVS_RET_3M	Excess return on a simple variance swap on the S&P 500 over a trailing three-month window (realized simple variance minus the SVIX ² strike).	Martin [2017]
V2X_FUT	Volatility factor (euro area). Excess return on a constant-maturity (one-month) rolling long position in VSTOXX futures (VSTOXX Short-Term Futures index), in excess of the euro risk-free.	Eraker and Wu [2017]
VIX_FUT	Volatility factor (U.S.). Excess return on a constant-maturity (one-month) rolling long position in VIX futures (S&P 500 VIX Short-Term Futures Excess-Return index).	Eraker and Wu [2017]

Panel F: Interest Rate Factors

5Y5Y_INFL	Forward inflation factor (euro area). Mark-to-market excess return on a long-inflation position in a 5y–5y forward zero-coupon inflation swap (annuity-weighted change in the forward breakeven).	Haubrich et al. [2012]
5Y5Y_INFL_US	Forward inflation factor (United States), expressed in euro. Mark-to-market excess return on a long-inflation position in a 5y–5y forward zero-coupon inflation swap.	Haubrich et al. [2012]
CURV_2S5S10S_EU	Yield-curve curvature factor. Excess return on a DV01-neutral butterfly in German government bond futures, long the 5-year belly (Bobl) and short the 2-year (Schatz) and 10-year (Bund) wings.	Litterman and Scheinkman [1991]
CURV_2S5S10S_US	Yield-curve curvature factor. Excess return on a DV01-neutral butterfly in U.S. Treasury futures, long the 5-year belly and short the 2- and 10-year wings.	Litterman and Scheinkman [1991]

Note: Factors are organized by category. Each series is reconstructed consistently with the original reference. Total: 75 factors.

Table 8 (continued)

Factor ID	Description	Reference
FI_CARRY	Fixed-income carry factor. Self-financing long-short return on developed-market government bonds sorted on term-structure carry (yield plus roll-down); converted to euro.	Ilmanen et al. [2021]
FI_DEF	Fixed-income defensive factor. Self-financing long-short return on developed-market government bonds tilted toward low-risk exposures; converted to euro.	Ilmanen et al. [2021]
FI_MOM	Fixed-income momentum factor. Self-financing long-short return on developed-market government bonds sorted on their past 12-month return; converted to euro.	Ilmanen et al. [2021]
FI_VAL	Fixed-income value factor. Self-financing long-short return on a cross-section of developed-market government bonds sorted on a value signal; converted to euro.	Ilmanen et al. [2021]
GLOBAL_TERM	Excess return over cash on the Bloomberg Global Treasury Index (EUR-hedged).	Brooks et al. [2020]
GOVT_EU	Euro-area term-premium factor. Excess return over cash on the broad German government bond aggregate (all maturities), the euro-area analogue of the US Term factor.	Brooks et al. [2020]
INFL_LINK	Excess return on global inflation-linked bonds over cash or nominal Treasuries.	Brooks et al. [2020]
R10_EU	Excess return on a European 10-year government bond portfolio.	Duarte et al. [2007]
R2_EU	Excess return on a European 2-year government bond portfolio.	Duarte et al. [2007]
R5_EU	Excess return on a European 5-year government bond portfolio.	Duarte et al. [2007]
SLOPE_10S30S_EU	Yield-curve slope factor. Excess return on a DV01-matched steepener in German government bond futures, long the 10-year (Bund) and short the 30-year (Buxl).	Litterman and Scheinkman [1991]
SLOPE_2S10S_EU	Yield-curve slope factor. Excess return on a DV01-matched steepener in German government bond futures, long the 2-year (Schatz) and short the 10-year (Bund).	Litterman and Scheinkman [1991]
SLOPE_2S10S_US	Yield-curve slope factor. Excess return on a DV01-matched steepener in U.S. Treasury futures, long the 2-year and short the 10-year.	Litterman and Scheinkman [1991]
SS10Y	Swap-spread factor at the 10-year point. Excess return on a DV01-matched position long the German Bund future and pay-fixed on a 10-year constant-maturity EUR swap.	Koijen et al. [2018], Klingler and Sundaresan [2019]
SS2Y	Swap-spread factor at the 2-year point. Excess return on a DV01-matched position long the German Schatz future and pay-fixed on a 2-year constant-maturity EUR swap.	Koijen et al. [2018], Klingler and Sundaresan [2019]
SS5Y	Swap-spread factor at the 5-year point. Excess return on a DV01-matched position long the German Bobl future and pay-fixed on a 5-year constant-maturity EUR swap.	Koijen et al. [2018], Klingler and Sundaresan [2019]
TERM_EU	Bond term structure factor. Excess total return on long-term German Bunds over the 1-month German government bill return.	Fama and French [1993]
TERM_US	Bond term structure factor. Excess total return on long-term U.S. government bonds over the 1-month U.S. Treasury bill return.	Fama and French [1993]

A.6 PCA Robustness

A.6.1 Robustness across the Number of Components

Table 9 reports alpha and $\bar{R}^2$ across the number of retained components K . Results are robust to the choice of K

Table 9
PCA Spanning Robustness across the Number of Components

K	Btp Italia		Cds Bond Basis		Itraxx Combined		Expl. Var.
	α (% p.a.)	$\bar{R}^2$	α (% p.a.)	$\bar{R}^2$	α (% p.a.)	$\bar{R}^2$	(%)
5	1.50*** (3.21)	0.089	5.37*** (6.00)	0.107	2.21*** (5.20)	0.050	55.0
8	1.44*** (3.24)	0.088	5.36*** (6.01)	0.102	2.21*** (5.68)	0.134	64.7
11	1.47*** (3.50)	0.086	5.35*** (5.87)	0.102	2.21*** (5.70)	0.136	71.8
13	1.75*** (3.97)	0.144	5.35*** (6.00)	0.098	2.21*** (5.82)	0.146	75.6

Full-sample contemporaneous spanning $r_{i,t} = \alpha_i + \beta_i' \mathbf{PC}_t + \varepsilon_{i,t}$ for K principal components (baseline $K = 8$). Newey–West HAC t in parentheses. Expl. Var. is the cumulative variance captured by the K components. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

A.6.2 Subperiod and Regime Analysis

Half-sample splits and regime-dependent alpha confirm that PCA-based alpha is stable across subperiods and increases during stress.

Table 10
PCA Subperiod Analysis (Contemporaneous)

	Full Sample	First Half	Second Half
BTP Italia			
α (% p.a.)	1.44*** (3.24)	1.27** (1.97)	1.38** (2.37)
β_{PC1}	-0.031** (-2.13)	-0.009 (-0.46)	-0.017 (-0.75)
β_{PC2}	-0.025**	-0.003	-0.040**

continued on next page

(continued)

	Full Sample	First Half	Second Half
	(-2.04)	(-0.18)	(-2.54)
β_{PC3}	0.030	-0.003	0.099*
	(1.06)	(-0.08)	(1.79)
β_{PC4}	0.049	0.114	0.039
	(1.42)	(1.64)	(1.09)
β_{PC5}	0.002	0.045	-0.019
	(0.03)	(0.62)	(-0.32)
β_{PC6}	-0.025	-0.033	0.009
	(-0.61)	(-0.77)	(0.18)
β_{PC7}	-0.038	-0.051	0.021
	(-1.45)	(-1.12)	(0.75)
β_{PC8}	-0.022	-0.023	-0.031
	(-1.23)	(-0.42)	(-1.21)
$\bar{R}^2$	0.088	0.009	0.147
N	157	78	79

CDS-Bond

α (% p.a.)	5.36***	7.04***	3.35***
	(6.01)	(4.17)	(4.53)
β_{PC1}	0.035*	0.003	0.122***
	(1.75)	(0.18)	(4.85)
β_{PC2}	0.064***	0.081*	0.034*
	(2.94)	(1.73)	(1.84)
β_{PC3}	-0.085***	-0.056*	-0.107*
	(-2.65)	(-1.83)	(-1.72)
β_{PC4}	-0.016	0.013	-0.002
	(-0.66)	(0.24)	(-0.06)
β_{PC5}	-0.014	-0.007	-0.049
	(-0.43)	(-0.18)	(-0.86)
β_{PC6}	-0.001	-0.052	-0.045
	(-0.03)	(-0.75)	(-1.63)
β_{PC7}	0.045	0.064	0.127**
	(1.34)	(1.40)	(2.54)
β_{PC8}	0.032	0.054	0.089
	(0.80)	(1.20)	(1.55)
$\bar{R}^2$	0.102	0.033	0.268
N	197	98	99

Index Skew

α (% p.a.)	2.21***	2.47***	1.75***
	(5.68)	(4.61)	(3.38)
β_{PC1}	-0.006	0.001	0.013
	(-0.86)	(0.17)	(0.79)
β_{PC2}	0.012	0.001	0.000
	(0.93)	(0.05)	(0.01)
β_{PC3}	0.037***	0.025***	0.120**
	(3.01)	(3.04)	(2.47)
β_{PC4}	0.018	-0.011	0.074***
	(0.94)	(-0.62)	(2.73)

continued on next page

(continued)

	Full Sample	First Half	Second Half
β_{PC5}	0.042*** (2.65)	0.021 (1.57)	0.145*** (2.98)
β_{PC6}	-0.075*** (-3.02)	-0.038* (-1.67)	-0.043* (-1.66)
β_{PC7}	0.020 (1.11)	0.002 (0.06)	0.019 (0.79)
β_{PC8}	0.042* (1.68)	-0.007 (-0.27)	0.097*** (2.79)
$\bar{R}^2$	0.134	0.005	0.253
N	209	104	105


Note: First/Second Half: temporal split at sample midpoint; stress regimes (LOW/MEDIUM/HIGH) are reported in Table 11. α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 11
PCA Alpha by Regime

Strategy	Full Sample	LOW	MEDIUM	HIGH
<i>BTP Italia</i>	1.44*** (3.24)	0.60 (1.10)	1.89*** (3.80)	3.25** (2.35)
<i>CDS–Bond</i>	5.36*** (6.01)	2.58*** (2.67)	4.45*** (3.64)	8.47*** (4.84)
<i>Index Skew</i>	2.21*** (5.68)	1.20 (1.57)	2.20*** (4.88)	3.05*** (3.54)

Note: Regimes on iTraxx Main 5Y: LOW < 60, MEDIUM $\in [60,100)$, HIGH ≥ 100 bps. α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

**Rolling 36-Month Alpha
(Shaded = iTraxx Main > 100 bps)**



A.6.3 Conditional Alpha: Ferson–Schadt and Threshold Robustness (PCA)

Table 12 reports Ferson–Schadt conditional alpha using PCA factors. Table 13 confirms robustness across stress thresholds (80, 100, 120 bps).

Table 12

Conditional Alpha (PCA, Contemporaneous): Christopherson, Ferson, and Glassman (1998) Model

	BTP Italia	CDS–Bond Basis	iTraxx Combined
α_0 (ann. %)	+1.32***	+5.39***	+2.18***
α_1 (conditional)	+0.1111***	+0.2386***	+0.1162***
t -stat	(3.17)	(3.99)	(3.97)
p (bootstrap)	0.0021	0.0032	0.0015
p (HAC, robustness)	0.0015	0.0001	0.0001
$\alpha(z = +1\sigma)$ ann.	+2.73%	+7.92%	+3.67%
$\alpha(z = -1\sigma)$ ann.	+0.06%	+2.20%	+0.88%
T	156	196	208

Note: Christopherson, Ferson, and Glassman (1998) full conditional model with rolling PCA factors (alpha and loadings both conditioned on z_t). Timing: Contemporaneous ($PC_t \rightarrow R_{t+1}$ if predictive). z_t = standardized iTraxx Main 5Y. Significance of α_1 from a moving-block bootstrap ($B = 9,999$, $H_0: \alpha_1 = 0$, block 9); Newey–West HAC t -statistics in parentheses; the HAC p -value is a robustness check. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 13

Conditional Alpha (PCA, Contemporaneous): Stress-Regime Dummies (common β)

Regime	BTP Italia		CDS–Bond Basis		iTraxx Combined	
	α (ann.)	t	α (ann.)	t	α (ann.)	t
LOW	+0.37%	0.65	+2.93***%	3.19	+1.03*%	1.95
MEDIUM	+1.68***%	3.17	+4.87***%	4.04	+2.23***%	4.62
HIGH	+3.27*%	1.82	+9.04***%	5.08	+3.45***%	4.32

Note: Common- β regime dummies (no intercept): $r_t = \sum_g a_g D_{g,t} + \sum_j \beta_j PC_{jt} + \varepsilon_t$, regimes on iTraxx Main 5Y (LOW < 60, MEDIUM $\in$ [60,100), HIGH $\geq$ 100 bps). a_g is the regime- g alpha (% p.a.). t -statistics (Newey–West HAC). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

A.6.4 Variance Diagnostics

Figure 2 shows the time-varying cumulative variance explained.




Figure 2
Time Series of Cumulative Variance Explained by Rolling PCA.

A.7 AEN Robustness

A.7.1 Sensitivity to Hyperparameters

We verify that factor selection is stable to the correlation pre-cleaning threshold ($\rho \in \{0.85, 0.90, 0.95\}$) and the adaptive weight exponent ($\gamma \in \{1, 2, 3\}$). The AEN selects the same set of factors across all configurations, with the post-selection alpha essentially unchanged. No selected factor belongs to a pair correlated above the baseline threshold, and the Stage 1 elastic net coefficients exhibit sufficient separation between relevant and irrelevant factors to make the adaptive weighting insensitive to the exponent.

A.7.2 Factor Selection Across Methods

Table 14 reports the factors selected by each penalised method for each strategy. LASSO, Elastic Net, and Adaptive LASSO produce broadly overlapping selections, while AEN—which retains only factors exceeding the 80% bootstrap stability threshold—yields the sparsest models. ALASSO-Chen selects substantially more factors across all strategies, reflecting its use of AIC for tuning λ_1 [Chen et al., 2025], a less stringent criterion than the HQC adopted for the baseline AEN.

Table 14
Selected Factors by Method and Strategy

Method	BTP Italia	CDS–Bond Basis	iTraxx Combined
LASSO	SMB_EU, UMD_EU, ILLIQ, LIBOR_OIS, CDX_IG, SS10Y, INFL_LINK	CRED_SPR_US, ΔUF, RI_EU, PB_EU_CDS_1Y, EMERG_FX, EP_SVIX_1M, SS5Y, PTFSFX, LIBOR_OIS, SMB_EU, ΔVIX	ΔUR, EP_SVIX_1M, ΔV2X, LIBOR_OIS, YSP_EU, PB_US_CDS_5Y, CMA_EU, GMOM_EU, SS10Y, ΔFAILS_PCT_TSY, HML_EU, MOVE, LIBOR_REPO_SHOCK
Elastic Net	SMB_EU, UMD_EU, ILLIQ, LIBOR_OIS, CDX_IG, SS10Y, INFL_LINK	CRED_SPR_US, ΔUF, RI_EU, PB_EU_CDS_1Y, EMERG_FX, EP_SVIX_1M, SS5Y, PTFSFX, LIBOR_OIS, SMB_EU, ΔVIX	ΔUR, EP_SVIX_1M, ΔV2X, LIBOR_OIS, YSP_EU, PB_US_CDS_5Y, CMA_EU, GMOM_EU, SS10Y, ΔFAILS_PCT_TSY, HML_EU, MOVE, LIBOR_REPO_SHOCK
Ada. LASSO	SMB_EU, UMD_EU, ILLIQ, LIBOR_OIS, CDX_IG, SS10Y, INFL_LINK	CRED_SPR_US, ΔUF, RI_EU, PB_EU_CDS_1Y, EMERG_FX, EP_SVIX_1M, SS5Y, PTFSFX	ΔUR, EP_SVIX_1M, ΔV2X, LIBOR_OIS, YSP_EU, PB_US_CDS_5Y, CMA_EU, GMOM_EU, SS10Y, ΔFAILS_PCT_TSY, HML_EU

(continued)

Method	BTP Italia	CDS–Bond Basis	iTraxx Combined
AEN	SMB_EU, UMD_EU, ILLIQ, LIBOR_OIS, CDX_IG, SS10Y	CRED_SPR_US, ΔUF, RI_EU, PB_EU_CDS_1Y, EMERG_FX, EP_SVIX_1M, SS5Y, PTFSFX	ΔUR, EP_SVIX_1M, ΔV2X, LIBOR_OIS, ΔFAILS_PCT_TSY
ALASSO-Chen	SMB_EU, UMD_EU, ILLIQ, LIBOR_OIS, TERM_EU, TERM_US, ATM_IV_CDX, ΔSLOPE_EU, Δ10Y_YIELD_US, LIBOR_REPO_SHOCK, ΔSLOPE_US, TED_SHOCK_US, ΔFAILS_PCT_TSY, DEF_US, CREDIT_EU, EPU_EU, ΔUF, YSP_US, EPU_US, MKT_EU, BAB_US, BAB_EU, PB_EU_CDS_1Y, PB_US_CDS_1Y, PB_EU_CDS_5Y, HPW_NOISE, GC-REPO_T-BILL, GMOM_EU, SILLIQ, SNR-MAIN, MAIN_5/3_FLATTENER, ITRX_MAIN, XOVER/MAIN, RMW_EU, HML_EU, SS5Y, BFCI_EU, IV_BUND, SS2Y, EURUSD_3M_IV, PTFSBD, PTFSKOM, 5Y5Y_INFL, PTFSSTK, R5_EU, RI_EU, R10_EU, GLOBAL_AGG, CORP_CREDIT, EMERG_DEBT, EMERG_FX	CRED_SPR_US, ΔUF, RI_EU, PB_EU_CDS_1Y, LIBOR_OIS, SMB_EU, CREDIT_US, DEF_US, EP_SVIX_3M, TERM_US, ΔV2X, CRED_SPR_EU, EPU_US, MKT_EU, BAB_US, TED_SHOCK_EU, PB_US_CDS_1Y, LIBOR_REPO_SHOCK, HML_EU, HPW_NOISE, CMA_EU, SS10Y, XOVER/MAIN, MOVE, EBP, PTFSBD, EURUSD_3M_IV, EURIBOR_OIS, GLOBAL_AGG	ΔUR, EP_SVIX_1M, ΔV2X, LIBOR_OIS, YSP_EU, PB_US_CDS_5Y, CMA_EU, GMOM_EU, SS10Y, HKM_IC, EP_SVIX_3M, TERM_US, DEF_EU, ΔSLOPE_EU, TED_SHOCK_EU, PB_EU_CDS_5Y, CDX_IG, SS5Y, EMERG_DEBT

Note: Each cell lists the factors selected (non-zero coefficient) by the corresponding method. Ridge excluded (selects all). AEN denotes the Adaptive Elastic Net with bootstrap stability selection ($\hat{\pi}_j \geq 80\%$).

A.7.3 Conditional Alpha: Detailed Results (AEN)

Table 15 reports the full Ferson–Schadt conditional alpha model with AEN-selected factors. Table ?? confirms robustness across stress thresholds (80, 100, 120 bps).

Table 15
Conditional Alpha: Christopherson, Ferson, and Glassman (1998) Model

	BTP Italia	CDS–Bond Basis	iTraxx Combined
<i>Panel A: Conditional alpha</i>			
α_0 (ann. %)	+1.83***	+4.76***	+2.69***
α_1 (conditional)	+0.0395	+0.1475*	+0.0477
t -stat	(0.99)	(2.58)	(1.38)
p (bootstrap)	0.2326	0.0917	0.2379
p (HAC, robustness)	0.3199	0.0098	0.1674
$\bar{R}^2$ (unconditional)	0.2442	0.2306	0.1759
$\bar{R}^2$ (conditional)	0.2926	0.2917	0.2010
$\alpha(z = +1\sigma)$ ann.	+2.42%	+6.96%	+3.29%
$\alpha(z = -1\sigma)$ ann.	+1.47%	+3.42%	+2.14%
T	156	196	210
<i>Panel B: Conditional betas (δ_j)</i>			
5Y5Y _{INFL_US} $\times z_t$	−0.043	—	—
ATM _{IV_CDX} $\times z_t$	—	—	−0.001
ATM _{IV_ITRX} $\times z_t$	—	—	−0.266
BAB _{EU} $\times z_t$	—	—	−0.006
BTP _{BUND_2Y} $\times z_t$	—	+0.134**	—
COM _{CARRY} $\times z_t$	—	−0.004	—
CRED _{SPR_US} $\times z_t$	+0.080	—	—
DEF _{US} $\times z_t$	−0.073	−0.008	—
FI _{DEF} $\times z_t$	−0.066	—	+0.030
FI _{MOM} $\times z_t$	—	—	+0.014
HML _{EU} $\times z_t$	—	—	−0.008
MKT _{EU} $\times z_t$	+0.014	—	—
PB _{CDS_1Y_EU} $\times z_t$	—	+0.306	—
PTFSFX $\times z_t$	—	+0.002	—
QMJ _{EU} $\times z_t$	−0.003	—	—
SS10Y $\times z_t$	+0.022*	—	—
SS5Y $\times z_t$	−0.018	+0.020***	—
SVS _{RET_1M} $\times z_t$	—	−1.145	−0.094
TERM _{US} $\times z_t$	—	−0.009	—
XOVER _{MAIN} $\times z_t$	—	—	+0.020

Note: Christopherson, Ferson, and Glassman (1998) full conditional performance model: both the alpha and the factor loadings are conditioned on z_t , the standardized iTraxx Main 5Y spread (mean zero, unit variance). $\alpha_1 > 0$ indicates alpha increases with credit stress. Significance of α_1 is from a moving-block bootstrap ($B = 9,999$, $H_0: \alpha_1 = 0$ imposed, block length 9); Newey–West HAC t -statistics in parentheses are Newey–West HAC (lag 3). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. Factor definitions are provided in Table 8.

Table 16
Conditional Alpha by Stress Regime

Strategy	LOW (< 60 bps)	MEDIUM ([60, 100) bps)	HIGH (≥ 100 bps)
BTP Italia	+0.96*% (1.90)	+2.08***% (3.44)	+3.33**% (2.26)
CDS–Bond Basis	+2.60**% (2.57)	+4.73***% (3.78)	+8.39***% (5.74)
iTraxx Combined	+1.85***% (3.33)	+2.87***% (4.21)	+3.54***% (4.16)

Note: $r_t = \alpha_{\text{LOW}}D_{\text{LOW},t} + \alpha_{\text{MED}}D_{\text{MED},t} + \alpha_{\text{HIGH}}D_{\text{HIGH},t} + \sum_j \beta_j X_{jt} + \varepsilon_t$ (no intercept; the three regime dummies partition the sample, betas common across regimes). Regimes on iTraxx Main 5Y: LOW < 60 , MEDIUM $\in [60, 100)$, HIGH ≥ 100 bps. Annualized α ; t -statistics in parentheses. Factors from best-subset selection. HAC (Newey–West) SE. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

A.7.4 Rolling Alpha

Figure 3 plots rolling 36-month alpha with regime shading for the AEN-selected factors. Alpha spikes during stress episodes.



Figure 3

Rolling Alpha with Regime Shading (AEN Stable Factors). Grey bands: 95% confidence intervals. Red shading: iTraxx Main 5Y > 100 bps. Rolling window: 36 months.

A.8 Subperiod and Rolling Window Analysis

This appendix reports sub-period splits and threshold robustness for the benchmark factor models of Section 3. Each table contains Panel A (full sample, first half, second half, HIGH, and NORMAL regimes) and Panel B (stress threshold robustness at 80, 100, and 120 bps of iTraxx Main 5Y).

A.8.1 Duarte et al. (2007)

Table 17 reports sub-period splits and stress regime alpha for the Duarte et al. [2007] specification.

Table 17
Sub-Period and Regime Analysis: Duarte et al. (2007)

<i>Panel A: Sub-period alpha (% p.a.)</i>									
Period	BTP Italia			iTraxx Combined			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
Full Sample	+1.62***	3.29	163	+2.19***	4.90	243	+4.80***	5.42	207
First Half	+1.11*	1.96	81	+2.10***	4.56	121	+6.75***	4.73	103
Second Half	+1.79**	2.33	82	+2.60***	3.41	122	+2.63***	3.00	104

<i>Panel B: Alpha by stress regime (% p.a.)</i>									
Regime	BTP Italia			iTraxx Combined			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
LOW	+0.78	1.24	62	+0.97**	2.22	91	+2.22**	2.30	63
MEDIUM	+1.97***	3.15	80	+2.69***	3.65	99	+4.41***	3.69	94
HIGH	+3.05*	1.70	21	+3.27***	4.89	53	+8.57***	4.80	50

Notes: Factor model: Duarte et al. (2007). EUR factors. Monthly frequency. Panel A: OLS with HAC (Newey–West) standard errors on time sub-samples. Panel B: regime alphas from common-beta dummy intercepts $r_t = \sum_g \alpha_g D_{g,t} + \sum_j \beta_j X_{jt} + \varepsilon_t$, with $g \in \{\text{LOW, MEDIUM, HIGH}\}$ defined on the iTraxx Main 5Y level (LOW < 60 bps, MEDIUM [60, 100), HIGH $\geq$ 100). N = months per regime; HAC (Newey–West) t -statistics. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.




Figure 4

Rolling 36-Month Alpha with Regime Shading — Duarte et al. (2007). 95% confidence bands. Shading denotes the iTraxx Main 5Y regime: green (LOW, <60 bps), amber (MEDIUM, 60–100 bps), red (HIGH, ≥ 100 bps).

A.8.2 Brooks et al. (2020) — Active Fixed Income

Table 18 reports the same analysis under the Brooks et al. [2020] specification.

Table 18
Sub-Period and Regime Analysis: Active FI (Brooks et al. (2020))

<i>Panel A: Sub-period alpha (% p.a.)</i>									
Period	BTP Italia			iTraxx Combined			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
Full Sample	+1.98***	3.83	163	+1.84***	4.40	243	+4.90***	5.53	207
First Half	+1.26	1.63	81	+1.69***	3.88	121	+7.33***	4.86	103
Second Half	+2.15***	3.11	82	+1.84***	2.92	122	+2.03***	3.43	104

<i>Panel B: Alpha by stress regime (% p.a.)</i>									
Regime	BTP Italia			iTraxx Combined			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
LOW	+1.03**	2.02	62	+0.79**	2.05	91	+2.59***	2.91	63
MEDIUM	+2.31***	3.74	80	+2.15***	3.11	99	+4.49***	3.70	94
HIGH	+3.95*	1.85	21	+3.21***	3.99	53	+9.11***	5.03	50

Notes: Factor model: Active FI (Brooks et al. (2020)). EUR factors. Monthly frequency. Panel A: OLS with HAC (Newey–West) standard errors on time sub-samples. Panel B: regime alphas from common-beta dummy intercepts $r_t = \sum_g \alpha_g D_{g,t} + \sum_j \beta_j X_{jt} + \varepsilon_t$, with $g \in \{\text{LOW, MEDIUM, HIGH}\}$ defined on the iTraxx Main 5Y level (LOW < 60 bps, MEDIUM [60, 100), HIGH $\geq$ 100). N = months per regime; HAC (Newey–West) t -statistics. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.




Figure 5
Rolling 36-Month Alpha with Regime Shading — Brooks et al. (2020). Same format as Figure 4.

A.8.3 Fung & Hsieh (2004)

Table 19 reports results for the Fung and Hsieh [2004] specification.

Table 19
Sub-Period and Regime Analysis: Fung & Hsieh (2004)

<i>Panel A: Sub-period alpha (% p.a.)</i>									
Period	BTP Italia			iTraxx Combined			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
Full Sample	+1.81***	4.08	163	+1.84***	4.49	243	+4.53***	5.24	207
First Half	+1.56**	2.04	81	+1.96***	4.40	121	+6.64***	4.57	103
Second Half	+1.66***	3.00	82	+1.73***	3.12	122	+2.19***	2.81	104

<i>Panel B: Alpha by stress regime (% p.a.)</i>									
Regime	BTP Italia			iTraxx Combined			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
LOW	+0.95	1.63	62	+0.50	1.44	91	+2.34**	2.45	63
MEDIUM	+2.05***	3.81	80	+2.23***	3.16	99	+4.03***	3.35	94
HIGH	+3.33*	1.74	21	+3.34***	4.21	53	+8.16***	4.45	50

Notes: Factor model: Fung & Hsieh (2004). EUR factors. Monthly frequency. Panel A: OLS with HAC (Newey–West) standard errors on time sub-samples. Panel B: regime alphas from common-beta dummy intercepts $r_t = \sum_g \alpha_g D_{g,t} + \sum_j \beta_j X_{jt} + \varepsilon_t$, with $g \in \{\text{LOW, MEDIUM, HIGH}\}$ defined on the iTraxx Main 5Y level (LOW < 60 bps, MEDIUM [60, 100), HIGH $\geq$ 100). N = months per regime; HAC (Newey–West) t -statistics. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.




Figure 6
Rolling 36-Month Alpha with Regime Shading — Fung & Hsieh (2004). Same format as Figure 4.

A.9 Cross-Strategy Interdependencies: Additional Tests

This appendix collects robustness checks and supplementary analyses for Section 5. All tests use the same mispricing magnitudes $m_{i,t}$ and regime definitions as the main text.

A.9.1 Purged Correlations

To verify that cross-strategy co-movement is not driven solely by common risk-factor exposures, we regress each Δm_i on the AEN-selected factors from Section 4 and compute correlations on the residuals. Ta-

ble 20 shows that the pairwise correlations retain their sign after purging—positive within the institutional pair (+0.13) and negative against the retail-held BTP Italia (−0.27 and −0.09)—though those involving BTP Italia attenuate appreciably. The persistence of the sign pattern indicates that the interdependencies documented in Section 5 reflect capital-market linkages beyond shared factor exposures, while the attenuation shows that part of the raw co-movement does operate through common factor exposure.

Table 20
Correlations After Purging Common Factors

Pair	Raw	Purged	Reduction
CDS–Bond Basis vs BTP Italia	−0.298	−0.318	−6.69%
CDS–Bond Basis vs iTraxx Combined	0.191	0.078	59.31%
BTP Italia vs iTraxx Combined	−0.100	−0.058	42.25%

Note: Correlations of residuals after regressing each strategy’s Δm on common factors (EURIBOR-OIS, HKM intermediary capital, HPW noise). Reduction shows the percentage decline in absolute correlation.

A.9.2 DCC-GARCH

We estimate bivariate DCC-GARCH(1,1) models on each pair of Δm series. Figure 7 plots the time-varying conditional correlations. Spikes in conditional correlation coincide with stress episodes (COVID-19, ECB hiking cycle), corroborating the regime-dependent results in Section 5.




Figure 7

Dynamic Conditional Correlations (DCC-GARCH). Bivariate DCC(1,1) estimates for each pair of Δm . Shaded areas: HIGH stress regime (iTraxx Main > 100 bps).

A.9.3 Granger Causality: Lag Robustness

Table 21 reports Granger causality tests at lag lengths $p = 1, 2, 3$. The main result — iTraxx Combined Granger-causes both the CDS-Bond basis and BTP Italia — is robust across all specifications.

Table 21
Granger Causality: Robustness to Lag Length

Cause	Effect	$p = 1$		$p = 2$		$p = 3$	
		F	p -val	F	p -val	F	p -val
BTP Italia	CDS–Bond Basis	0.05	0.8176	0.55	0.5790	1.61	0.1858
iTraxx Combined	CDS–Bond Basis	6.52**	0.0110	3.97**	0.0195	2.14*	0.0946
CDS–Bond Basis	BTP Italia	3.08*	0.0802	1.40	0.2465	1.53	0.2071
iTraxx Combined	BTP Italia	7.39***	0.0068	4.16**	0.0162	3.04**	0.0289
CDS–Bond Basis	iTraxx Combined	0.01	0.9043	0.16	0.8537	0.13	0.9445
BTP Italia	iTraxx Combined	0.48	0.4897	1.22	0.2970	0.87	0.4589

Note: Granger causality F -tests from VAR(p) on Δm , estimated separately at $p = 1, 2, 3$ lags. Lag selection: HQC selects $p = 1$; AIC/FPE select $p = 2$; BIC selects $p = 0$ (forced to 1). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

A.10 VIF Diagnostics

Tables 22–24 report Variance Inflation Factors for the EUR factor specifications used in the benchmark regressions of Section 3. VIF is computed on the factor matrix for each strategy’s sample period; because the factors are identical across strategies, differences in VIF reflect only the different estimation windows.

All equity and bond-market factors exhibit VIF below 10, confirming acceptable multicollinearity levels. The Duarte specification shows elevated VIF for the term-structure factors R2, R5, and R10 (VIF ≈ 15 –55), which is expected given the high correlation structure of yield-curve returns and does not affect inference on alpha. The Fung & Hsieh specification shows VIF below 5 for all factors. The Active FI specification shows VIF below 5 for most factors, with elevated values for Global_Term and Global_Aggregate (VIF ≈ 40 –50) reflecting the strong correlation between global bond benchmarks.

Table 22
Variance Inflation Factors — Duarte et al. (2007), EUR Factors

Factor	BTP Italia	CDS–Bond Basis	iTraxx Combined
Mkt-RF	4.86	5.27	4.97
SMB	1.11	1.18	1.13
HML	3.15	3.14	3.01
UMD	1.72	1.85	1.81
RS	7.29	8.25	8.10
RI	16.07	5.23	5.08
RB	11.83	2.82	2.77
R2	16.02	14.42	15.93
R5	55.30	51.71	55.21
R10	26.56	24.19	24.01

Note: VIF computed on EUR factor matrix. VIF > 10 indicates serious multicollinearity; VIF > 5 moderate concerns. Elevated VIF for R2, R5, R10 reflects the inherent correlation of Treasury return factors and does not affect inference on alpha.

Table 23
Variance Inflation Factors — Brooks et al. (2020), EUR Factors

Factor	BTP Italia	CDS–Bond Basis	iTraxx Combined
Term	6.61	6.03	5.53
Global_Term	60.48	51.95	50.57
Global_Aggregate	66.17	55.82	53.25
Inflation_Linkers	5.04	4.10	3.84
Corporate_Credit	4.01	3.75	3.64
Emerging_Debt	8.22	6.04	5.60
Emerging_Currency	2.12	2.29	2.12
UST_Volatility	1.53	1.47	1.42

Note: VIF computed on EUR factor matrix. VIF > 10 indicates serious multicollinearity; VIF > 5 moderate concerns.

Table 24
Variance Inflation Factors — Fung & Hsieh (2004), EUR Factors

Factor	BTP Italia	CDS–Bond Basis	iTraxx Combined
SNPMRF	2.22	2.49	2.22
SCMLC	1.07	1.15	1.10
PTFSBD	1.45	1.40	1.37
PTFSFX	1.46	1.50	1.46
PTFSCOM	1.19	1.27	1.24
R10	3.28	2.60	2.45
BAAMTSY	4.10	4.36	4.04

Note: VIF computed on EUR factor matrix. VIF > 10 indicates serious multicollinearity; VIF > 5 moderate concerns.

Tables

Table 25
Summary Statistics for Fixed Income Arbitrage Strategies

This table reports summary statistics for the signal-weighted arbitrage strategies, net of transaction costs. Panel A reports monthly basis levels (in basis points), with P_{10} and P_{90} the 10th and 90th percentiles; Panel B reports monthly percentage excess returns. The basis for each strategy is constructed as described in Section 2.1. In Panel B, the mean and standard deviation are annualized and the Sharpe ratio is their ratio; kurtosis is excess kurtosis; the t -statistic for the mean is corrected for serial correlation (Newey–West); and %Neg is the proportion of negative monthly excess returns. All series end in May 2025; start dates differ across strategies.

Panel A: Basis Levels (bps)										
Strategy	Mean	Median	Std Dev	Min	Max	Skew	Kurt	AC(1)	P_{10}	P_{90}
<i>BTP Italia</i>	44.7	46.3	28.6	-66.5	101.0	-1.43	3.29	0.81	14.4	73.8
<i>CDS-Bond Basis</i>	6.9	9.8	18.5	-73.9	39.5	-0.97	1.69	0.83	-18.7	30.5
<i>iTraxx Combined</i>	0.2	1.1	6.7	-37.5	21.9	-1.92	7.76	0.66	-6.1	6.6

Panel B: Strategy Returns (%)											
Strategy	n	Mean (% p.a.)	t -Stat	Std Dev (% p.a.)	Min (%)	Max (%)	Skew	Kurt	%Neg	AC(1)	Sharpe
<i>BTP Italia</i>	163	1.59	3.52	1.62	-1.25	1.81	0.81	1.77	36.8	0.08	0.98
<i>CDS Bond Basis</i>	207	5.16	5.95	3.07	-3.31	3.81	0.91	4.08	24.6	0.00	1.68
<i>iTraxx Combined</i>	243	1.91	4.92	1.53	-2.06	2.50	1.01	7.37	38.3	0.01	1.25

Table 26
Manipulation-Proof Performance Measure by Risk Aversion

MPPM is the manipulation-proof performance measure of Goetzmann et al. (2007), computed from monthly returns; ρ is the coefficient of relative risk aversion and the risk-free rate is 1-month Euribor. Returns are net of transaction costs.

Strategy	$\rho = 2$ (% p.a.)	$\rho = 3$ (% p.a.)	$\rho = 4$ (% p.a.)
BTP-Italia	1.50	1.48	1.47
CDS-Bond Basis	4.99	4.95	4.90
iTraxx <i>combined</i>	1.88	1.87	1.86

Table 27
Duarte Factor Model Regressions

	α (%)	β_{Mkt}	β_{SMB}	β_{HML}	β_{UMD}	β_{R_S}	β_{R_I}	β_{R_B}	β_{R_2}	β_{R_5}	$\beta_{R_{10}}$	$\bar{R}^2$
Panel A: US Factors												
<i>BTP Italia</i>	2.02*** (4.17)	-0.05*** (-3.49)	-0.02* (-1.72)	-0.02 (-1.60)	0.01 (0.44)	0.03*** (3.19)	0.03 (0.37)	-0.09 (-1.19)	-0.02 (-0.07)	-0.23 (-0.62)	0.21 (0.97)	0.11
<i>CDS-Bond</i>	5.14*** (6.01)	-0.03 (-0.90)	0.01 (0.48)	-0.00 (-0.09)	-0.01 (-1.10)	0.00 (0.11)	0.26*** (3.30)	-0.06 (-1.56)	0.12 (0.24)	-0.65* (-1.69)	0.25 (1.25)	0.16
<i>Index Skew</i>	1.86*** (4.71)	0.00 (0.31)	-0.02* (-1.68)	-0.02 (-1.38)	-0.01* (-1.92)	-0.00 (-0.09)	0.02 (0.50)	-0.02 (-1.04)	0.22 (0.74)	-0.07 (-0.31)	0.01 (0.10)	0.02
Panel B: EUR Factors												
<i>BTP Italia</i>	1.62*** (3.27)	-0.04* (-1.84)	-0.06*** (-2.85)	-0.03 (-1.43)	0.02 (1.16)	0.01 (1.03)	-0.03 (-0.27)	0.04 (0.46)	0.03 (0.06)	-0.04 (-0.14)	-0.02 (-0.16)	0.11
<i>CDS-Bond</i>	4.80*** (5.63)	0.00 (0.11)	0.03 (0.79)	0.03 (1.03)	-0.02 (-0.92)	-0.03* (-1.71)	0.47*** (3.95)	-0.06 (-1.54)	-0.29 (-0.45)	0.00 (0.00)	-0.13 (-0.79)	0.14
<i>Index Skew</i>	2.19*** (4.96)	-0.02 (-1.25)	-0.03** (-2.04)	-0.07** (-2.28)	-0.02* (-1.92)	0.01 (1.55)	0.05 (0.73)	-0.02 (-0.75)	0.06 (0.16)	0.05 (0.21)	-0.06 (-0.78)	0.06

Note: Monthly excess returns regressed on the ten Duarte et al. (2007) equity and bond risk factors. α is annualized (multiplied by 12). *t*-statistics in parentheses (Newey–West HAC). *Mkt*: market excess return (Fama–French). *R_S*: bank stock excess return. *R_I*, *R_B*: A/BAA-rated industrial and bank bond excess returns. *R₂*, *R₅*, *R₁₀*: maturity-sorted Treasury portfolio excess returns. Sample: BTP Italia 163 months, CDS–Bond 207 months, Index Skew 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 28
Fung & Hsieh Factor Model Regressions

	α (%)	$\beta_{S\&P}$	β_{SC-LC}	β_{BdOpt}	β_{FXOpt}	β_{ComOpt}	β_{10Y}	$\beta_{CredSpr}$	$\bar{R}^2$
Panel A: US Factors									
<i>BTP Italia</i>	1.84*** (3.82)	-0.02* (-1.94)	-0.02* (-1.89)	0.00 (0.90)	-0.00 (-0.00)	0.00 (0.86)	-0.02 (-0.50)	-0.04 (-1.05)	0.10
<i>CDS-Bond</i>	4.71*** (5.26)	-0.01 (-0.70)	0.01 (0.48)	0.00 (1.13)	-0.01* (-1.95)	0.00 (0.03)	0.13*** (3.44)	0.14*** (3.29)	0.11
<i>Index Skew</i>	1.75*** (4.44)	0.01 (0.65)	-0.02** (-2.02)	0.00 (0.39)	0.00 (1.03)	0.00 (0.52)	0.03 (1.11)	0.01 (0.51)	0.02
Panel B: EUR Factors									
<i>BTP Italia</i>	1.81*** (4.09)	-0.03** (-2.52)	-0.06*** (-2.71)	0.00 (1.22)	0.00 (0.03)	0.00 (1.08)	-0.01 (-0.27)	0.02 (0.31)	0.10
<i>CDS-Bond</i>	4.53*** (5.50)	-0.02 (-1.35)	0.05 (1.13)	0.00 (1.01)	-0.01* (-1.93)	-0.00 (-0.46)	0.18*** (3.39)	0.19*** (2.76)	0.10
<i>Index Skew</i>	1.84*** (4.58)	0.00 (0.40)	-0.02 (-1.45)	0.00 (0.62)	0.00 (1.00)	0.00 (0.58)	0.03 (0.84)	0.02 (0.61)	0.00

Note: Monthly excess returns regressed on the seven Fung & Hsieh (2004) hedge fund risk factors. α is annualized (multiplied by 12). t -statistics in parentheses (Newey–West HAC). S&P: S&P 500 excess return. SC–LC: Wilshire small-minus-large-cap equity spread. BdOpt, FXOpt, ComOpt: lookback-straddle returns on bond, currency, and commodity futures. 10Y: change in 10-year constant-maturity yield. CredSpr: change in Moody’s Baa minus 10-year Treasury spread. Sample: BTP Italia 163 months, CDS–Bond 207 months, Index Skew 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 29
Active Fixed Income Factor Model Regressions

	α (%)	β_{Term}	$\beta_{\text{Global Term}}$	$\beta_{\text{Global Agg}}$	$\beta_{\text{Infl. Linkers}}$	$\beta_{\text{Corp Credit}}$	$\beta_{\text{Emerg Debt}}$	$\beta_{\text{Emerg Curr}}$	$\beta_{\text{UST Vol}}$	$\bar{R}^2$
Panel A: US Factors										
<i>BTP Italia</i>	1.99*** (4.56)	0.25** (2.50)	0.09 (0.28)	-0.35 (-0.90)	-0.06* (-1.85)	-0.04 (-0.79)	0.06 (1.25)	-0.01 (-0.35)	-11.05 (-0.57)	0.11
<i>CDS-Bond</i>	4.51*** (5.60)	-0.19 (-1.14)	-0.07 (-0.11)	0.40 (0.54)	-0.05 (-0.70)	-0.04 (-0.74)	0.14*** (3.36)	-0.01 (-0.33)	25.11 (0.85)	0.11
<i>Index Skew</i>	1.86*** (4.26)	0.07 (0.87)	-0.22 (-0.90)	0.26 (1.01)	-0.02 (-0.48)	0.06** (2.04)	-0.08** (-2.17)	0.01 (0.52)	34.05* (1.90)	0.06
Panel B: EUR Factors										
<i>BTP Italia</i>	1.98*** (3.82)	-0.01 (-0.29)	-0.05 (-0.16)	0.12 (0.45)	-0.06* (-1.88)	-0.11** (-2.12)	0.06 (1.21)	-0.02 (-0.69)	-0.63 (-0.03)	0.07
<i>CDS-Bond</i>	4.90*** (5.75)	0.07 (1.57)	-0.38 (-0.62)	0.26 (0.46)	-0.06 (-0.89)	-0.03 (-0.66)	0.14*** (2.89)	0.02 (0.43)	13.51 (0.48)	0.10
<i>Index Skew</i>	1.84*** (4.48)	-0.04* (-1.87)	-0.23 (-1.00)	0.45* (1.88)	0.01 (0.22)	0.02 (1.45)	-0.08** (-2.26)	0.01 (0.26)	38.28** (2.01)	0.06

Note: Monthly excess returns regressed on the eight Active Fixed Income premia of Brooks et al. (2020). α is annualized (multiplied by 12). t -statistics in parentheses (Newey–West HAC). Term: Treasury excess returns (US or Eurozone). Global Term/Agg: Bloomberg Global Treasury/Aggregate (hedged). Infl. Linkers: Bloomberg Global Inflation-Linked. Corp Credit: high-yield corporate bond return in excess of cash (leveraged-loan leg dropped). Emerg Debt/Curr: EM bond and FX baskets. UST Vol: delta-hedged straddles on 10Y Treasury futures. Sample: BTP Italia 163 months, CDS-Bond 207 months, Index Skew 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 30
Alpha Estimates Across Benchmark Factor Models

Panel A: Full-sample alpha (% p.a.)

Strategy	Duarte et al. (2007)			Brooks et al. (2020)			Fung & Hsieh (2004)		
	α	t	N	α	t	N	α	t	N
<i>BTP Italia</i>	+1.62***	3.29	163	+1.98***	3.83	163	+1.81***	4.08	163
<i>iTraxx Combined</i>	+2.19***	4.90	243	+1.84***	4.40	243	+1.84***	4.49	243
<i>CDS–Bond Basis</i>	+4.80***	5.42	207	+4.90***	5.53	207	+4.53***	5.24	207

Panel B: Alpha by stress regime (% p.a.)

	Duarte et al. (2007)			Brooks et al. (2020)			Fung & Hsieh (2004)		
	LOW	MED	HIGH	LOW	MED	HIGH	LOW	MED	HIGH
<i>BTP Italia</i>	+0.78 (1.24)	+1.97*** (3.15)	+3.05* (1.70)	+1.03** (2.02)	+2.31*** (3.74)	+3.95* (1.85)	+0.95 (1.63)	+2.05*** (3.81)	+3.33* (1.74)
<i>iTraxx Combined</i>	+0.97** (2.22)	+2.69*** (3.65)	+3.27*** (4.89)	+0.79** (2.05)	+2.15*** (3.11)	+3.21*** (3.99)	+0.50 (1.44)	+2.23*** (3.16)	+3.34*** (4.21)
<i>CDS–Bond Basis</i>	+2.22** (2.30)	+4.41*** (3.69)	+8.57*** (4.80)	+2.59*** (2.91)	+4.49*** (3.70)	+9.11*** (5.03)	+2.34** (2.45)	+4.03*** (3.35)	+8.16*** (4.45)

Panel C: Christopherson, Ferson, and Glassman (1998) conditional alpha (% p.a.)

	Duarte et al. (2007)			Brooks et al. (2020)			Fung & Hsieh (2004)		
	α_0	α_1	$\bar{R}^2$	α_0	α_1	$\bar{R}^2$	α_0	α_1	$\bar{R}^2$
<i>BTP Italia</i>	+1.58*** (3.78)	+1.20** (2.25)	0.23	+1.74*** (3.13)	+0.85 (1.02)	0.13	+1.70*** (5.24)	+0.50 (1.22)	0.19
<i>iTraxx Combined</i>	+1.94*** (4.75)	+0.93** (2.55)	0.11	+1.93*** (4.43)	+1.47*** (3.41)	0.15	+1.89*** (4.85)	+1.51*** (4.66)	0.05
<i>CDS–Bond Basis</i>	+5.27*** (6.37)	+2.73*** (3.34)	0.26	+5.35*** (6.85)	+3.75*** (4.27)	0.16	+4.53*** (5.86)	+2.33** (3.01)	0.18

Panel D: Out-of-sample alpha (% p.a.)

	Duarte et al. (2007)			Brooks et al. (2020)			Fung & Hsieh (2004)		
	α	t	N	α	t	N	α	t	N
<i>BTP Italia</i>	+1.30**	2.20	102	+2.11***	3.73	103	+1.65**	2.41	102
<i>iTraxx Combined</i>	+2.30***	4.85	181	+1.97***	4.08	183	+3.01***	4.96	135
<i>CDS–Bond Basis</i>	+4.07***	4.88	145	+3.69***	4.00	147	+3.46***	5.18	135

Note: Panel A reports annualized alpha (% p.a.) from monthly OLS regressions with Newey–West HAC standard errors using EUR factors. Panel B reports regime alphas from common-beta dummy intercepts, with LOW / MEDIUM / HIGH defined on the 5-year iTraxx Europe Main level (LOW < 60, MEDIUM [60, 100), HIGH ≥ 100 bps). Panel C reports the intercept α_0 and the stress-loading coefficient α_1 from the Christopherson et al. [1998] full conditional model $r_t = \alpha_0 + \alpha_1 z_t + (\beta + \delta z_t)' X_t + \varepsilon_t$, in which both the alpha and the factor loadings are conditioned on z_t , the standardized iTraxx Main 5Y spread; $\bar{R}^2$ is the adjusted R^2 . Significance of α_1 is assessed by a moving-block bootstrap ($B = 9,999$, $H_0: \alpha_1 = 0$ imposed, block length 9); Newey–West HAC t -statistics in parentheses are Newey–West HAC. Panel D reports out-of-sample alpha:

Table 31
PCA Spanning Regressions

	α (%)	β_{PC1}	β_{PC2}	β_{PC3}	β_{PC4}	β_{PC5}	β_{PC6}	β_{PC7}	β_{PC8}	$\bar{R}^2$	N
Contemporaneous ($PC_t \rightarrow R_t$)											
<i>BTP Italia</i>	1.44*** (3.24)	-0.031** (-2.13)	-0.025** (-2.04)	0.030 (1.06)	0.049 (1.42)	0.002 (0.03)	-0.025 (-0.61)	-0.038 (-1.45)	-0.022 (-1.23)	0.088	157
<i>CDS-Bond</i>	5.36*** (6.01)	0.035* (1.75)	0.064*** (2.94)	-0.085*** (-2.65)	-0.016 (-0.66)	-0.014 (-0.43)	-0.001 (-0.03)	0.045 (1.34)	0.032 (0.80)	0.102	197
<i>Index Skew</i>	2.21*** (5.68)	-0.006 (-0.86)	0.012 (0.93)	0.037*** (3.01)	0.018 (0.94)	0.042*** (2.65)	-0.075*** (-3.02)	0.020 (1.11)	0.042* (1.68)	0.134	209

Note: Monthly strategy excess returns regressed on the first 8 principal components extracted from 24-month rolling windows. α is annualized (multiplied by 12). t -statistics in parentheses (Newey–West HAC). The first 8 PCs explain on average 64.7% of factor variance (min 64.7%, max 64.7%). Sample: BTP Italia 157 months, CDS–Bond 197 months, Index Skew 209 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 32
Top Factor Loadings by Principal Component

PC1		PC2		PC3	
Factor	Loading	Factor	Loading	Factor	Loading
CRED_SPR_EU	+0.203	GLOBAL_AGG	+0.300	SNRFIN_MAIN	+0.303
DEF_US	+0.197	GLOBAL_TERM	+0.282	BTP_BUND_2Y	+0.212
DEF_EU	+0.195	R10_EU	+0.281	PB_CDS_5Y_EU	+0.202
CRED_SPR_US	+0.194	GOVT_EU	+0.281	CREDIT_US	-0.197
ITRX_MAIN	+0.192	TERM_EU	+0.280	BTP_BUND	+0.192
HY_CORP	+0.189	RI_EU	+0.268	FI_VAL	+0.187

Note: Loadings are time-series averages of the eigenvector entries across all 24-month rolling windows. For each principal component, the 6 factors with the largest absolute loading are reported. Sign indicates the direction of the relationship.

Table 33
Out-of-Sample PCA Spanning Alpha

Strategy	Expanding (Panel D)	Rolling-60 (App. A.7)
btp italia	+1.61*** (2.86, IR +0.97)	+1.65*** (3.00, IR +1.02)
cds bond basis	+4.86*** (5.08, IR +1.62)	+4.10*** (3.83, IR +1.30)
itraxx combined	+1.85*** (3.75, IR +1.11)	+1.96*** (3.97, IR +1.13)

Table 34
Adaptive Elastic Net: Post-Selection OLS on Stable Factors

<i>Panel A: BTP Italia</i>		
	Coefficient	<i>t</i> -stat
α (ann. %)	+1.39***	3.34
SMB_EU	−0.044**	−2.30
ILLIQ	+0.638**	2.50
SS10Y	−0.021***	−3.26
UMD_EU	+0.030**	1.96
CDX_IG	+0.009	1.60
LIBOR_OIS	+0.490*	1.90
$T = 157, k = 6, \bar{R}^2 = 0.2173, DW = 2.0351$		
<i>Panel B: CDS–Bond Basis</i>		
α (ann. %)	+3.94***	3.66
$\Delta U F$	−4.271***	−2.61
RI_EU	+0.160***	3.27
PB_EU_CDS_1Y	+0.011**	2.02
EMERG_FX	+0.061**	1.99
SS5Y	+0.018**	2.33
CRED_SPR_US	−0.510**	−2.20
EP_SVIX_1M	+3.245***	2.82
PTFSFX	−0.450*	−1.74
$T = 197, k = 8, \bar{R}^2 = 0.2670, DW = 2.0233$		
<i>Panel C: iTraxx Combined</i>		
α (ann. %)	+0.47	1.04
EP_SVIX_1M	+3.306***	2.89
LIBOR_OIS	+0.484**	2.21
$\Delta V2X$	−0.017**	−2.55
ΔUR	+3.015**	2.35
$\Delta FAILS_PCT_TSY$	+0.402***	3.20
$T = 237, k = 5, \bar{R}^2 = 0.2359, DW = 2.2489$		

Note: Post-selection OLS on factors identified by bootstrap stability selection ($\hat{\Pi}_j \geq 0.80$). Factors are in original (un-standardized) units. α is annualized (% p.a.). Coefficients rounded to 3 decimals. *t*-statistics based on Newey–West HAC standard errors (6 lags). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. Factor definitions are provided in Table 8.

Table 35
Alpha by Stress Regime: PCA and Best-Subset

Strategy	PCA			Best-Subset		
	LOW	MEDIUM	HIGH	LOW	MEDIUM	HIGH
<i>BTP Italia</i>	+0.37 (0.65)	+1.68*** (3.17)	+3.27* (1.82)	–	–	–
<i>CDS–Bond Basis</i>	+2.93*** (3.19)	+4.87*** (4.04)	+9.04*** (5.08)	–	–	–
<i>iTraxx Combined</i>	+1.03* (1.95)	+2.23*** (4.62)	+3.45*** (4.32)	–	–	–

Note: Annualized alpha (% p.a.) per stress regime, from a single regression with three intercept dummies and betas common across regimes (no constant), Newey–West HAC standard errors. PCA: spanning regression on 8 rolling principal components. Best-Subset: post-selection OLS on the ℓ_0 -selected factors. Regimes on the 5-year iTraxx Europe Main spread: LOW < 60, MEDIUM $\in [60, 100)$, HIGH ≥ 100 bps. *t*-statistics in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 36
Factor Selection Method Comparison

Method	BTP Italia			CDS–Bond Basis			iTraxx Combined		
	k	α	$\bar{R}^2$	k	α	$\bar{R}^2$	k	α	$\bar{R}^2$
Lasso	7	+1.39*** (3.39)	0.2203	11	+3.91*** (3.44)	0.2594	13	+0.67 (1.58)	0.3016
Elastic Net	7	+1.39*** (3.39)	0.2203	11	+3.91*** (3.44)	0.2594	13	+0.67 (1.58)	0.3016
Adaptive Lasso	7	+1.39*** (3.39)	0.2203	8	+3.94*** (3.66)	0.2670	11	+0.70 (1.63)	0.3056
AEN	6	+1.39*** (3.34)	0.2173	8	+3.94*** (3.66)	0.2670	5	+0.47 (1.04)	0.2359
ALASSO-Chen	51	+0.98 (1.62)	0.4818	29	+4.97*** (2.90)	0.3509	19	+0.74* (1.68)	0.3571

Note: α is annualized (% p.a.) from post-selection OLS. k denotes the number of selected factors. t -statistics (Newey–West HAC) in parentheses. AEN, LASSO, and Elastic Net use the same tuning criterion (HQC). Adaptive LASSO (Chen) uses AIC per Chen and Chen (2008). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 37
Unconditional and Regime-Dependent Correlations

Panel A: Strategy Returns

Pair	Unconditional		By Regime		
	Pearson	<i>t</i> -stat	LOW	MED	HIGH
CDS–Bond Basis vs BTP Italia	−0.183	−1.51	−0.124	−0.284**	−0.130
CDS–Bond Basis vs iTraxx Combined	−0.019	−0.16	−0.080	−0.023	−0.030
BTP Italia vs iTraxx Combined	0.217***	2.59	0.151	0.251**	0.154

Panel B: Mispricing Changes (Δm)

Pair	Unconditional		By Regime		
	Pearson	<i>t</i> -stat	LOW	MED	HIGH
CDS–Bond Basis vs BTP Italia	−0.298***	−3.01	−0.134	−0.324***	−0.465**
CDS–Bond Basis vs iTraxx Combined	0.191*	1.70	−0.208*	0.201*	0.316**
BTP Italia vs iTraxx Combined	−0.100	−1.07	0.150	−0.147	−0.144

Note: Pearson correlations (unconditional) with HAC *t*-statistics (Newey–West). Regime columns: Pearson correlations in LOW, MEDIUM, and HIGH stress (iTraxx Main 5Y thresholds). Stars on unconditional: HAC inference. Stars on regime correlations: two-sided *t*-test with $n - 2$ degrees of freedom. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 38
Cross-Strategy Spanning Regressions (Δm)

Dependent	Regressor	β	t -stat	p	$\bar{R}^2$
CDS–Bond Basis	BTP Italia_L0	−0.148**	−2.20	0.0280	0.0913
	BTP Italia_L1	−0.010	−0.21	0.8314	
	BTP Italia_L2	−0.039	−1.18	0.2379	
	iTraxx Combined_L0	+0.370	0.91	0.3625	
	iTraxx Combined_L1	+0.710*	1.94	0.0528	
	iTraxx Combined_L2	+0.030	0.08	0.9366	
BTP Italia	iTraxx Combined_L0	−0.687	−1.40	0.1622	0.1000
	iTraxx Combined_L1	−0.950	−1.38	0.1664	
	iTraxx Combined_L2	+0.391	0.47	0.6377	
	CDS–Bond Basis_L0	−0.424**	−2.44	0.0147	
	CDS–Bond Basis_L1	+0.215	1.59	0.1124	
	CDS–Bond Basis_L2	+0.072	0.63	0.5295	
iTraxx Combined	BTP Italia_L0	−0.010	−0.76	0.4472	0.0103
	BTP Italia_L1	+0.016	0.80	0.4235	
	BTP Italia_L2	+0.028*	1.90	0.0576	
	CDS–Bond Basis_L0	+0.017	0.54	0.5921	
	CDS–Bond Basis_L1	+0.019	0.76	0.4478	
	CDS–Bond Basis_L2	+0.008	0.40	0.6924	

Note: $\Delta m_{i,t} = \alpha + \beta \Delta m_{j,t} + \varepsilon_t$. Newey–West HAC standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 39
Stress Amplification: Interaction Regressions

Dependent	Variable	Coefficient	<i>t</i> -stat	<i>p</i>	$\bar{R}^2$
CDS–Bond Basis	Δm_j (BTP Italia)	−0.144**	−2.06	0.039	0.1221
	$\Delta m_j \times z_t$	−0.086	−1.50	0.133	
	z_t (stress)	+1.591***	2.61	0.009	
BTP Italia	Δm_j (CDS–Bond Basis)	−0.464***	−3.20	0.001	0.1286
	$\Delta m_j \times z_t$	−0.637***	−2.68	0.007	
	z_t (stress)	+1.502	0.97	0.332	
CDS–Bond Basis	Δm_j (iTraxx Combined)	+0.482	1.07	0.284	0.0531
	$\Delta m_j \times z_t$	+0.016	0.08	0.939	
	z_t (stress)	+0.934**	2.11	0.035	
iTraxx Combined	Δm_j (CDS–Bond Basis)	+0.030	1.14	0.254	0.0731
	$\Delta m_j \times z_t$	+0.079*	1.72	0.086	
	z_t (stress)	−0.076	−0.27	0.784	
BTP Italia	Δm_j (iTraxx Combined)	−0.389	−0.66	0.509	0.0154
	$\Delta m_j \times z_t$	−0.618	−0.92	0.360	
	z_t (stress)	−0.128	−0.08	0.939	
iTraxx Combined	Δm_j (BTP Italia)	−0.010	−0.83	0.408	0.0201
	$\Delta m_j \times z_t$	−0.020	−1.00	0.316	
	z_t (stress)	−0.045	−0.19	0.850	

Note: $\Delta m_{i,t} = \beta_0 \Delta m_{j,t} + \beta_1 (\Delta m_{j,t} \times z_t) + \gamma z_t + \varepsilon_t$, where z_t is the standardized iTraxx Main 5Y spread. $\beta_1 > 0$: interdependence increases under stress. Newey–West HAC standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 40
AR(1) Persistence of Mispricing Levels

Strategy	ϕ	HAC t -stat	Half-life (months)	N
CDS–Bond Basis	0.7913***	14.52	3.0	148
BTP Italia	0.7272***	13.92	2.2	148
iTraxx Combined	0.7258***	9.96	2.2	148

Note: AR(1) regression of mispricing levels $m_{i,t} = c_i + \phi_i m_{i,t-1} + u_{i,t}$, monthly frequency. Half-life $h_i^* = \ln(0.5) / \ln(\phi_i)$ measures the number of months for the mispricing level to halve following a shock. Newey–West HAC standard errors with six lags. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Table 41
Common Factor (PC1) and Funding Stress

Factor	Coefficient	<i>t</i> -stat	<i>p</i> -value
Panel A: Core Intermediary Proxies			
$T = 149, \bar{R}^2 = 0.230$			
HKM	-4.236**	-2.21	0.027
LIBOR_OIS	+3.344***	3.08	0.002
PB_CDS_5Y_EU	+0.347*	1.93	0.054
ILLIQ	+1.740***	2.78	0.005
Panel C: Macroeconomic Placebo			
$T = 151, \bar{R}^2 = 0.105$			
$\Delta U M$	+0.269	0.02	0.983
$\Delta U R$	+10.366	1.23	0.220
5Y5Y_INFL	-0.606	-0.63	0.530
EPU_EU	+0.003	1.29	0.196

Note: OLS regression of the first principal component of the three Δm series on intermediary stress proxies and on a macro placebo set. Panel A: one proxy per Duffie channel (intermediary capital, funding cost, dealer health, market illiquidity). Panel B: all available funding and liquidity proxies. Panel C: macroeconomic placebo set (macro and real-activity uncertainty, long-term inflation expectations, economic policy uncertainty). Significant coefficients in Panels A and B indicate that the common factor driving correlated mispricings is linked to intermediary balance-sheet conditions [Duffie, 2010, Brunnermeier and Pedersen, 2009, He et al., 2017]. Non-significant coefficients in Panel C confirm that the common factor is not driven by macroeconomic fundamentals. Newey–West HAC standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.




Figure 8

The TIPS–Treasury Basis, 2005–2025. The series is the average TIPS–Treasury mispricing across matched bond pairs, in basis points, following Fleckenstein et al. [2014]. The vertical dashed line marks the end of their sample (19 November 2009). The basis exceeded 140 bps during the 2008–09 crisis; that extreme dislocation has not recurred, and the basis has since compressed, turning slightly negative in 2021–22.




Figure 9

Market-Level Basis Time Series for Fixed Income Arbitrage Strategies. For BTP Italia, the basis is the cross-sectional mean of all outstanding issues with at least six months to maturity. For iTraxx Combined, it is the mean of the four on-the-run sub-index skews. For CDS-Bond Basis, it is the cross-sectional median of the single-name bases in the iTraxx Main universe. Dashed lines indicate entry thresholds.




Figure 10

Cumulative Return Indices and Drawdowns of Fixed Income Arbitrage Strategies (Signal-Weighted). Each row shows one strategy: the upper sub-panel reports the cumulative return index from inception, and the lower sub-panel reports the running drawdown (peak-to-trough). Sample periods differ across strategies: BTP Italia from April 2012, iTraxx Combined from September 2005, CDS-Bond Basis from September 2008. All series end in May 2025 and are net of transaction costs.




Figure 11

Rolling 36-Month Alpha Across Benchmark Factor Models. Each panel shows one strategy. Lines represent rolling alpha under Duarte et al. (2007, blue), Brooks et al. (2020, orange), and Fung & Hsieh (2004, green). Red shading: HIGH stress regime (iTraxx Main 5Y > 100 bps).




Figure 12

Scree Plot: Cumulative Variance Explained by Principal Components. Dashed line indicates 80% threshold. Baseline specification uses $K = 8$ components.




Figure 13

Pairwise Correlations by Stress Regime. Bars show Pearson correlations of Δm in LOW, MEDIUM, and HIGH stress regimes defined by iTraxx Main 5Y thresholds (60/120 bps).




Figure 14

Generalized Impulse Response Functions. Responses of Δm to one-standard-deviation shocks. Shaded areas: 95% bootstrap confidence intervals (1000 replications).