

Exploring Excess Return Generation through European Fixed Income Arbitrages

Alessio Ottaviani

EDHEC Business School

Supervisor: Prof. Riccardo Rebonato

Abstract:

Contents

List of Tables

List of Figures

Introduction

? pose the question that motivates our analysis: do fixed-income arbitrage strategies earn genuine alpha, or are their returns simply compensation for systematic risk—“picking up nickels in front of a steamroller”? Using a ten-factor spanning regression, they find that the more complex strategies produce significant alpha that survives equity and bond market risk controls. We take this question to a new setting—European fixed-income markets—and to three strategies that have not previously been studied jointly as systematic arbitrages, namely the BTP Italia inflation-linked basis, the European corporate CDS–bond basis, and the iTraxx index skew.

A second set of questions then arises. Does this alpha withstand a far more demanding test of risk—and which risks does each strategy ultimately bear? Is the alpha that survives roughly constant, or time- and state-dependent? And, perhaps above all, why does it persist? The presidential address of ? provides a framework: when intermediary balance-sheet capacity is the binding constraint on arbitrage, mispricings across different markets should co-move, the co-movement should intensify in stress, and the transmission should follow a liquidity pecking order from more liquid to less liquid markets. ? provide the mechanism—shocks to funding liquidity and market liquidity reinforce each other in a spiral that forces simultaneous deleveraging across positions—and ? document the empirical counterpart, showing that arbitrage crashes propagate across strategies as capital withdraws.

Each of the three arbitrage strategies we study has an established literature, but none has been analysed in the way we propose. The inflation-linked bond basis was documented by ? in the context of U.S. Treasury Inflation-Protected Securities (TIPS), who found an average mispricing of 54.5 basis points between TIPS and nominal Treasuries over the 2004 to 2009 period, with peaks exceeding 200 basis points during the 2008 to 2009 financial crisis (Figure ??). This basis has since compressed substantially, oscillating in a 0–30 bps band post-2009, reflecting broader institutional participation in the U.S. inflation-linked market—and, as ? show, the basis narrows as hedge fund capital flows into the trade, providing direct evidence for the slow-moving capital channel. In their

Section VII, they further document that changes in the TIPS–Treasury mispricing are correlated with changes in the corporate CDS–Bond basis and the CDX index–component basis, establishing the cross-strategy co-movement that we extend to the European setting in Section ???. In contrast, the analogous basis on Italian BTP Italia bonds persists to date. To the best of our knowledge, the BTP Italia has not been the subject of any prior academic study, and its inflation-linked basis has not previously been analysed as a systematic arbitrage strategy. Its persistence is the sharper puzzle because the BTP Italia eliminates the ballooning risk that characterises TIPS and other conventional inflation-linked sovereigns,¹ and because its retail-dominated investor base creates a structural segmentation that insulates the basis from the institutional capital flows that closed the U.S. mispricing.

Moving to the second strategy we analyse, the CDS–Bond basis is the difference between the CDS spread and the bond credit spread for the same reference entity, a relationship that should be zero by no arbitrage in a frictionless market. It was first documented by ? and ? in its pre-crisis positive form. ? show that during 2007 to 2009, tightening financial constraints drove the basis deeply negative, and ? provide a comprehensive analysis of its cross-sectional determinants and its connection to limits of arbitrage. However, this literature focuses predominantly on U.S. corporate bond markets, and little attention has been paid to the euro-area corporate counterpart, which is the second focus of this study.

This brings us to the third strategy of interest, the CDS index skew, which is defined as the deviation of the traded index spread from its intrinsic, no-arbitrage value implied by the single-name constituents. In a U.S. context, this strategy has been studied by ?, who use it as a market-wide liquidity measure, and by ?, who document that post-crisis regulatory capital requirements have made basis trades between the index and single-name CDS “significantly less economical,” leading to persistently wider skews; both papers focus on the U.S. CDX index. To the best of our knowledge, neither the negative basis trade on European investment-grade corporates in the

¹Conventional inflation-linked bonds adjust the principal upward with realised inflation; in deflationary scenarios or when inflation expectations spike, the redemption value can grow to many multiples of par (the so-called *ballooning* risk), introducing path-dependent valuation risk. The BTP Italia instead pays inflation revaluation as part of each semiannual coupon and redeems at par regardless of the inflation path, eliminating this source of uncertainty.

iTraxx Main universe nor the index skew across the four iTraxx sub-indices has been studied as a systematic arbitrage strategy in the academic literature.

Beyond being the first study to analyse these three strategies jointly—and, for the BTP Italia, the first to analyse it at all—our paper extends the ? framework in three respects:

1. Each of our three strategies has a deterministic payoff at maturity: if the position can be held to maturity, the profit is certain regardless of the path of market variables in the interim. The strategies in ?, by contrast—swap spreads, yield curve trades, mortgage hedging, volatility selling, capital structure trades—do not share this property: their payoffs depend on models, mean-reversion assumptions, or stochastic processes (prepayment, realised volatility) that leave the outcome uncertain even at maturity. This yields a cleaner test of the alpha-versus-risk question: when the terminal payoff is certain, a positive average return cannot be compensation for payoff risk, and any surviving alpha points to limits to arbitrage rather than to a missing risk premium.
2. We subject the alpha estimates to a progressive battery of factor controls: three classical benchmark specifications from the fixed-income hedge fund literature (???), a full-sample principal component analysis following ??, and best-subset selection (?), which picks for each strategy the eight most informative factors from a candidate universe of 75 tradable risk factors. Every layer is re-estimated out of sample, with the factor hedge formed from past-only loadings, so the surviving alpha is an abnormal return implementable in real time rather than an artefact of in-sample factor combination; the selected models are re-run under five alternative selectors—from the adaptive elastic net (?) to a nonlinear machine-learning screen and model-X knockoffs (?), which at a controlled false discovery rate select no factor at all—with honest split-sample inference, and the alpha verdict is invariant to the choice. Using the conditional framework of ? we show that the alpha rises systematically during periods of intermediary funding stress.
3. We test the cross-strategy implications of slow-moving capital systematically, drawing on ?, ?, ?, and ??: whereas the existing literature has studied each of these arbitrages in isolation, we examine all three jointly in the European market and ask whether their mispricings co-

move, whether the co-movement intensifies in stress, and whether a common factor tied to intermediary capital—rather than to macroeconomic fundamentals—explains it. The European setting provides a discriminating test of the mechanism: all three arbitrages are run by the same type of institutional arbitrageur, but the underlying instruments differ in who holds them. The CDS–Bond and iTraxx legs are held and arbitrated by leveraged institutions; the BTP Italia basis, by contrast, rests on an inflation-linked bond held predominantly by unlevered retail households, who face no margin calls and do not unwind their positions under funding stress. We relate the resulting evidence to the structural limits to arbitrage that keep the mispricings open, including the post-crisis regulatory capital costs documented by ?.

We find that all three strategies generate positive and economically significant excess returns net of transaction costs, with annualised Sharpe ratios from 0.98 (BTP Italia) to 1.56 (CDS–Bond Basis), positive skewness, and contained drawdowns. The alpha survives every layer of controls: the three benchmark frameworks, the principal components, and the best-subset models, which explain up to 31% of monthly return variation yet leave annualised alphas of +2.2%, +4.8%, and +2.9%, all significant at the 1% level, in and out of sample. The three selected factor sets are pairwise disjoint, so no single common exposure is responsible; the iTraxx alpha loads on the traded intermediary capital factor of ?, paying off when dealer equity underperforms. The residual alpha rises during periods of intermediary funding stress, in a manner consistent with the slow-moving capital mechanism of ?, and the iTraxx skew serves as the leading mispricing in the cross-strategy hierarchy, Granger-causing both the CDS–Bond ($p = 0.011$) and the BTP Italia ($p = 0.007$) mispricings. The co-movement of the three mispricings is governed by a common factor that loads on intermediary balance-sheet variables—the ? capital ratio, the Libor–OIS spread, and the ? illiquidity measure explain 23.3% of its variation—and not on macroeconomic fundamentals: a placebo regression on macro indicators yields uniformly insignificant coefficients and an adjusted R^2 of 10.6%. The persistence of the mispricing levels partially mirrors the funding-dependence of each arbitrage trade: the CDS–Bond basis, which requires continuous repo financing of the cash bond, has the longest half-life at approximately three months, roughly one-third longer than

the BTP Italia and the iTraxx skew, which mean-revert within two months. In monthly mispricing changes, the institutional pair (CDS–Bond and iTraxx) co-moves ever more tightly as stress builds, with correlations rising monotonically from -0.21 in low-stress to $+0.32$ in high-stress regimes. The retail-held BTP Italia basis moves the other way: its correlation with the CDS–Bond basis is significantly negative on the full sample at -0.30 , deepening to -0.47 in high-stress months. This opposing behaviour provides independent evidence of investor-base segmentation: a single intermediary-capital shock pushes the two institutional mispricings wider and the unlevered, retail-held BTP Italia tighter.

The paper is organized as follows. Section ?? describes the three strategies, their construction, and their performance. Section ?? evaluates the alpha against three classical benchmark frameworks. Section ?? extends the analysis to a candidate universe of 75 tradable factors using principal components and best-subset selection. Section ?? tests the cross-strategy predictions of the slow-moving capital literature. Section ?? concludes.

I Data, Strategy Construction, and Performance

This section introduces the three European fixed-income arbitrage strategies that the paper studies: the inflation-linked basis on Italian government bonds (BTP Italia), the negative basis between corporate bonds and their credit default swaps (CDS–Bond Basis), and the CDS index skew—the difference between the traded iTraxx spread and the intrinsic spread implied by its single-name constituents (iTraxx Skew). Section ?? describes the data and the construction of each trade, with particular attention to the two features of the BTP Italia—the absence of ballooning risk and the retail bondholder base—on which the identification strategy of Section ?? rests. Section ?? aggregates individual trades into investable indices, compares equal-weighted and signal-weighted schemes, and introduces the regime-dependent transaction costs charged throughout the paper. Section ?? measures the performance of the resulting indices and subjects it to manipulation-proof and volatility-managed diagnostics.

A Data and Strategy Construction

Data are sourced from the MOT market operated by Borsa Italiana (BTP Italia and nominal BTP prices), Bloomberg (inflation swap rates and corporate bond spreads), CMA/CME Group (single-name CDS spreads), and J.P. Morgan DataQuery (iTraxx index and constituent data, cross-validated against Barclays Live). The iTraxx skew sample begins in June 2004 with the launch of the first on-the-run series of the iTraxx Europe Main 5-year index, the CDS–Bond sample begins in September 2008, and the BTP Italia sample begins in April 2012, the month after the first issuance. All three samples extend to May 2025. Following ?, each individual trade is modelled as a separate fund that opens when the basis crosses an entry threshold and closes when the basis reverts through an exit threshold or the underlying instrument matures. Entry and exit rules are summarised in Table ??; trade-level statistics are reported in Table ?. Over the full sample, the strategies generate 260 trades for BTP Italia, 931 for CDS–Bond Basis, and 206 for iTraxx Skew. Figure ?? plots the market-level basis time series for each strategy.

BTP Italia: The Inflation-Linked Basis. The BTP Italia is an Italian government bond indexed to domestic inflation, measured by the ISTAT FOI consumer price index (excluding tobacco). It differs in two economically meaningful ways from conventional inflation-linked sovereigns such as the U.S. TIPS studied by ?.

The first is the absence of *ballooning risk*. A conventional inflation-linked bond accumulates the inflation adjustment in the principal over the life of the bond, so that the notional outstanding at time t equals $100 \times I_t$, where I_t is the cumulative inflation factor since issuance. In a credit event, the recovery rate (RR) is applied to the original face value of 100, not to the inflation-adjusted notional: an investor holding a TIPS with $I_t = 1.25$ recovers $RR \times 100$ and loses the entire accumulated revaluation of 25 percentage points. The BTP Italia eliminates this exposure. The capital revaluation is paid out semiannually alongside the coupon, so the inflation factor resets to unity at every coupon date and the notional exposed to credit risk remains at par throughout the life of the bond. The two legs of the arbitrage therefore carry identical sovereign credit exposure, which

isolates a pure relative-value component in the basis and removes the differential-default contamination that affects the TIPS–Treasury basis in stress periods. Appendix ?? reports the technical specification of the indexation coefficient and of the deflation floors on coupon and principal.

The second feature is the investor base. The Italian Treasury introduced the BTP Italia in 2012 explicitly as a retail instrument, structuring its primary placement in two sequential phases: an initial phase reserved exclusively for retail investors, which typically absorbs the bulk of each issuance, followed by a shorter institutional phase. Investors who purchase at issuance and hold to maturity receive an additional loyalty bonus (the *premio fedeltà*) that further incentivises buy-and-hold behaviour. The BTP Italia is therefore held predominantly by households with long holding periods and no regulatory balance-sheet constraints. This segmentation matters because, as we document in Section ??, it is what makes the BTP Italia basis behave differently from the two institutional credit strategies during stress episodes: retail holders face no margin calls and do not unwind their positions when intermediary capital withdraws.

We construct the BTP Italia basis following the replication logic of ?, adapted to the Italian instrument. The inflation-linked cash flows of the BTP Italia are converted into fixed-rate equivalents using the term structure of zero-coupon inflation swaps, yielding a fixed-rate equivalent yield y_t^{IL} . The basis is the difference between this yield and the yield to maturity of the maturity-matched nominal BTP y_t^{NOM} :

$$\text{Basis}_t^{\text{BTP}} = y_t^{\text{IL}} - y_t^{\text{NOM}}. \quad (1)$$

A positive basis indicates that the BTP Italia, once converted to fixed, yields more than the nominal bond, and vice versa; the trade is executed accordingly through an asset swap (ASW) on the BTP Italia against an offsetting position in the maturity-matched nominal BTP. Both the BTP Italia and the nominal BTP are accepted as general collateral in the euro repo market, so the two legs carry identical funding costs, which rules out funding-cost differentials as an explanation for the observed basis.

CDS–Bond Basis: The Negative Basis Trade. The CDS–Bond basis for issuer i at date t is defined as the difference between the CDS spread and the z-spread of a reference corporate bond:

$$\text{Basis}_{i,t} = \text{CDS spread}_{i,t} - \text{z-spread}_{i,t}, \quad (2)$$

where the z-spread is the constant spread that, when added to the bootstrapped euro zero-coupon swap curve, equates the discounted bond cash flows to the observed market price. When the basis is negative, credit protection via CDS is cheaper than the credit spread embedded in the bond price. The negative basis trade exploits this gap: the arbitrageur purchases the bond (earning the z-spread) and simultaneously buys CDS protection (paying the lower CDS spread), locking in a positive carry that persists until convergence or bond maturity. We restrict the strategy to the negative basis trade; the reverse trade (short bond, sell protection) requires sourcing the corporate bond through a reverse repo, which for corporate names is subject to recall risk—particularly during periods of stress, when the securities lender may recall the borrowed bond and force premature unwinding of the position (?). In a default event, the bondholder recovers $\text{RR} \times 100$ on par while the CDS contract pays $(1 - \text{RR}) \times 100$, so the combined position returns par regardless of realised recovery—provided the CDS maturity is no shorter than the bond’s, which we ensure by matching hedge tenors to bond residual life (Appendix ??). The tradable universe is restricted to single-name CDS whose reference entity belongs to the current on-the-run iTraxx Europe Main 5-year index, which ensures liquidity and provides a transparent filter for investment-grade credit quality. When a name exits the on-the-run index at a roll date, no new trades are initiated on that name; existing positions are carried to convergence or maturity. All CDS are centrally cleared through LCH or ICE, eliminating direct counterparty credit risk on the protection leg.

iTraxx Skew: The CDS Index Skew. The CDS index skew is the difference between the traded spread of an iTraxx index and its intrinsic spread—the no-arbitrage level implied by the con-

stituents' CDS curves under standard market conventions:

$$\text{Skew}_t = S_t^{\text{index}} - S_t^{\text{intrinsic}}. \quad (3)$$

A positive skew indicates that the index trades wide relative to the sum of its parts, and vice versa.² Three forces sustain a non-zero equilibrium skew. First, the index is far more liquid than its single-name constituents, especially in stress (?). Second, the two legs face different sources of demand: institutional investors use the index as the natural macro credit hedge, while dealer banks trade single names for more granular hedging needs, including the hedging of counterparty CVA exposure—forces that can push the two legs in non-synchronised directions and make the skew bidirectional. Third, executing the replicating basket of single-name positions carries operational costs that the index trade does not. Our strategy trades the skew across the four iTraxx sub-indices—Main, Senior Financials, Subordinated Financials, and Crossover—restricting new positions to the on-the-run series at the time of entry. Entry thresholds are calibrated to the historical volatility of each sub-index skew (Table ??). Unlike the CDS–Bond basis, the iTraxx skew is a derivative-on-derivative trade with no cash leg financed in repo. The two legs are traded under a single netting set—either centrally cleared at the same CCP or bilaterally with the same dealer under an ISDA Master Agreement and CSA—so that long and short positions offset for daily variation margin and balance-sheet purposes.

B Index Construction and Weighting Schemes

Following ?, individual trades are treated as separate fictional hedge funds, each with its own entry date, exit date, and daily return stream. The strategy index on any given day is the weighted average of the returns of the funds that are active on that day. The equal-weighted (EW) scheme of ? assigns

²Each constituent's survival curve is bootstrapped from its CDS quotes, and the index's intrinsic spread is the single running spread that reprices the resulting portfolio under the standard quoting convention (flat hazard rate, 40% recovery; see ?).

identical weight $1/K_t$ to every active trade:

$$r_t^{\text{EW}} = \frac{1}{K_t} \sum_{k=1}^{K_t} r_{k,t}. \quad (4)$$

The signal-weighted (SW) scheme, inspired by the curvature signal of θ , weights each trade proportionally to the absolute distance between the basis at entry and its historical equilibrium level,

$$r_t^{\text{SW}} = \sum_{k=1}^{K_t} \frac{|b_{k,0} - \theta_{t_k}|}{\sum_{j=1}^{K_t} |b_{j,0} - \theta_{t_j}|} r_{k,t}, \quad (5)$$

where $b_{k,0}$ is the basis at the entry date t_k of trade k , θ_{t_k} is the expanding-window average of the cross-sectional basis from the start of the sample up to t_k , and $r_{k,t}$ is the signed return of trade k , positive when the position profits, with the direction—long or short—fixed at entry. The absolute value makes the weight reflect the magnitude of the dislocation rather than its sign, which $r_{k,t}$ already carries. The rationale follows θ : a basis that deviates further from its historical equilibrium represents a richer opportunity and warrants a larger allocation conditional on convergence. Signal-weighted indices dominate equal-weighted indices on both Sharpe and manipulation-proof performance for all three strategies (Appendix ??, Table ??), and we adopt the SW scheme as the baseline throughout the paper; EW results in the appendix confirm that all findings are robust to this choice.

Market conditions are classified into three regimes based on the level of the iTraxx Europe Main 5-year spread: LOW (below 60 bps), MEDIUM (60 to 100 bps), and HIGH (100 bps and above). The same classification recurs in Section ??, where we decompose alpha by regime. Transaction costs are applied at the individual trade level as entry and exit fees proportional to DV01, and fee levels are regime-dependent: bid-offer spreads and fees rise monotonically from LOW to HIGH, reflecting the empirical widening of dealer quotes during market dislocation. This tiered structure makes costs conservative precisely in the stress periods when arbitrage returns are largest. The full fee schedule is reported in Table ??; from this point forward, all strategy returns are net of these regime-dependent costs.

C Performance Analysis

Table ?? reports monthly summary statistics for the three strategy indices. Throughout the paper, t -statistics use Newey–West heteroskedasticity- and autocorrelation-consistent standard errors (?), with the truncation lag set by the rule of ?, $m = \lfloor 4(T/100)^{2/9} \rfloor$, which equals four months for every sample we use. All strategy returns are excess returns by construction: the BTP Italia trade is self-financing (long and short legs are both repoed at the same general collateral rate), the CDS–Bond trade earns the z-spread net of the CDS premium (where the z-spread is itself measured over the swap curve, which proxies for the funding cost), and the iTraxx skew trade requires no capital beyond initial margin. Panel A describes the basis levels: all three series are persistent, with monthly first-order autocorrelations between 0.66 and 0.83, consistent with the slow-moving capital literature (???) that we examine formally in Section ?. Panel B describes the strategy returns. The CDS–Bond basis produces the highest annualised return, 4.61% ($t = 5.87$), and Sharpe ratio, 1.56, with an annualised volatility of 2.96%. iTraxx Skew returns 1.91% p.a. ($t = 4.92$) with a Sharpe ratio of 1.25 over the longest sample of the three. BTP Italia earns 1.59% p.a. ($t = 3.52$) with a Sharpe ratio of 0.98 over the shortest. All three strategies exhibit positive return skewness—0.81 for BTP Italia, 0.79 for the CDS–Bond basis, and 1.01 for iTraxx Skew—which cuts against the popular characterisation of fixed-income arbitrage as “picking up nickels in front of a steamroller” (?). Maximum drawdowns are moderate and short-lived, as Figure ?? illustrates: the deepest drawdown is 4.4% for the CDS–Bond basis during the 2008 to 2009 credit crisis, 3.0% for BTP Italia, and 2.5% for iTraxx Skew, and all three are recovered within months.

Manipulation-Proof Performance. Sharpe ratios are vulnerable to manipulation by strategies whose dynamic trading rules alter the shape of the return distribution, which is precisely the structure that defines the indices we construct. Table ?? therefore reports the Manipulation-Proof Performance Measure (MPPM) of ? for risk-aversion parameters $\rho \in \{2, 3, 4\}$: a single-valued score that is independent of dollar scale, cannot be enhanced by information-free trading or leverage, and

is consistent with standard financial market equilibrium—the appropriate benchmark for strategies, like ours, that use dynamic entry and exit rules.

The MPPM is nearly invariant to the risk-aversion level. For BTP Italia, the measure is 1.50, 1.48, and 1.47% per year at $\rho = 2, 3, 4$; for CDS–Bond Basis, 4.43, 4.39, and 4.35; for iTraxx Skew, 1.88, 1.87, and 1.86. Because the concave transformation $(1 + r_t)^{1-\rho}$ penalises strategies that earn their Sharpe ratio through occasional large losses, this stability—read together with the positive return skewness documented above—shows that the strategies’ risk-adjusted performance does not rely on tail events: an investor with $\rho = 4$ evaluates these strategies almost identically to an investor with $\rho = 2$.

Volatility-Managed Portfolios. A second concern the literature has raised about fixed-income arbitrage is that unconditional Sharpe ratios may average over a time-varying risk-return trade-off, with periods of attractive compensation alternating with periods in which risk is poorly priced. ? propose a simple test. The managed portfolio scales exposure inversely to the previous month’s realised variance,

$$f_{t+1}^{\sigma} = \frac{c}{\hat{\sigma}_t^2} f_{t+1}, \quad (6)$$

where $\hat{\sigma}_t^2$ is the realised variance over month t and the scaling constant c matches the unconditional standard deviation of the unmanaged portfolio. A significant managed-portfolio alpha relative to buy-and-hold is diagnostic of a time-varying ratio μ_t/σ_t^2 : an investor earns additional risk-adjusted returns by taking less risk precisely when volatility is high, which can only happen if the risk-return trade-off itself varies over time.

Table ?? reports the results. The CDS–Bond basis produces a positive managed-portfolio alpha of 1.39% per year for the SW index, significant at the 10% level, consistent with a time-varying funding-liquidity mechanism: when volatility is low, intermediary balance sheets are healthy and arbitrage capital is available to exploit the negative basis; when volatility spikes, capital withdraws and the trade-off deteriorates. For BTP Italia and iTraxx Skew, the managed-portfolio alphas are slightly negative (−0.22% and −0.16% p.a., both insignificant), so the realised-variance signal

identifies no exploitable variation in the risk-return trade-off for these two strategies. The contrast between the one strategy that finances a cash bond and the two that do not is itself informative: Section ?? shows that, once benchmark factor exposures are controlled for, the funding-liquidity pattern documented here for the CDS–Bond basis generalises into a regime effect common to all three strategies.

II Benchmark Factor Models

This section tests whether the alpha of the three strategies survives the canonical benchmark spanning regressions of the fixed-income arbitrage literature—?, ?, and ?—and whether the residual alpha varies with intermediary funding stress, as the slow-moving capital framework of ? predicts. Section ?? describes the three specifications and the construction of their euro-denominated counterparts. Section ?? presents the regression results and three complementary evaluations: the alpha by stress regime, the conditional alpha model of ?, and an out-of-sample re-estimation in which the benchmark hedge is formed from past-only loadings.

A Benchmark Specifications

Three benchmark specifications are widely used in the empirical literature on fixed-income arbitrage and active fixed-income management, and they provide the natural reference points against which our strategies should be evaluated.

?, in the paper that motivates our research design, test whether fixed-income arbitrage excess returns are compensation for exposures to equity, banking-sector, and term-structure risk:

$$r_{i,t} = \alpha_i + \beta_1 Mkt_t + \beta_2 SMB_t + \beta_3 HML_t + \beta_4 UMD_t + \beta_5 RS_t + \beta_6 RI_t + \beta_7 RB_t + \beta_8 R2_t + \beta_9 R5_t + \beta_{10} R10_t + \varepsilon_{i,t}, \quad (7)$$

where Mkt , SMB , HML , and UMD are the Fama–French–Carhart equity factors (market, size, value, and momentum), RS is the S&P bank stock index excess return, RI and RB are A/BAA-

rated industrial and bank bond index excess returns, and $R2$, $R5$, $R10$ are the CRSP Fama bond portfolios at two-, five-, and ten-year maturities.

? propose a seven-factor *asset-based style* model that tests whether hedge fund returns are explained by equity, fixed-income, and trend-following exposures:

$$r_{i,t} = \alpha_i + \beta_1 S\&P_t + \beta_2 SC-LC_t + \beta_3 BdOpt_t + \beta_4 FXOpt_t + \beta_5 ComOpt_t + \beta_6 10Y_t + \beta_7 CredSpr_t + \varepsilon_{i,t}, \quad (8)$$

where $S\&P$ is the S&P 500 excess return, $SC-LC$ is the small-minus-large-cap total return, $BdOpt$, $FXOpt$, and $ComOpt$ are the three trend-following factors constructed as lookback straddle returns on bond, currency, and commodity futures, $10Y$ is the excess return on a 10-year government bond total-return index, and $CredSpr$ is the return spread between BAA-rated corporate bonds and 10-year government bonds.

? take a different angle, decomposing the returns of active fixed-income managers into eight *traditional risk premia* organised along three channels of risk-taking—duration, credit, and volatility selling:

$$r_{i,t} = \alpha_i + \beta_1 Term_t + \beta_2 GlTerm_t + \beta_3 GlAgg_t + \beta_4 InflLink_t + \beta_5 CorpCred_t + \beta_6 EmDebt_t + \beta_7 EmCurr_t + \beta_8 RatesVol_t + \varepsilon_{i,t}, \quad (9)$$

where $Term$, $GlTerm$, $GlAgg$, and $InflLink$ are the total returns on the Treasury, the currency-hedged global Treasury, global aggregate, and global inflation-linked bond indices, each in excess of cash, $CorpCred$ is a high-yield corporate bond index in excess of cash, $EmDebt$ is the excess return on a hard-currency emerging-market debt index, $EmCurr$ is the return on an equal-weighted basket of emerging-market currencies, and $RatesVol$ is the excess return on a one-month variance swap on the ten-year Bund futures.

Two ingredients of the original specifications are replaced with tradable counterparts. In ?, the bond and credit factors are yield and spread *changes*; a yield change is not the return on an in-

vestable position, so we use the corresponding total-return series (the *10Y* and *CredSpr* definitions above). In ?, *CorpCred* blends high yield with leveraged loans; loan indices are quote-based and not investable at the monthly horizon, so we retain the high-yield leg only. With every regressor a realised excess return, the intercept in each specification keeps its interpretation as a Jensen alpha.

Each benchmark is estimated in two parallel panels. The first uses the original U.S.-denominated factors as specified by the respective authors, which is the natural way to replicate each framework as it was first published. The second uses euro equivalents, since all three strategies trade European instruments and may therefore be better explained by European factor exposures. The euro panel is the primary specification and is reported in the main text; the US-factor estimates are reported in Appendix ??.

For the euro panel, the Fama–French factors—used directly in ? and indirectly in ?—are sourced from the Kenneth French data library; because these are reported in USD, we convert them to EUR following ?. For long factors (market excess return):

$$F_t^{\text{EUR}} = \frac{1 + F_t^{\text{USD}} + r_{f,t}^{\text{USD}}}{1 + r_t^{\text{USD/EUR}}} - 1 - r_{f,t}^{\text{EUR}}, \quad (10)$$

and for long-short factors (SMB, HML, UMD):

$$LS_t^{\text{EUR}} = \frac{LS_t^{\text{USD}}}{1 + r_t^{\text{USD/EUR}}}, \quad (11)$$

where F_t^{USD} and LS_t^{USD} are the original factors in US dollars, $r_{f,t}^{\text{USD}}$ and $r_{f,t}^{\text{EUR}}$ are the US and euro risk-free rates (Fama–French T-bill rate and one-month German bill, respectively), and $r_t^{\text{USD/EUR}}$ is the period return on the dollar-to-euro exchange rate. All remaining factors enter the euro panel with native European counterparts: European bank stocks, euro-area corporate bond indices, and German government bonds in the Duarte specification; euro-area term, credit, and inflation-linked premia in Brooks et al.; and the 10-year Bund total-return index and the BBB–Bund return spread in Fung and Hsieh. The Fung and Hsieh trend-following lookback straddle factors are currency-

independent by construction (D. Hsieh, private communication) and enter both panels identically. The complete factor list is reported in Appendix ??.

All specifications are estimated by ordinary least squares with Newey–West HAC standard errors. Variance inflation factors for each specification are reported in Appendix ?. Two groups of regressors are highly collinear—the maturity-sorted bond portfolios $R2$, $R5$, $R10$ in Duarte and the global bond benchmarks in Brooks et al.—which inflates the standard errors of those loadings but leaves inference on the intercept unaffected.

B Results and Discussion

Factor Model Regressions. Tables ??, ??, and ?? report the full euro-panel regression results; the US-factor counterparts are Tables ??, ??, and ?? in Appendix ?. Across all three frameworks, all three strategies, and both currency denominations, the estimated intercept is positive and significant at the 1% level. The annualised alphas are economically meaningful: +4.0 to +4.2% p.a. for the CDS–Bond basis, +1.7 to +2.2% for the iTraxx skew, and +1.6 to +2.1% for BTP Italia.

The adjusted R^2 is low across all specifications, ranging from zero to 18%, so the standard factor sets leave most of the return variation unexplained—a question we take up in Section ?. The significant loadings that do emerge differ across strategies and are economically intuitive. The CDS–Bond basis loads on credit and rates returns: the A-rated industrial bond return is positive and significant in the Duarte specification, and the 10-year government bond and credit-spread returns are positive and significant in Fung–Hsieh, in both currency denominations—the natural exposures of a long-credit, cash-financed position. BTP Italia and iTraxx Skew load instead on equity factors, with small negative coefficients: market and size for BTP Italia, and size, value, and momentum for iTraxx Skew, in both the Duarte and the Fung–Hsieh specifications. The Brooks et al. specification, which uses broad risk premia rather than individual asset returns, produces fewer significant loadings but delivers comparable alpha estimates, reinforcing that the intercepts are not an artefact of any single factor set.

Cross-Model Comparison and Regime Analysis. Table ?? consolidates the alpha estimates from the three frameworks. Panel A reports the unconditional alphas side by side: for each strategy, all three frameworks deliver positive and significant alpha, and the estimates are consistent across specifications despite their different factor structures.

The discrete regime decomposition—alpha estimated separately in the LOW, MEDIUM, and HIGH stress regimes on the iTraxx Europe Main 5-year spread, with the same 60 and 100 basis-point cutoffs that govern the transaction-cost tiers of Section ?? (Appendix ??)—is reported for each framework in Appendix ?. In every cell of those tables, alpha rises monotonically from the LOW to the HIGH regime, and the HIGH-regime alpha is three to seven times the LOW-regime alpha across the nine strategy–framework combinations. The mechanism is the one ? formalises: when intermediary capital is abundant, arbitrage forces close mispricings quickly and compress the residual premium; when it is scarce—as it is during episodes of widening credit spreads—arbitrageurs retreat or are forced to unwind, mispricings widen, and the premium for those who can still deploy capital rises. This regime dependence motivates the cross-strategy analysis of Section ??, where we test whether the same intermediary balance-sheet shocks drive co-movement of mispricings across the three strategies.

Conditional Alpha. Panel B of Table ?? replaces the discrete regime split with the conditional performance evaluation framework of ?, in which the alpha varies continuously with a predetermined, publicly observable state variable. We estimate the full conditional model

$$r_{i,t} = \alpha_0 + \alpha_1 z_{t-1} + (\beta + \delta z_{t-1})' \mathbf{X}_t + \varepsilon_t, \quad (12)$$

where $\mathbf{X}_t$ collects the factor returns of the chosen benchmark and z_{t-1} is the one-month-lagged level of the iTraxx Europe Main 5-year spread, standardised to mean zero and unit variance over the estimation sample. Lagging the state variable keeps it in the investor’s information set, as the conditional framework requires, and conditioning the loadings through δ prevents time-varying betas from masquerading as conditional alpha. The intercept α_0 is then the alpha at the average

stress level, and α_1 is the additional alpha per one-standard-deviation increase in funding stress. Because fixed-bandwidth HAC t -tests over-reject in conditional models at these sample sizes (??), significance of α_1 is assessed by a moving-block residual bootstrap with the null imposed ($B = 9,999$, block length nine months); HAC inference, reported as a robustness check, gives the same pattern.

The estimates of α_1 in Panel B are positive in all nine strategy–framework cells, and the intercept α_0 remains significant in every case, so a baseline alpha persists at the average stress level. For the two institutionally held strategies, α_1 is significant in every framework: a one-standard-deviation increase in the iTraxx Main spread adds 1.7 to 3.2 percentage points of annual alpha for the CDS–Bond basis and 0.9 to 1.5 for iTraxx Skew. For BTP Italia, α_1 is significant only under the Duarte specification—a weaker stress sensitivity consistent with the retail holder base that Section ?? develops. Across the nine cells, the implied alpha at $z = +1\sigma$ is two to nine times the alpha at $z = -1\sigma$: the bulk of the unconditional alpha is earned when funding stress is elevated, the signature of the slow-moving capital mechanism.

Out-of-Sample Alpha. A residual concern, central to the recent literature on performance evaluation (??), is that the benchmark loadings in Tables ??–?? are estimated in sample, so the surviving alpha could reflect an in-sample-optimal combination of the factors rather than an abnormal return an investor could have earned in real time. Panel C of Table ?? addresses this directly. For each strategy and each benchmark, we re-estimate the loadings on an expanding past-only window in the recursive out-of-sample design of ?, form the factor hedge from the lagged slopes, and define the out-of-sample alpha as the mean of the realised hedged return $e_{i,t} = r_{i,t} - \hat{\beta}'_{i,t-1} \mathbf{f}_t$, tested with the ? HAC correction. The expanding window is initialised with a 60-month minimum, so that the first fit has at least five observations per parameter under the ten-factor ? specification, and windows with near-singular factor designs are skipped. The out-of-sample alpha remains positive and significant for all three strategies and all three frameworks. The point estimates are close to the in-sample alphas of Panel A, with only modest attenuation, and the ranking across strategies is

preserved, with the CDS–Bond basis the largest. The rolling-window counterpart, which additionally relaxes the assumption of constant loadings, is reported in Appendix ?? and delivers the same conclusion.

Rolling Alpha. A natural concern with regime-based or conditional analyses is that the estimated relationship between alpha and stress may be driven by a single sub-period—the global financial crisis of 2008 to 2009, the European sovereign debt crisis of 2011 to 2012, or the COVID-19 shock of March 2020—rather than by a stable feature of the data. Figure ?? plots rolling 36-month alpha estimates for each strategy under each of the three benchmark specifications. The rolling alpha is positive across essentially the whole sample for all three strategies, so the estimates are not the product of a single favourable sub-period, and the three benchmark estimates move broadly together within each panel, reinforcing the cross-model robustness of Panel A. Detailed sub-period splits for each framework are reported in Appendix ??.

Taken together, the five sets of results in this section establish two findings that are central to the rest of the paper. First, the alpha is robust: it survives three widely used benchmark frameworks, in both currency denominations, in every sub-period we examine, and when each benchmark hedge is re-estimated out of sample. Second, the alpha is not constant: it rises systematically when intermediary capital is scarce, whether stress enters as a discrete regime or as a continuous predetermined state, and the high-stress alpha is a multiple of the calm-market alpha. The first finding says the observed excess returns are not explained by the systematic exposures the prior literature has identified; the second points to an economic mechanism but not yet to the specific risk factors behind the residual variation. Section ?? takes up that question on a much larger candidate factor universe.

III Factor Selection and Spanning Regressions

The benchmark factor models of Section ?? are designed to capture known systematic exposures, but they leave open the question of whether a broader set of risk factors can better explain strategy excess returns and absorb the residual alpha. Starting from the candidate universe of 75 tradable

risk factors described in Section ??, we proceed in two steps. Section ?? extracts a small number of latent factors by principal components, following ??, and tests whether these diffuse combinations span strategy excess returns. Section ?? turns to a supervised alternative—best-subset selection (??)—that picks the individual factors most relevant for each strategy from the same candidate set. As in Section ??, each layer is subjected to the conditional alpha model (Section ??) and to an out-of-sample re-estimation of the hedge, and Section ?? consolidates the alpha estimates across all layers of the paper.

A Candidate Factor Universe

We assemble a candidate set of 75 risk factors from a comprehensive review of the empirical asset pricing, intermediary, and macro-finance literatures. Each factor is selected on the basis of a published theoretical or empirical link to fixed-income arbitrage returns, intermediary balance-sheet constraints, or broad market risk conditions. The factors are organised into six economic categories: credit risk (Panel A), liquidity and funding (Panel B), equity (Panel C), volatility (Panel D), macro-financial (Panel E), and interest rates and term structure (Panel F). The complete list with definitions, data sources, and references is reported in Table ?? (Appendix ??). This candidate universe is the common input to both the PCA of Section ?? and the best-subset selection of Section ??.

Every candidate is tradable: each series is the excess return on an implementable position—an index, a long-short portfolio, a derivative overlay, or a DV01-neutral curve trade. State variables enter through traded counterparts rather than as raw proxies; the intermediary capital factor of ?, for example, enters as the excess equity return on the primary-dealer holding companies (correlation 0.92 with the underlying capital ratio), and aggregate liquidity enters through the traded liquidity portfolio of ?. The restriction matters for interpretation: with every regressor a realised excess return, the intercept α_i in the spanning regressions of this section is a Jensen alpha—the return on an implementable strategy hedged with implementable positions—rather than the residual of a projection on non-traded state variables.

B Principal Component Analysis

When the number of candidate factors is large relative to the sample size, regressing strategy returns on the full set is unreliable: the strong collinearity of the candidates inflates the variance of the coefficient estimates and overfits the sample. The standard response is to summarise the common variation in a small number of principal components and test whether these diffuse combinations span strategy returns; estimating latent factors from a large panel and using them to measure the abnormal return of a managed position is the performance-measurement framework of ??.

We estimate the components from the full estimation sample rather than from a trailing window. For an in-sample spanning test—as distinct from a return *forecast*—using the entire sample to estimate the loadings is both standard and desirable: it delivers the most precise estimate of the common-factor space, and with the cross-sectional dimension N and the time dimension T both large it yields consistent and asymptotically normal second-stage estimates of α_i and β_i (?). We standardise each factor to zero mean and unit variance over the estimation sample and retain the first $K = 8$ components, the number selected by the ? IC_{p2} criterion (the less parsimonious IC_{p1} indicates thirteen). The eight components capture 64.7% of the variance in the balanced factor panel (Figure ??), and robustness of the alpha to $K \in \{5, 8, 11, 13\}$ is reported in Appendix ??.

The spanning regression takes the form

$$r_{i,t} = \alpha_i + \beta_i' \mathbf{PC}_t + \varepsilon_{i,t},$$

where $r_{i,t}$ is the monthly excess return of strategy i and $\mathbf{PC}_t$ is the $K \times 1$ vector of contemporaneous principal component scores. Inference uses Newey–West HAC standard errors (?). Because the regressors are estimated principal components, the second-stage standard errors are in principle subject to a generated-regressor adjustment; under $N, T \rightarrow \infty$ with $\sqrt{T}/N \rightarrow 0$ the adjustment is asymptotically negligible (?), and the out-of-sample design below, which re-estimates the components on past-only data, provides the finite-sample check.

Table ?? reports the estimates. The intercept remains positive and significant at the 1% level for all three strategies: +1.44% p.a. for BTP Italia, +4.78% for CDS–Bond Basis, and +2.21% for iTraxx Skew. The adjusted $\bar{R}^2$ is low throughout, reaching at most 13.4%: the diffuse principal components do not span strategy excess returns. Table ?? reports, for each retained component, the candidate factors with the largest absolute loading. The components admit a clear economic reading: PC1 is dominated by credit and equity-market factors, PC2 by global rates and bond returns, and PC3 by swap spreads and dealer CDS.

Out-of-Sample Alpha. As in Section ??, the full-sample loadings could reflect an in-sample-optimal combination of the factors rather than a return an investor could have earned in real time (??). Panel A of Table ?? addresses this directly. At each evaluation date we re-estimate the principal components *and* the spanning betas on a past-only window—expanding (initialised at 60 months) as the baseline and a 60-month rolling window as a robustness check—and define the out-of-sample alpha as the mean of the hedged return, tested with the ? HAC correction. Because the components are themselves re-estimated within each window, the design is invariant to the sign and rotation indeterminacy of principal components. The out-of-sample alpha remains positive and significant at the 1% level for all three strategies: +1.61% p.a. for BTP Italia, +4.55% for CDS–Bond Basis, and +1.85% for iTraxx Skew under the expanding window, close to the in-sample estimates of Table ??, so the alpha is not an artefact of full-sample factor estimation.

Because PCA is an unsupervised method—it extracts factors that maximise variance in the candidate set, not factors that explain variation in strategy returns—it cannot identify strategy-specific risk drivers. This motivates the supervised selection of the next subsection.

C Best-Subset Selection

Preprocessing. The candidate factors are preprocessed for selection. Six candidates are near-duplicates of another factor in the universe, with absolute correlation above 0.95; one member of each pair is removed, leaving $p = 69$ candidates in the selection design. Each factor is standardised

to unit L_2 -norm, so that which factors enter is invariant to scale; point estimates and inference are subsequently computed on the original units. Even at $p = 69$, the space of possible factor subsets— $2^{69} \approx 6 \times 10^{20}$ —rules out exhaustive search and calls for a selection method with an explicit cardinality control.

Selection. For each strategy we solve the best-subset problem of ?,

$$\min_{\alpha, \beta} \frac{1}{T} \sum_{t=1}^T (r_t - \alpha - \beta' \mathbf{f}_t)^2 \quad \text{subject to} \quad \|\beta\|_0 \leq k, \quad (13)$$

computed with the splicing algorithm of ? polished against a greedy forward search, retaining whichever feasible k -subset attains the lower in-sample residual sum of squares. The ℓ_0 constraint has two properties that penalised methods lack and that matter in our design. First, a factor that enters is not shrunk, so among a cluster of correlated candidates the constraint keeps the most informative member and drains its near-duplicates, instead of splitting a common signal across the cluster or shrinking it toward zero as the LASSO does. Second, the cardinality is controlled directly: we fix $k = 8$ for parsimony and cross-strategy comparability, rather than selecting the count by an information criterion or cross-validation, whose objective is nearly flat—and whose argmin is therefore ill-determined—on alpha-dominated series with weak factor signal. Table ?? (Appendix ??) re-estimates the selection at $k \in \{5, 8, 10\}$: the alpha is flat in k , so the conclusions do not hinge on the choice. All-subsets model comparison has asset-pricing precedent in ?.

Post-Selection Inference. Given the selected set $\hat{S}_i$, we re-estimate the spanning regression by unrestricted OLS on the original (un-standardised) factors,

$$r_{i,t} = \alpha_i + \beta_i' \mathbf{f}_{\hat{S}_i,t} + \varepsilon_{i,t}, \quad (14)$$

with Newey–West HAC standard errors. Because the set is selected on the same data, we complement the analytic t -tests with a stationary block bootstrap (?) of the post-selection regression (9,999

replications, expected block length nine months) and report the percentile 95% confidence interval for alpha. The post-double-selection control set of $\hat{\alpha}$, which guards the alpha against omitted-control bias from the selection step, is reported alongside the selected sets in Appendix ??.

Results. Table ?? reports the post-selection estimates. Alpha is significant at the 1% level for all three strategies: +2.2% p.a. for BTP Italia (bootstrap CI [1.3, 3.0]), +4.8% for CDS–Bond Basis (CI [2.9, 6.6]), and +2.9% for iTraxx Skew (CI [1.5, 3.8]). With eight factors each, the models deliver adjusted $\bar{R}^2$ of 31%, 21%, and 25%—well above any benchmark framework of Section ?? and the PCA above—yet absorb none of the three alphas.

The selected sets differ across strategies in ways that mirror each trade’s structure, and all but two of the twenty-four selected loadings are significant. The BTP Italia set combines defensive and value fixed-income styles with inflation and equity-market exposure: the set is anchored by the fixed-income defensive factor ($-0.21, t = -4.9$), with long-horizon inflation expectations ($-0.23, t = -2.7$) and the equity market ($-0.04, t = -4.0$) carrying the next-strongest loadings, and size, fixed-income value, the euro default premium, traded liquidity, and the five-year swap-spread trade completing the set. The CDS–Bond set maps onto the credit and intermediary channels: the strategy loads negatively on the prime-broker CDS factor ($-0.22, t = -2.3$)—it pays off when dealer credit deteriorates—and positively on the default premium ($+0.11, t = 3.2$) and the industrial bond return ($+0.22, t = 3.1$), the natural exposures of a leveraged long-credit position, with smaller loadings on high yield, emerging-market debt, and the curvature trade. The iTraxx set is anchored by the traded intermediary capital factor of $\hat{\alpha}$ ($-0.03, t = -4.1$)—the strategy pays off when dealer equity underperforms—alongside betting-against-beta ($-0.04, t = -4.3$), global and fixed-income momentum (-0.05 and $-0.04, t = -3.1$ and -3.9), the one-month variance-swap return ($-0.86, t = -2.5$), value, and the CDX investment-grade position. The three selected sets are pairwise disjoint: no factor appears in more than one strategy, so the alpha that survives is not compensation for a single common exposure, and its origin is strategy-specific.

Out-of-Sample Alpha. Panel B of Table ?? applies the recursive design of Section ?? to the selected models: the best-subset factor set is frozen, only the hedge loadings are re-estimated on the past-only window, and the realised hedged return is tested with the HAC correction—the real-time abnormal return of the selected model, with no per-window re-selection. The out-of-sample alpha is positive and significant at the 1% level for all three strategies under both schemes: +2.05% p.a. for BTP Italia, +3.90% for CDS–Bond Basis, and +2.50% for iTraxx Skew on the expanding window, against in-sample estimates of +2.2, +4.8, and +2.9% in Table ??.

Selector Robustness. Table ?? shows that neither the ℓ_0 machinery nor the post-selection inference drives the result. Five alternative selectors span the methodological space: a one-per-risk-category variant of the best subset, the adaptive elastic net of ?, the relaxed lasso of ?, a nonlinear machine-learning screen (random forest permutation importance, ?), and model-X knockoffs (?) with the false discovery rate controlled at 10%. For every selector the table reports the honest split-sample alpha—factors selected on the first half of the sample, alpha estimated on the second—and the frozen-set out-of-sample alpha. All are significant at the 1% level, with the honest estimates spanning +1.5 to +2.0% p.a. for BTP Italia, +2.5 to +3.4% for the CDS–Bond basis, and +2.1 to +2.5% for iTraxx Skew; for the primary selector the out-of-sample column coincides with Panel B of Table ??. The knockoffs row is the sharpest statement of the two-sided message: at a controlled false discovery rate, no individual candidate carries enough signal to be selected at all, and the resulting no-controls alpha is the raw strategy alpha—so the data neither identify a stable factor set nor produce one that absorbs the premium. The identity of the selected factors is unstable across selectors and sample halves, but the exposures that recur are the consensus ones—the fixed-income defensive factor for BTP Italia, emerging-market debt and the industrial bond return for the CDS–Bond basis, the intermediary and defensive-style block for iTraxx Skew—and the double-selection control set is a strict subset of exactly these (Appendix ??). All primary alpha estimates also clear the $t > 3.0$ hurdle that ? propose for factors discovered in multiple-testing environments.

D Conditional Alpha

Table ?? decomposes alpha across the three stress regimes on the iTraxx Main 5-year level—LOW, MEDIUM, and HIGH, with the same 60 and 100 basis-point cutoffs used throughout the paper—under both the PCA and the best-subset controls. In every cell, alpha rises monotonically from the LOW to the HIGH regime. The CDS–Bond basis shows the sharpest gradient—from about +2.8% in LOW to +7.6% in HIGH under either set of controls—and BTP Italia the familiar pattern of an alpha concentrated away from calm markets: under the PCA controls it rises from +0.4% (insignificant) in LOW to +3.3% in HIGH; under the best-subset controls the whole profile shifts up, from +1.3% to +3.8%. For iTraxx Skew the alpha is significant in every regime under both specifications, rising from +1.0% to +3.5% under the PCA controls, with a higher and flatter profile (+2.5% to +3.4%) under the best-subset controls.

The ? conditional model of Section ??—the full specification of equation (??), with the lagged standardised stress state conditioning both the alpha and the loadings, and α_1 tested by the moving-block bootstrap—is re-estimated on each factor layer in Appendices ?? and ?. Under the PCA controls, α_1 is significant at the 1% level for all three strategies. Under the best-subset controls, it remains large and significant for the CDS–Bond basis (+2.55% per standard deviation, $t = 4.31$), while for BTP Italia and iTraxx Skew the continuous slope is absorbed once the loadings themselves are allowed to vary with stress—the regime gradient of Table ?? survives, but it is carried by the conditional loadings rather than by a conditional intercept. Sub-period stability of the alpha estimates is confirmed in Appendix ??, and rolling 36-month alpha estimates with regime shading are plotted in Appendix ??.

E Alpha Across Layers

Table ?? consolidates the answer to the question this section and the previous one pose. Each cell reports the alpha and adjusted $\bar{R}^2$ from one layer of controls: the three benchmark frameworks of Section ??, the principal components, and the best-subset selection. Moving down the table, the explanatory power rises—from $\bar{R}^2$ between 0% and 15% under the benchmarks to 21–31% under

the best-subset models—while the alpha barely moves: +1.4 to +2.2% p.a. for BTP Italia, +4.0 to +4.8% for the CDS–Bond basis, and +1.8 to +2.9% for iTraxx Skew, significant at the 1% level in every cell. A supervised search over 75 tradable candidates, free to pick whichever eight factors best explain each strategy, absorbs no more of the alpha than the fixed benchmarks do.

Two conclusions follow. First, the excess returns documented in Section ?? are not compensation for systematic risk exposures, whether those exposures are specified ex ante, extracted as latent components, or selected by the data. Second, the alpha that survives is state-dependent, rising when intermediary funding stress is elevated, under every layer of controls. The natural question is why such a premium persists—and whether the three mispricings that generate it are connected through the balance sheets of the intermediaries that trade them. Section ?? takes up that question.

IV Cross-Strategy Interdependencies

The preceding sections have established that all three strategies generate alpha that survives both classical benchmark and data-driven factor controls, and that this alpha rises during periods of intermediary funding stress. A natural question follows: are the three mispricings linked to one another, and if so, what economic mechanism ties them together? The slow-moving capital literature introduced above—?, ?, ?, ?, and ?—yields five testable predictions about how arbitrage mispricings behave when they draw on a common pool of slow-moving intermediary capital: (P1) mispricings serviced by a common pool of intermediary capital should co-move (?); (P2) the co-movement should intensify when that capital is scarce (??); (P3) mispricing changes in one strategy should contain information for predicting those of others, with effects propagating asymmetrically along the liquidity gradient (??); (P4) the persistence of mispricing levels should reflect the funding-dependence of each arbitrage trade, with strategies that require continuous balance-sheet funding displaying the longest half-lives (?); and (P5) a common factor extracted from the co-movement should be explained by intermediary balance-sheet variables rather than by macroeconomic fundamentals (??).

A feature of our three-strategy setting gives these tests their identifying power. All three arbitrages are run by the same type of leveraged institutional arbitrageur, but the instruments differ in who holds them: the CDS–Bond and iTraxx legs trade among institutions, whereas the BTP Italia inflation-linked bond is held predominantly by retail investors. This heterogeneity in the holder base is the source of identification we exploit throughout—it lets us separate the slow-moving-capital mechanism, under which co-movement depends on sharing a capital pool, from a common macroeconomic driver, under which all three mispricings would move together regardless of who holds the instruments.

We devote one subsection to each prediction, stating in each case the prediction, the test, and the verdict where the evidence is presented. Section ?? first defines the mispricing series and the stress regimes. Sections ??–?? then test P1 through P5 in turn. Section ?? interprets the evidence as a whole, explains why these arbitrage opportunities persist in post-crisis European credit markets, and collects the verdicts in a summary mapping.

A Mispricing Construction and Stress Regimes

Following ?, we work with mispricing magnitudes $m_{i,t}$ rather than strategy returns, because it is the level of the mispricing—the distance of the basis from its no-arbitrage value—that measures the degree to which capital is failing to close the gap. For each strategy, $m_{i,t}$ is defined as the absolute value of the basis, sign-adjusted so that an increase $\Delta m > 0$ denotes a widening of the mispricing for all three strategies. The BTP Italia mispricing is the cross-sectional mean of the absolute basis across all outstanding issues with at least six months to maturity; the CDS–Bond mispricing is the cross-sectional mean of the absolute value of all negative single-name bases in the iTraxx Main universe on each day; and the iTraxx Skew mispricing is the cross-sectional mean of the absolute skew across the four on-the-run sub-indices. All three series are aggregated to monthly frequency. Market conditions are classified into three regimes—LOW, MEDIUM, and HIGH—based on the iTraxx Europe Main 5-year spread, with thresholds at 60 and 100 basis points. These are the same thresholds that define the stress-dependent transaction-cost tiers applied to every return series in the

paper (Appendix ??) and that demarcate the HIGH stress regime in the benchmark and conditional-alpha analyses of Sections ?? and ??.

B P1: Cross-Market Co-Movement

The first prediction is that mispricings serviced by a common pool of intermediary capital should co-move: when the capital that closes one arbitrage is withdrawn, the same withdrawal leaves the others unhedged, so their mispricings widen together. We test this with pairwise Pearson correlations of mispricing changes Δm , computed unconditionally and within each stress regime (Table ??, Panel B). Panel A shows why the test is run on mispricings rather than on returns: strategy returns, which mix the mispricing with carry and hedging components, display none of the structure below—the institutional-pair return correlation is -0.02 —so the capital channel operates through the mispricing levels.

The prediction holds for the pair whose underlying instruments are all institutionally held and traded—the CDS–Bond basis and the iTraxx skew, hereafter the “institutional pair.” Their mispricing changes correlate at $+0.19$ unconditionally, significant at the 10% level, and the correlation rises monotonically across regimes from -0.21 in the LOW regime to $+0.32$ in the HIGH regime, with the LOW-to-HIGH movement significant at the 1% level under a Fisher z -test. The β correction for the heteroskedastic increase in mispricing volatility under stress attenuates the shift but leaves it significant at the 10% level, so the intensification reflects a genuine increase in the cross-market link rather than the mechanical inflation of correlations that volatility differences alone would produce. When dealer balance sheets are healthy the two markets price almost independently; when balance sheets are depleted, common funding shocks force the simultaneous unwinding of positions in both, and the link turns strongly positive. P1 is confirmed for the institutional pair.

The BTP Italia pairs move in the opposite direction, and it is worth being explicit that this does not contradict P1. The CDS–Bond versus BTP Italia correlation is -0.30 unconditionally, significant at the 1% level, and deepens from -0.13 in LOW to -0.47 in HIGH, the latter significantly

different from zero at the 5% level. The BTP Italia versus iTraxx correlation is negative but indistinguishable from zero on the full sample (-0.10), and reverses sign across regimes, from $+0.15$ in LOW to -0.14 in HIGH. The dynamic conditional correlations estimated by a bivariate GARCH–DCC model (Appendix ??, Figure ??) show that the negative sign of both BTP Italia pairs is a persistent feature of the data—the conditional correlation between CDS–Bond and BTP Italia remains below zero in every month of the sample, with mean -0.29 , and that between BTP Italia and iTraxx is negative throughout, with mean -0.10 —not an artefact of averaging over offsetting positive and negative episodes.

P1 does not predict that all mispricings co-move—only that those drawing on a common capital pool do. The BTP Italia inflation-linked bond is retail-held and does not draw on the institutional pool, so P1 predicts no positive co-movement, and there is none. This is exactly the holder-base heterogeneity that identifies the mechanism: the sign of the co-movement turns on whether the capital pool is shared—positive within the institutional pool, negative against the retail-held basis—rather than being common to all three, as a macroeconomic driver would require. We take up the economics of the negative sign in Section ??.

C P2: Stress Amplification

The second prediction sharpens the first: the co-movement should not merely be present but should intensify as intermediary capital becomes scarce. The regime correlations of the previous subsection already provide one reading—co-movement within the institutional pair rises from -0.21 in LOW to $+0.32$ in HIGH—but they treat stress as a discrete state and cannot say whether transmission rises continuously with it, nor separate the incremental effect of stress from the average level of transmission. We therefore test the prediction parametrically with an interaction regression:

$$\Delta m_{i,t} = \alpha + \beta_0 \Delta m_{j,t} + \beta_1 (\Delta m_{j,t} \times z_t) + \gamma z_t + \varepsilon_t, \quad (15)$$

where z_t is the contemporaneous standardised iTraxx Main 5-year spread, a proxy for intermediary funding stress; the state enters contemporaneously because the object here is amplification, not predetermined conditional performance. The coefficient β_1 captures the incremental transmission during stress: $\beta_1 > 0$ means that a given shock to Δm_j has a larger effect on Δm_i when balance sheets are strained. The theory is explicit on the mechanism: ? notes that a depleted intermediary is “less able or predisposed to absorb supply and demand shocks,” and ? show that the destabilising effects of funding constraints are “largest when traders are near their constraint.”

Table ?? reports the results. Within the institutional pair, the interaction coefficient on iTraxx mispricing carries the predicted positive sign, and the direct effect of stress on the CDS–Bond basis is significant in both interaction regressions, at the 1% level ($\gamma = +1.59$) when BTP Italia mispricing is the explanatory variable and at the 5% level ($\gamma = +0.93$) when iTraxx mispricing is. Three pieces of evidence converge to support P2 within the institutional pair. First, the regime correlation rises monotonically across stress states (Section ??). Second, the Forbes–Rigobon correlation confirms that this rise reflects a genuine increase in the cross-market link rather than a heteroskedasticity artefact. Third, the within-pair interaction term β_1 is correctly signed and, in the CDS–Bond $\rightarrow$ iTraxx direction, significant at the 5% level ($\beta_1 = +0.08$). All three readings point the same way, and P2 is confirmed for the institutional pair. For the pairs involving BTP Italia the interaction coefficient is significantly negative ($\beta_1 = -0.64$ for CDS–Bond $\rightarrow$ BTP Italia, significant at the 1% level): stress widens, rather than amplifies, the link between the institutional pair and the retail-held basis—the mirror image of P2 that the holder-base heterogeneity of Section ?? explains.

D P3: Lead–Lag Transmission

The third prediction concerns timing: if capital moves slowly across markets, a mispricing change in one strategy should help predict future changes in another, and the transmission should run from more liquid to less liquid markets as capital is reallocated across the dealer network. We test

predictability with spanning regressions in the design of ?,

$$\Delta m_{i,t} = \alpha + \sum_{j \neq i} \sum_{l=0}^L \beta_{j,l} \Delta m_{j,t-l} + \varepsilon_t, \quad (16)$$

with $L = 2$ lags and Newey–West HAC standard errors. The objects of interest are the lagged coefficients: a significant loading on $\Delta m_{j,t-1}$ means that last month’s mispricing in strategy j predicts this month’s in strategy i , a genuine lead–lag relation rather than contemporaneous co-movement.

Table ?? reports the results. The significant coefficients trace a liquidity pecking order: one-month-lagged iTraxx changes predict the widening of the CDS–Bond mispricing with coefficient $+0.71$ at the 10% level, while CDS–Bond changes enter the BTP Italia equation contemporaneously with coefficient -0.42 at the 5% level, with the mirror-image loading of BTP Italia changes in the CDS–Bond equation (-0.15) significant at the 5% level as well. The two institutional mispricings thus lead the BTP Italia basis with a delay of zero to one month, the timing dictated by the speed at which dealers can rotate inventory across markets.

The spanning regressions establish predictability but not direction. To fix the direction we embed the three Δm series in a first-order VAR and test for Granger causality. The hierarchy is sharp. The iTraxx index Granger-causes both the CDS–Bond basis ($p = 0.011$) and the BTP Italia basis ($p = 0.007$), and the link survives at two and three lags. Neither direction reverses: BTP Italia does not Granger-cause either institutional mispricing at conventional levels, nor does CDS–Bond Granger-cause iTraxx. The directional link from iTraxx to the CDS–Bond basis strengthens to the 1% level on the longer bivariate sample available for those two series alone. Shocks thus originate in the most liquid, standardised index market, where dealer inventory adjusts fastest, and reach both the less liquid cash basis and the retail-held inflation-linked basis with a delay. The generalised impulse responses of ?, invariant to the ordering of variables in the VAR, are consistent with the same picture (Figure ??). A one-standard-deviation shock to the iTraxx index produces a positive response in the CDS–Bond basis that peaks at one month, and a negative response in the BTP Italia basis that peaks at one month; both decay within two to three months, with confidence

bands that narrow rapidly toward zero. Shocks in the reverse direction—originating in either of the less liquid bases and propagating back to the iTraxx skew—are indistinguishable from zero across the full horizon, with bootstrap bands that contain zero throughout. The asymmetric pattern matches the forced-deleveraging mechanism of ?, by which intermediaries adjust their most liquid positions first when funding contracts and their less liquid positions only with a delay. P3 is confirmed: predictability, direction, and dynamic response all run from the liquid index to the less liquid bases. Robustness to lag length ($p = 1, 2, 3$) appears in Appendix ??.

E P4: Persistence and Funding-Dependence

The fourth prediction links the persistence of each mispricing to how funding-dependent its trade is: the more a strategy relies on continuously financing a cash position on a dealer balance sheet, the longer a shock to its mispricing should take to decay, because the capital that closes it arrives only slowly. ? makes this quantitative—the half-life of the reversal that follows a shock is decreasing in the mean rate at which arbitrage capital reaches the trade—so the cross-strategy ordering of half-lives should track the ordering of funding frictions. We estimate first-order autocorrelations of the mispricing levels,

$$m_{i,t} = c_i + \phi_i m_{i,t-1} + u_{i,t}, \quad (17)$$

with Newey–West standard errors, and read off the implied half-life $h_i^* = \ln(0.5)/\ln(\phi_i)$.

Table ?? reports the estimates. The CDS–Bond basis is the most persistent ($\phi = 0.79$, half-life ≈ 3.0 months), with the BTP Italia and the iTraxx skew essentially identical at $\phi \approx 0.73$ and half-life ≈ 2.2 months, all three coefficients significant at the 1% level. The CDS–Bond basis requires the arbitrageur to hold the cash bond on a repo book and finance it continuously, so a widening persists until fresh balance-sheet capacity arrives—the longest half-life of the three, roughly one-third longer than its counterparts. The two remaining mispricings do not separate: the iTraxx skew is a pure derivative position, but the on-balance-sheet exposure of trading the index against its constituents carries non-trivial repo and counterparty costs that the half-life of the basis itself reflects; the BTP Italia is refreshed by primary auctions reserved to households and stabilised by the loyalty

bonus, both of which slow its mean reversion below what a pure-derivative basis would otherwise deliver. P4 is therefore confirmed at the upper end of the funding-intensity spectrum—the cash-bond basis is the most persistent—but not in the finer ordering below it: the differential between the two less funding-intensive strategies is too small to register in monthly persistence. The skew would, on funding grounds alone, be expected to revert fastest, yet its half-life matches the BTP Italia’s.

The impulse responses corroborate this reading from a second angle. The own-responses on the diagonal of Figure ?? decay within the first month: innovations to Δm dissipate quickly precisely because the level m mean-reverts over a multi-month horizon, at a speed governed by funding-dependence. This is the “sharp reaction . . . and a subsequent and more extended reversal” that ? identifies as the signature of slow-moving capital: the impact is absorbed within a month, while the level reverts over the horizon set by the rate of capital arrival.

F P5: The Common Factor

The fifth prediction concerns the source of the co-movement. If slow-moving intermediary capital is the binding constraint, the common factor behind the three mispricings should load on intermediary balance-sheet variables and not on macroeconomic fundamentals; ? identifies “the level of capital commonly available to bear losses” as the link tying mispricings across markets, and ? tie the commonality of liquidity to the speculators’ shadow cost of capital. We extract the first principal component (PC1) of the three Δm series and regress it on two sets of variables (Table ??): intermediary stress proxies, and a placebo set of macroeconomic indicators that should not explain PC1 if the slow-moving capital interpretation is correct.

The contrast is sharp. In the intermediary panel, one proxy per channel of the literature: a fall in the ? capital ratio (−4.28, significant at the 5% level), a rise in the Libor–OIS funding spread (+3.35, significant at the 1% level), and a rise in the ? illiquidity measure (+1.77, significant at the 1% level) all predict wider mispricings, each with the sign the theory assigns to its channel. The fourth proxy, the excess return on protection sold on the European prime brokers, enters pos-

itively and significantly at the 5% level: when dealer credit deteriorates, the mispricings widen. The model explains 23.3% of the variation in the common factor. The placebo panel reverses the picture: regressed on macroeconomic and real-activity uncertainty, long-term inflation expectations, and economic policy uncertainty, the common factor yields four uniformly insignificant coefficients and an adjusted R^2 that falls by more than half, to 10.6%. A factor loading on macroeconomic fundamentals would have reversed this ranking. P5 is confirmed: the force driving the three mispricings to co-move is the state of intermediary balance sheets, not of the macroeconomy.

G Discussion: Why These Arbitrages Persist

Holder-Base Heterogeneity as the Source of Identification. The evidence of the preceding subsections is unified by the single structural feature introduced at the outset: although all three strategies are run by the same type of institutional arbitrageur, the instruments differ in who holds them, and it is this heterogeneity that identifies the mechanism. The CDS–Bond basis and the iTraxx skew trade in instruments whose holder base is overwhelmingly institutional—hedge funds, dealer proprietary desks, and leveraged accounts—so that a spike in volatility inflates the Value-at-Risk of the combined position, exhausts the desk’s risk budget, and forces simultaneous liquidation across both books (?), pushing both bases wider at once. The BTP Italia is different: only its inflation-linked leg is structurally retail-dominated. Introduced in 2012 with a primary auction reserved to households and a loyalty bonus (*premio fedeltà*) for buy-and-hold investors, and serving as the simplest inflation hedge available to an Italian household, the bond is held by preferred-habitat investors in the sense of ?: they use no leverage, face no margin calls, and do not unwind when dealer capital is scarce. The arbitrageur of the BTP Italia basis is itself institutional and faces the same funding environment as the others—so the strategy is fully part of the experiment, not a bystander to it—but the absence of forced selling on the bondholder side keeps the basis out of the institutional fire sales, so that stress drives it opposite to the institutional pair rather than alongside it. This is what gives the design its identifying power: a common macroeconomic factor would move all three mispricings in the same direction regardless of who holds the instruments,

whereas the sign of the co-movement turns entirely on whether the holder base is shared—the pattern documented under P1 and P2.

Why the Mispricings Persist. Why have these opportunities not been competed away? The answer lies in structural limits to arbitrage that differ across strategies but share a common root in the scarcity of intermediary capital. For the BTP Italia, the binding constraint is the segmentation itself: institutional arbitrageurs can exploit the basis, but the retail holders who form the natural counterparty neither monitor nor arbitrage it (?). For the CDS–Bond basis, the limits are those of ? and ?: the trade must fund the corporate bond on a dealer balance sheet, the repo market for European investment-grade corporates is thin and subject to recall risk, and the counterparty risk on the CDS leg consumes regulatory capital—so capacity is procyclical and the mispricing countercyclical, widening precisely when the trade is most attractive but least feasible. For the iTraxx skew, the constraint is post-crisis regulation: ? show that the supplementary leverage ratio and other Basel III requirements raised the capital charge on index-versus-constituent trades—as required leverage ratios rose from roughly 2–3% pre-crisis to 5–6% post-crisis, the return on equity fell by a factor of two to three—rendering them “significantly less economical” and producing “potentially ‘new normal’ levels for both the CDS-bond and the CDX-CDS bases.” The persistence of alpha documented in Sections ?? and ?? is therefore not a puzzle calling for a risk-based explanation; it is the equilibrium of a market in which the cost of intermediary capital has been permanently raised, and slow-moving capital, in the sense of ?, keeps mispricings open longer than frictionless models would predict (?).

Summary: The Five Predictions.

P1. Mispricings drawing on a common capital pool co-move. The institutional pair correlates at +0.19, while the retail-held BTP Italia, outside that pool, correlates significantly negatively with the CDS–Bond basis at –0.30 unconditionally (Table ??). The sign turns on the holder base, exactly as the common-capital mechanism requires. Confirmed.

- P2.** The co-movement intensifies in stress. Institutional co-movement rises monotonically from -0.21 in LOW to $+0.32$ in HIGH, a LOW-to-HIGH increase significant at the 1% level under a Fisher z -test (Tables ??, ??). Confirmed.
- P3.** Mispricing changes transmit along the liquidity gradient. The iTraxx index Granger-causes both other mispricings ($p = 0.007$ and 0.011), with impulse responses tracing the same liquid-to-illiquid path (Table ??, Figure ??). Confirmed.
- P4.** Persistence reflects funding-dependence. The CDS–Bond basis, the only strategy requiring continuous cash-bond financing on a dealer balance sheet, has a half-life about one-third longer than the iTraxx skew and the BTP Italia (Table ??); the two less funding-intensive strategies are essentially identical. Partially confirmed.
- P5.** The common factor is intermediary capital, not macro fundamentals. PC1 loads on the balance-sheet proxies with the predicted signs ($\bar{R}^2 = 23.3\%$) but not on a macroeconomic placebo set, where every coefficient is insignificant and $\bar{R}^2$ falls to 10.6% (Table ??). Confirmed.

Four of the five predictions find unambiguous support in the data; only P4—the cross-strategy ordering of persistences—is partially supported, holding at the upper end of the funding-intensity spectrum but not within the institutional pair. Taken together, the evidence locates the geography of slow-moving capital not in asset-class labels but in the network of intermediaries that service these markets—and ultimately in who holds the bonds.

V Conclusion

We document positive and economically significant alpha in three European fixed-income arbitrage strategies—the BTP Italia inflation-linked basis, the corporate CDS–bond basis, and the iTraxx index skew—and show that this alpha is not compensation for the identified systematic risk exposures, whether the controls are entered as classical benchmark factors or as data-driven selections from a universe of 75 tradable candidates. It survives out-of-sample re-estimation from past-only

loadings and is invariant to the selection method: a formal false-discovery control selects no factor at all. The mispricings are linked by a common factor that loads on intermediary balance-sheet variables (and not on macroeconomic fundamentals), and their persistence reflects the funding-dependence of the arbitrage trade: the CDS–Bond basis, the only strategy requiring continuous repo financing of the cash bond, is the most persistent, while the BTP Italia and the iTraxx skew mean-revert more rapidly. The three correlations exhibit an asymmetric pattern that mirrors the structural divide between institutional and retail holder bases: positive and rising in stress within the institutional pair (CDS–Bond and iTraxx), and negative across holder bases (BTP Italia versus either institutional strategy). The slow-moving capital literature (???), combined with the post-crisis regulatory environment documented by ?, provides an economic explanation for why these mispricings persist.

The findings have implications for two strands of the literature. For the limits-to-arbitrage literature, the European setting provides a natural test of the slow-moving capital mechanism in a market with a structurally heterogeneous holder base, and the investor-base-segmentation pattern in cross-strategy correlations offers independent evidence that the geography of arbitrage capital—rather than asset-class labels—determines cross-market co-movement. For the fixed-income asset pricing literature, supervised selection on a large candidate universe recovers strategy-specific risk drivers that the standard benchmark factor sets miss—including a direct intermediary-capital exposure for the iTraxx skew—yet economically significant alpha persists for all three strategies, indicating that the residual return is genuinely unspanned rather than compensation for an overlooked systematic factor.

A Appendix

A.1 BTP Italia: Indexation Coefficient and Deflation Floors

The BTP Italia indexes both coupons and principal to the Italian consumer price index for blue- and white-collar worker households, excluding tobacco (FOI ex tabacchi) published by Istat. At each coupon date t , the indexation coefficient CI_t is the ratio of the relevant FOI reference index to its value at the previous coupon date $t - 6m$, so that CI_t measures the cumulative inflation over the six-month coupon period. The semiannual coupon and the capital revaluation payment are computed as

$$C_t = \frac{1}{2} c N \tilde{CI}_t, \quad R_t = N (\tilde{CI}_t - 1), \quad (18)$$

where c is the fixed real annual coupon rate set at issuance, N is the subscribed nominal capital, and $\tilde{CI}_t$ is the *modified* indexation coefficient that incorporates a semiannual deflation floor:

$$\tilde{CI}_t = \max(CI_t, 1). \quad (19)$$

If the FOI ex tabacchi falls between consecutive coupon dates, the indexation coefficient is reset to one for that semester, the coupon reverts to the fixed real rate $c/2 \cdot N$, and no capital revaluation is paid; in the following semester, the base index is reset to the higher of the previous semester's level and the highest reading observed in earlier semesters. The principal is redeemed at par at maturity (N), regardless of the cumulative inflation path. Two features of this design are relevant for the arbitrage. First, the semiannual reset of $\tilde{CI}_t$ to unity at every coupon date means that the notional exposed to default risk is always the original face value N , eliminating the path-dependent contamination that affects conventional inflation-linked bonds such as the U.S. TIPS or the euro-area BTP€i. Second, the deflation floor on the coupon guarantees a strictly positive nominal payment in every period, which is absent from instruments with floor only at maturity.

A.2 Trade Construction Details

A.2.1 Entry/Exit Rules and Thresholds

Table ?? summarises entry and exit thresholds for each strategy. All strategies require a minimum of six months to maturity at entry and maintain at least one open trade after the first signal. The CDS–Bond universe is restricted to single-name CDS whose reference entity belongs to the current on-the-run iTraxx Main series; names that exit the index are no longer traded. All CDS are centrally cleared. The iTraxx skew strategy trades only on-the-run series across Main, SnrFin, SubFin, and Xover.

Table A.I
Entry and Exit Thresholds by Strategy

All thresholds in bps. CDS–Bond Basis is a one-directional negative basis trade (long bond, buy CDS protection). Minimum 6 months to maturity at entry for all strategies. Trades shorter than 3 trading days excluded ex-post.

Strategy	Entry Long	Entry Short	Exit Long	Exit Short
BTP Italia	> 40	< –50	< 10	> –10
CDS–Bond Basis	< –40	–	> 0	–
iTraxx Main	> 5	< –10	< –5	> 10
iTraxx SnrFin	> 8	< –15	< –5	> 5
iTraxx SubFin	> 10	< –20	< –10	> 10
iTraxx Xover	> 15	< –25	< –15	> 10

A.2.2 Trade-Level Statistics

Table ?? reports trade-level statistics for each strategy, including the number of trades, average duration, average profit per trade, and the distribution of exit reasons.

Table A.II
Trade-Level Statistics by Strategy

This table reports trade-level statistics for each arbitrage strategy. P&L is the cumulative return per trade in percentage points. Duration is in calendar days. Exit reasons may not sum to 100% due to rounding.

	BTP-Italia	iTraxx <i>skew</i>	CDS-Bond Basis
Total trades	260	206	931
Long	224	146	931
Short	36	60	0
Avg duration (days)	386	225	240
Median duration (days)	212	97	119
Avg P&L per trade (%)	2.079	1.491	2.463
Median P&L per trade (%)	1.960	1.395	1.927
<i>Exit reasons (%)</i>			
Target hit	92.3	97.1	90.7
Maturity	6.9	0.5	1.0
End of sample	0.8	2.4	8.4

A.2.3 Transaction Costs

Transaction costs are computed at the individual trade level and applied as entry and exit fees proportional to DV01 (bonds) or duration (CDS). Fee levels are regime-dependent: at each trade date, the prevailing level of the iTraxx Main 5Y index determines which of three tiers applies (Table ??). Thresholds are set at 60 and 100 bps, reflecting the empirical widening of bid-offer spreads in stressed markets. For trades held to maturity, the exit fee is waived.

Table A.III
Transaction Fee Schedule (bps per unit DV01 or duration)

The fee is applied as entry $DV01 \times \text{fee (bps)}$ at entry plus exit $DV01 \times \text{fee (bps)}$ at exit; the exit fee is waived for trades held to maturity. The regime is determined by the iTraxx Main 5Y level at the trade date: LOW (<60 bps), MEDIUM (60–100 bps), HIGH (≥ 100 bps).

	LOW (Main < 60)	MEDIUM (60–100)	HIGH (Main > 100)
BTP Italia	3.0	5.0	8.0
CDS–Bond Basis	4.0	7.0	10.0
iTraxx Main	0.5	0.75	1.0
iTraxx SnrFin	0.5	0.75	1.0
iTraxx SubFin	1.0	2.0	3.0
iTraxx Xover	3.0	5.0	7.0

A.3 Signal-Weighted vs Equal-Weighted Comparison

Table ?? compares the signal-weighted (SW, ?) and equal-weighted (EW, ?) indices. SW is adopted as the baseline in the main text; EW results confirm robustness to the weighting scheme. The equal-weighted counterpart of the manipulation-proof performance measure (Table ??) is reported in Table ??; the volatility-managed results for both weighting schemes are in Table ??.

Table A.IV
Equal-Weighted Strategy Returns

This table reports monthly percentage excess returns, net of transaction costs, for the equal-weighted indices, in which each active trade receives weight $1/K_t$; it is the equal-weighted counterpart of Panel B of Table ??, whose signal-weighted indices are the baseline of the paper. The mean and standard deviation are annualized and the Sharpe ratio is their ratio; kurtosis is excess kurtosis; the t -statistic for the mean is corrected for serial correlation (Newey–West); and %Neg is the proportion of negative monthly excess returns.

Strategy	n	Mean (% p.a.)	t -Stat	Std Dev (% p.a.)	Min (%)	Max (%)	Skew	Kurt	%Neg	AC(1)	Sharpe
BTP Italia	163	1.34	3.01	1.69	-1.25	1.92	0.82	1.88	39.9	0.06	0.79
CDS–Bond Basis	207	3.84	5.71	2.62	-3.37	3.54	0.49	5.28	28.0	-0.08	1.46
iTraxx Skew	243	1.37	5.30	1.05	-1.44	1.11	0.26	3.68	39.9	-0.04	1.30

Table A.V
Manipulation-Proof Performance Measure by Risk Aversion (Equal-Weighted)

MPPM is the manipulation-proof performance measure of Goetzmann et al. (2007), computed from monthly returns; ρ is the coefficient of relative risk aversion and the risk-free rate is 1-month Euribor. Returns are net of transaction costs.

Strategy	$\rho = 2$ (% p.a.)	$\rho = 3$ (% p.a.)	$\rho = 4$ (% p.a.)
BTP Italia	1.24	1.22	1.21
CDS–Bond Basis	3.66	3.63	3.60
iTraxx Skew	1.35	1.34	1.34

Table A.VI
Moreira–Muir Volatility-Managed Portfolio Performance

Following Moreira and Muir (2017). The volatility-managed strategy scales the position inversely to the previous month's realized variance, $f_{t+1}^\sigma = (c/\sigma_t^2) f_{t+1}$, with c chosen to match buy-and-hold volatility. Alpha is the intercept from regressing volatility-managed on buy-and-hold monthly returns, with Newey–West standard errors (4 lags). Panel A uses signal-weighted indices; Panel B, equal-weighted. Returns are net of transaction costs. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Strategy	Buy-and-Hold			Vol-Managed			Alpha (% p.a.)	$t(\alpha)$
	Ret	Vol	Sharpe	Ret	Vol	Sharpe		
	(% p.a.)	(% p.a.)		(% p.a.)	(% p.a.)			
<i>Panel A: Signal-Weighted</i>								
BTP Italia	1.52	1.60	0.95	0.37	1.60	0.23	-0.22	-0.50
CDS–Bond Basis	4.52	2.88	1.57	3.08	2.88	1.07	1.39*	1.90
iTraxx Skew	1.91	1.53	1.25	0.84	1.53	0.55	-0.16	-0.58
<i>Panel B: Equal-Weighted</i>								
BTP Italia	1.27	1.67	0.76	0.21	1.67	0.13	-0.33	-0.73
CDS–Bond Basis	3.73	2.53	1.48	2.45	2.53	0.97	0.93	1.38
iTraxx Skew	1.36	1.05	1.29	0.60	1.05	0.57	-0.01	-0.09

A.4 Benchmark Regressions with US Factors

For completeness, the three benchmark regressions are re-estimated using the original US factor proxies; the euro-area results in the main text (Tables ??, ??, and ??) remain the primary specification.

Table A.VII
Duarte Factor Model Regressions (US Factors)

This table reports estimates of the regression $r_{i,t} = \alpha_i + \beta_1 Mkt_t + \beta_2 SMB_t + \beta_3 HML_t + \beta_4 UMD_t + \beta_5 R_{S,t} + \beta_6 R_{I,t} + \beta_7 R_{B,t} + \beta_8 R_{2,t} + \beta_9 R_{5,t} + \beta_{10} R_{10,t} + \varepsilon_{i,t}$, where $r_{i,t}$ is the monthly excess return of strategy i and the regressors are the ten Duarte et al. (2007) US equity and bond risk factors. α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC). Mkt , SMB , HML , UMD : the Fama–French market, size, value, and momentum factors. R_S : U.S. bank stock excess return. R_I , R_B : A/BAA-rated U.S. industrial and bank bond excess returns. R_2 , R_5 , R_{10} : maturity-sorted U.S. Treasury portfolio excess returns. $\bar{R}^2$ is the adjusted R^2 . Sample: BTP Italia 163 months, CDS–Bond 207 months, iTraxx 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	α (% p.a.)	β_{Mkt}	β_{SMB}	β_{HML}	β_{UMD}	β_{R_S}	β_{R_I}	β_{R_B}	β_{R_2}	β_{R_5}	$\beta_{R_{10}}$	$\bar{R}^2$
<i>BTP Italia</i>	2.02*** (4.17)	−0.05*** (−3.49)	−0.02* (−1.72)	−0.02 (−1.60)	0.01 (0.44)	0.03*** (3.19)	0.03 (0.37)	−0.09 (−1.19)	−0.02 (−0.07)	−0.23 (−0.62)	0.21 (0.97)	0.11
<i>CDS–Bond</i>	4.58*** (5.87)	−0.03 (−0.90)	0.01 (0.45)	−0.00 (−0.11)	−0.01 (−1.01)	0.00 (0.16)	0.25*** (3.35)	−0.06 (−1.63)	0.19 (0.40)	−0.70* (−1.89)	0.26 (1.34)	0.18
<i>iTraxx</i> <i>Skew</i>	1.86*** (4.71)	0.00 (0.31)	−0.02* (−1.68)	−0.02 (−1.38)	−0.01* (−1.92)	−0.00 (−0.09)	0.02 (0.50)	−0.02 (−1.04)	0.22 (0.74)	−0.07 (−0.31)	0.01 (0.10)	0.02

Table A.VIII
Fung & Hsieh Factor Model Regressions (US Factors)

This table reports estimates of the regression $r_{i,t} = \alpha_i + \beta_1 S\&P_t + \beta_2 SC-LC_t + \beta_3 BdOpt_t + \beta_4 FXOpt_t + \beta_5 ComOpt_t + \beta_6 10Y_t + \beta_7 CredSpr_t + \varepsilon_{i,t}$, where $r_{i,t}$ is the monthly excess return of strategy i and the regressors are the seven Fung & Hsieh (2004) US hedge fund risk factors. α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC). S&P: S&P 500 excess return. SC–LC: Russell 2000 minus S&P 500 (small-minus-large) equity spread. BdOpt, FXOpt, ComOpt: trend-following returns from lookback straddles on government-bond, currency, and commodity futures. 10Y: excess return on a 10-year U.S. Treasury total-return index. CredSpr: return spread between BAA U.S. corporate bonds and 10-year U.S. Treasuries. $\bar{R}^2$ is the adjusted R^2 . Sample: BTP Italia 163 months, CDS–Bond 207 months, iTraxx 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	α (% p.a.)	$\beta_{S\&P}$	β_{SC-LC}	β_{BdOpt}	β_{FXOpt}	β_{ComOpt}	β_{10Y}	$\beta_{CredSpr}$	$\bar{R}^2$
<i>BTP Italia</i>	1.84*** (3.82)	−0.02* (−1.94)	−0.02* (−1.89)	0.00 (0.90)	−0.00 (−0.00)	0.00 (0.86)	−0.02 (−0.50)	−0.04 (−1.05)	0.10
<i>CDS–Bond</i>	4.14*** (5.05)	−0.01 (−0.64)	0.01 (0.46)	0.00 (1.17)	−0.01** (−1.97)	0.00 (0.08)	0.13*** (3.33)	0.14*** (3.40)	0.11
<i>iTraxx</i> <i>Skew</i>	1.75*** (4.44)	0.01 (0.65)	−0.02** (−2.02)	0.00 (0.39)	0.00 (1.03)	0.00 (0.52)	0.03 (1.11)	0.01 (0.51)	0.02

Table A.IX
Active Fixed Income Factor Model Regressions (US Factors)

This table reports estimates of the regression $r_{i,t} = \alpha_i + \beta_1 \text{Term}_t + \beta_2 \text{Global Term}_t + \beta_3 \text{Global Agg}_t + \beta_4 \text{Infl. Linkers}_t + \beta_5 \text{Corp Credit}_t + \beta_6 \text{Emerg Debt}_t + \beta_7 \text{Emerg Curr}_t + \beta_8 \text{UST Vol}_t + \varepsilon_{i,t}$, where $r_{i,t}$ is the monthly excess return of strategy i and the regressors are the eight Active Fixed Income risk premia of Brooks et al. (2020). α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC). Term: U.S. Treasury excess return over cash. Global Term: Bloomberg Global Treasury Index (hedged), in excess of cash. Global Agg: Bloomberg Global Aggregate Bond Index, in excess of cash. Infl. Linkers: global inflation-linked bonds, in excess of cash. Corp Credit: U.S. high-yield corporate bond index, in excess of cash. Emerg Debt: excess return on hard-currency emerging-market debt. Emerg Curr: excess return on a basket of emerging-market currencies against the U.S. dollar. UST Vol: excess return on a one-month variance swap on the ten-year U.S. Treasury future, in percentage points of annualised variance. $\bar{R}^2$ is the adjusted R^2 . Sample: BTP Italia 163 months, CDS–Bond 207 months, iTraxx 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	α (% p.a.)	β_{Term}	β_{GlobTerm}	β_{GlobAgg}	$\beta_{\text{InflLinkers}}$	$\beta_{\text{CorpCredit}}$	$\beta_{\text{EmergDebt}}$	$\beta_{\text{EmergCurr}}$	$\beta_{\text{UST Vol}}$	$\bar{R}^2$
<i>BTP Italia</i>	1.99*** (4.56)	0.25** (2.50)	0.09 (0.28)	-0.35 (-0.90)	-0.06* (-1.85)	-0.04 (-0.79)	0.06 (1.25)	-0.01 (-0.35)	-0.11 (-0.57)	0.11
<i>CDS–Bond</i>	3.98*** (5.49)	-0.20 (-1.22)	-0.15 (-0.25)	0.47 (0.66)	-0.05 (-0.60)	-0.04 (-0.80)	0.13*** (3.32)	-0.01 (-0.40)	0.28 (0.97)	0.11
<i>iTraxx Skew</i>	1.86*** (4.26)	0.07 (0.87)	-0.22 (-0.90)	0.26 (1.01)	-0.02 (-0.48)	0.06** (2.04)	-0.08** (-2.17)	0.01 (0.52)	0.34* (1.90)	0.06

A.5 Subperiod and Rolling Window Analysis

This appendix reports sub-period splits, stress-regime decompositions, and out-of-sample robustness for the benchmark factor models of Section ?? . Each table contains Panel A (full sample, first half, second half) and Panel B (alpha by LOW/MEDIUM/HIGH stress regime on the iTraxx Main 5Y level, thresholds at 60 and 100 bps).

A.5.1 Duarte et al. (2007)

Table ?? reports sub-period splits and stress regime alpha for the ? specification.

Table A.X
Sub-Period and Regime Analysis: Duarte et al. (2007)

Annualized alpha (% p.a.) of the three strategies under the Duarte et al. (2007) specification, EUR factors, monthly frequency. Panel A: OLS with Newey–West HAC standard errors (4 lags) on time sub-samples. Panel B: regime alphas from common-beta dummy intercepts $r_t = \sum_g \alpha_g D_{g,t} + \sum_j \beta_j X_{j,t} + \varepsilon_t$, with $g \in \{\text{LOW}, \text{MEDIUM}, \text{HIGH}\}$ defined on the iTraxx Main 5Y level (LOW < 60 bps, MEDIUM [60, 100), HIGH ≥ 100). N = months per sub-sample or regime. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

<i>Panel A: Sub-period alpha (% p.a.)</i>									
Period	BTP Italia			iTraxx Skew			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
Full Sample	+1.62***	3.29	163	+2.19***	4.90	243	+4.24***	5.27	207
First Half	+1.11*	1.96	81	+2.10***	4.56	121	+5.76***	4.39	103
Second Half	+1.79**	2.33	82	+2.60***	3.41	122	+2.50***	2.84	104

<i>Panel B: Alpha by stress regime (% p.a.)</i>									
Regime	BTP Italia			iTraxx Skew			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
LOW	+0.78	1.24	62	+0.97**	2.22	91	+2.15**	2.23	63
MEDIUM	+1.97***	3.15	80	+2.69***	3.65	99	+4.01***	3.53	94
HIGH	+3.05*	1.70	21	+3.27***	4.89	53	+7.11***	4.50	50

A.5.2 Brooks et al. (2020) — Active Fixed Income

Table ?? reports the same analysis under the ? specification.

Table A.XI
Sub-Period and Regime Analysis: Brooks et al. (2020)

Annualized alpha (% p.a.) of the three strategies under the Active FI (Brooks et al. (2020)) specification, EUR factors, monthly frequency. Panel A: OLS with Newey–West HAC standard errors (4 lags) on time sub-samples. Panel B: regime alphas from common-beta dummy intercepts $r_t = \sum_g \alpha_g D_{g,t} + \sum_j \beta_j X_{jt} + \varepsilon_t$, with $g \in \{\text{LOW, MEDIUM, HIGH}\}$ defined on the iTraxx Main 5Y level (LOW < 60 bps, MEDIUM [60, 100), HIGH $\geq$ 100). N = months per sub-sample or regime. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

<i>Panel A: Sub-period alpha (% p.a.)</i>									
Period	BTP Italia			iTraxx Skew			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
Full Sample	+2.12***	3.84	163	+1.75***	4.24	243	+4.20***	5.20	207
First Half	+1.53*	1.83	81	+1.67***	3.86	121	+6.52***	4.83	103
Second Half	+2.28***	3.13	82	+1.73**	2.49	122	+1.48***	2.83	104

<i>Panel B: Alpha by stress regime (% p.a.)</i>									
Regime	BTP Italia			iTraxx Skew			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
LOW	+1.12**	2.13	62	+0.68*	1.94	91	+2.45***	2.86	63
MEDIUM	+2.50***	3.86	80	+2.03***	2.96	99	+3.82***	3.53	94
HIGH	+4.16**	1.98	21	+3.25***	3.92	53	+7.51***	4.51	50

A.5.3 Fung & Hsieh (2004)

Table ?? reports results for the ? specification.

Table A.XII
Sub-Period and Regime Analysis: Fung & Hsieh (2004)

Annualized alpha (% p.a.) of the three strategies under the Fung & Hsieh (2004) specification, EUR factors, monthly frequency. Panel A: OLS with Newey–West HAC standard errors (4 lags) on time sub-samples. Panel B: regime alphas from common-beta dummy intercepts $r_t = \sum_g \alpha_g D_{g,t} + \sum_j \beta_j X_{jt} + \varepsilon_t$, with $g \in \{\text{LOW, MEDIUM, HIGH}\}$ defined on the iTraxx Main 5Y level (LOW < 60 bps, MEDIUM [60, 100), HIGH ≥ 100). N = months per sub-sample or regime. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

<i>Panel A: Sub-period alpha (% p.a.)</i>									
Period	BTP Italia			iTraxx Skew			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
Full Sample	+1.81***	4.08	163	+1.84***	4.49	243	+3.98***	5.05	207
First Half	+1.56**	2.04	81	+1.96***	4.40	121	+5.62***	4.34	103
Second Half	+1.66***	3.00	82	+1.73***	3.12	122	+2.08***	2.68	104

<i>Panel B: Alpha by stress regime (% p.a.)</i>									
Regime	BTP Italia			iTraxx Skew			CDS–Bond Basis		
	α	t	N	α	t	N	α	t	N
LOW	+0.95	1.63	62	+0.50	1.44	91	+2.27**	2.39	63
MEDIUM	+2.05***	3.81	80	+2.23***	3.16	99	+3.65***	3.19	94
HIGH	+3.33*	1.74	21	+3.34***	4.21	53	+6.70***	4.12	50

A.5.4 Out-of-Sample Alpha: Rolling-Window Robustness

Table ?? re-estimates the out-of-sample alpha of Panel C of Table ?? with the loadings fitted on a rolling 60-month window, which relaxes the constant-loadings assumption of the expanding scheme. The layout mirrors Panel C, so the two schemes can be compared cell by cell; the expanding estimates are not reproduced here.

Table A.XIII
Out-of-Sample Alpha: Rolling-Window Robustness

Out-of-sample alpha of the three benchmark hedges (EUR factors) with the loadings re-estimated on a rolling 60-month past-only window, which relaxes the constant-loadings assumption of the expanding scheme of Panel C of Table ??; the layout mirrors Panel C. At each month t , $\hat{\beta}_{i,t-1}$ is estimated by OLS of the strategy excess return $r_{i,t}$ on the benchmark factors $\mathbf{f}_t$ over $[t-60, t)$, and the realized abnormal return is $e_{i,t} = r_{i,t} - \hat{\beta}_{i,t-1}' \mathbf{f}_t$; the in-window intercept (the alpha under measurement) is excluded from the hedge, and rank-deficient or ill-conditioned windows (condition index $> 10^3$) are skipped. α is the annualized mean of $e_{i,t}$ (% p.a.), tested with Newey–West HAC standard errors; IR is the annualized information ratio $\bar{e}_i / \sigma(e_i) \times \sqrt{12}$; N is the number of out-of-sample months, common to the three benchmarks within each strategy. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	N	Duarte et al. (2007)			Brooks et al. (2020)			Fung & Hsieh (2004)		
		α	t	IR	α	t	IR	α	t	IR
<i>BTP Italia</i>	103	+1.24**	2.21	+0.75	+1.90***	3.17	+1.07	+1.56***	2.61	+0.94
<i>iTraxx Skew</i>	183	+2.43***	3.97	+1.18	+1.65**	2.49	+0.87	+2.01***	3.51	+1.06
<i>CDS–Bond Basis</i>	147	+3.35***	3.54	+1.15	+2.88***	2.94	+0.98	+3.11***	3.42	+1.05

A.6 VIF Diagnostics

Table ?? reports Variance Inflation Factors for the EUR factor specifications used in the benchmark regressions of Section ?. VIF is computed on the factor matrix for each strategy’s sample period; because the factors are identical across strategies, differences in VIF reflect only the different estimation windows.

Two groups of regressors are highly collinear. In the Duarte specification, the maturity-sorted bond portfolios and the bond index returns carry elevated VIF ($R5 \approx 52$ – 55 , $R10 \approx 24$ – 27 , $R2$ and $RI \approx 14$ – 16), the expected consequence of the common duration exposure of government and investment-grade bond returns. In the Brooks et al. specification, the two global bond benchmarks are near-duplicates (Global_Term and Global_Aggregate, $VIF \approx 47$ – 66). Collinearity among controls inflates the standard errors of those loadings but leaves inference on the intercept unaffected; all Fung–Hsieh factors have VIF below 4.

Table A.XIV
Variance Inflation Factors — Benchmark Factor Models, EUR Factors

Variance inflation factors of the EUR factors of the three benchmark specifications, computed on each factor matrix over each strategy's estimation sample; conventional caution thresholds are 5 and 10. Elevated VIF for R_2 , R_5 , R_{10} (Panel A) reflects the inherent correlation of maturity-sorted government bond returns; elevated VIF for Global_Term and Global_Aggregate (Panel C) reflects the strong correlation between global bond benchmarks. Neither affects inference on alpha.

Factor	BTP Italia	CDS–Bond Basis	iTraxx Skew
Panel A: Duarte et al. (2007)			
Mkt-RF	4.86	5.27	4.97
SMB	1.11	1.18	1.13
HML	3.15	3.14	3.01
UMD	1.72	1.85	1.81
RS	7.29	8.25	8.10
RI	16.07	5.23	5.08
RB	11.83	2.82	2.77
R2	16.02	14.42	15.93
R5	55.30	51.71	55.21
R10	26.56	24.19	24.01
Panel B: Fung & Hsieh (2004)			
SNPMRF	2.22	2.49	2.22
SCMLC	1.07	1.15	1.10
PTFSBD	1.45	1.40	1.37
PTFSFX	1.46	1.50	1.46
PTFSCOM	1.19	1.27	1.24
R10	3.28	2.60	2.45
BAAMTSY	4.10	4.36	4.04
Panel C: Brooks et al. (2020)			
Term	6.83	6.12	5.57
Global_Term	59.73	46.97	46.73
Global_Aggregate	65.80	50.63	49.48
Inflation_Linkers	5.14	4.19	3.90
Corporate_Credit	3.96	3.78	3.65
Emerging_Debt	8.21	6.10	5.64
Emerging_Currency	2.13	2.23	2.08
UST_Volatility	1.29	1.22	1.19

A.7 Candidate Factor List with Sources

Table ?? reports the full set of candidate risk factors used in the principal-component and best-subset analyses of Section ??, together with their data sources, transformations, and economic category.

Table A.XV
Candidate Risk Factors

Factor ID	Description	Reference
Panel A: Credit Risk Factors		
BTP_BUND	BTP-Bund spread factor (10-year). Duration-matched excess return of Italian over German government bonds, long the Italian and short the German 7-10Y total-return index.	?
BTP_BUND_2Y	BTP-Bund spread factor (2-year). Duration-matched excess return of Italian over German government bonds at the short end, long the Italian and short the German 1-3Y total-return index.	?
CDX_IG	Credit risk factor (U.S. investment grade). Excess return on a fully-collateralised sell-protection position in the 5-year CDX North American Investment Grade index.	??
CREDIT_EU	Bond credit risk factor. Return spread between euro-area BBB and AAA corporate bond total-return indices.	?
CREDIT_US	Bond credit risk factor. Return spread between BAA and AAA U.S. corporate bond total-return indices.	?
CRED_SPR_EU	Bond credit spread factor. Return spread between euro-area BBB corporate bonds and 10-year German Bunds.	?
CRED_SPR_US	Bond credit spread factor. Return spread between BAA U.S. corporate bonds and 10-year U.S. Treasuries.	?
DEF_EU	Bond default risk factor. Total return spread between long-term European corporate bonds and long-term German Bunds.	?
DEF_US	Bond default risk factor. Total return spread between long-term U.S. corporate bonds and long-term U.S. government bonds.	?
EMERG_DEBT	Excess return on hard-currency emerging market debt.	?
EMERG_FX	Return on a basket of emerging market currencies versus the USD (MSCI Emerging Markets Currency Index; currency weights from the MSCI EM Index), in excess of USD cash.	?
GLOBAL_AGG	Excess return on the Bloomberg Global Aggregate Bond Index.	?
HY_CORP	Sub-investment-grade corporate credit factor: excess return of a EUR-hedged Pan-European high-yield corporate bond index over cash.	?
ITRX_MAIN	Credit risk factor (European investment grade). Excess return on a fully-collateralised sell-protection position in the 5-year iTraxx Europe Main index.	??

Note: Factors are organized by category. Each series is reconstructed consistently with the original reference. Total: 75 factors.

Table A.XV (continued)

Factor ID	Description	Reference
PB_CDS_1Y_EU	Intermediary funding factor (Europe, short tenor). Excess return on an equally weighted basket of 1-year sell-protection positions in the senior CDS of five European prime brokers.	??
PB_CDS_1Y_US	Intermediary funding factor (United States, short tenor). Excess return on an equally weighted basket of 1-year sell-protection positions in the senior CDS of four U.S. prime brokers.	??
PB_CDS_5Y_EU	Intermediary funding factor (Europe). Excess return on an equally weighted basket of 5-year sell-protection positions in the senior CDS of five European prime brokers.	??
PB_CDS_5Y_US	Intermediary funding factor (United States). Excess return on an equally weighted basket of 5-year sell-protection positions in the senior CDS of four U.S. prime brokers.	??
RI_EU	Excess return on a European A-rated industrial corporate bond index.	?
SLOPE_3S5S_MAIN	Credit-curve slope factor. Excess return on a DV01-matched position long protection at the 5-year and short protection at the 3-year point of the iTraxx Europe Main curve.	??
SNRFIN_MAIN	Financial-sector stress factor. Excess return on a DV01-matched position long credit via the iTraxx Senior Financials index and short the iTraxx Europe Main index.	??
XOVER_MAIN	Credit-cycle factor. Excess return on a beta-neutral position long credit via the 5-year iTraxx Crossover index and short the 5-year iTraxx Europe Main index, the short leg scaled by a rolling 36-month spread-change beta.	?
Panel B: Liquidity Factors		
HKM_IC	Intermediary capital factor (tradable). Value-weighted equity return on the primary-dealer holding companies, in excess of cash and converted to euro; a tradable proxy (0.92 correlation) for the He-Kelly-Manela capital-ratio innovation.	?
LIQ_V	Value-weighted return on a 10–1 portfolio formed by sorting stocks on historical liquidity betas.	?
Panel C: Equity Factors		
BAB_EU	Betting Against Beta (Europe). Long low-beta and short high-beta European equities, each leg rescaled to unit beta; USD factor converted to euro.	?
CMA_EU	Investment factor (Conservative Minus Aggressive). Return spread between European conservative and aggressive investment portfolios.	?
COM_CARRY	Commodity carry factor. Self-financing long-short return on commodity futures sorted on the slope of the futures curve (carry); converted to euro.	?
COM_MOM	Commodity momentum factor. Self-financing long-short return on commodity futures sorted on their past 12-month return; converted to euro.	?

Note: Factors are organized by category. Each series is reconstructed consistently with the original reference. Total: 75 factors.

Table A.XV (continued)

Factor ID	Description	Reference
FX_CARRY	Currency carry factor. Self-financing long-short return, long high-interest-rate and short low-interest-rate developed-market (G10) currencies; converted to euro.	?
FX_MOM	Currency momentum factor. Self-financing long-short return on developed-market (G10) currencies sorted on their past 12-month return; converted to euro.	?
GMOM	Global all-asset momentum factor. Long-short momentum return averaged across the eight Value-and-Momentum-Everywhere strategies (stock selection in the U.S., U.K., Europe and Japan, plus equity-index, government-bond, currency and commodity allocations), converted to euro.	?
GVAL	Global all-asset value factor. Long-short value return averaged across the eight Value-and-Momentum-Everywhere strategies (stock selection in the U.S., U.K., Europe and Japan, plus equity-index, government-bond, currency and commodity allocations), converted to euro.	?
HML_EU	Value factor (High Minus Low). Return spread between European high and low book-to-market portfolios.	?
MKT_EU	Market excess return on the European value-weighted equity market portfolio over the 1-month German government bill.	?
QMJ_EU	Quality Minus Junk (Europe). Long high-quality and short low-quality (“junk”) European stocks, where quality aggregates profitability, growth, safety, and payout; USD factor converted to euro.	?
RB_EU	Excess return on a European bank bond index.	?
RMW_EU	Profitability factor (Robust Minus Weak). Return spread between European robust and weak operating profitability portfolios.	?
RS_EU	Excess return on the EuroStoxx Banks index.	?
SMB_EU	Size factor (Small Minus Big). Return spread between diversified European small- and large-cap portfolios.	?
UMD_EU	Momentum factor (Up Minus Down). Monthly return on a European zero-investment winners-minus-losers momentum portfolio.	?
Panel D: Volatility Factors		
ATM_IV_CDX	Excess return on a three-month variance swap on the CDX Investment Grade index spread: realized variance of the on-the-run index spread minus its at-the-money (100% moneyness) implied variance fixed at the start of the window (equivalently, the delta-hedged ATM payer/receiver straddle return).	?
ATM_IV_ITRX	Excess return on a three-month variance swap on the iTraxx Europe Main index spread: realized variance of the on-the-run index spread minus its at-the-money (100% moneyness) implied variance fixed at the start of the window (equivalently, the delta-hedged ATM payer/receiver straddle return).	?

Note: Factors are organized by category. Each series is reconstructed consistently with the original reference. Total: 75 factors.

Table A.XV (continued)

Factor ID	Description	Reference
IV_BUND	Excess return on a one-month variance swap on the 10-year German Bund future: realized variance of the future minus its at-the-money implied variance fixed at the start of the month (equivalently, the delta-hedged ATM straddle return).	?
IV_TSY	Excess return on a one-month variance swap on the 10-year U.S. Treasury future: realized variance of the future minus its at-the-money implied variance fixed at the start of the month (equivalently, the delta-hedged ATM straddle return).	?
PTFSBD	Return on the bond trend-following factor, constructed from lookback straddles on government bond futures.	?
PTFSCOM	Return on the commodity trend-following factor, constructed from lookback straddles on commodity futures.	?
PTFSFX	Return on the currency trend-following factor, constructed from lookback straddles on major currency futures.	?
PTFSIR	Return on the short-term interest rate trend-following factor, constructed from lookback straddles on 3-month interest rate futures.	?
PTFSSTK	Return on the equity index trend-following factor, constructed from lookback straddles on stock index futures.	?
SVS_RET_1M	Excess return on a one-month simple variance swap on the S&P 500: realized simple variance minus the option-implied SVIX ² strike fixed at the start of the month.	?
SVS_RET_3M	Excess return on a simple variance swap on the S&P 500 over a trailing three-month window (realized simple variance minus the SVIX ² strike).	?
V2X_FUT	Volatility factor (euro area). Excess return on a constant-maturity (one-month) rolling long position in VSTOXX futures (VSTOXX Short-Term Futures index), in excess of the euro risk-free.	?
VIX_FUT	Volatility factor (U.S.). Excess return on a constant-maturity (one-month) rolling long position in VIX futures (S&P 500 VIX Short-Term Futures Excess-Return index).	?

Panel F: Interest Rate Factors

5Y5Y_INFL	Forward inflation factor (euro area). Mark-to-market excess return on a long-inflation position in a 5y–5y forward zero-coupon inflation swap (annuity-weighted change in the forward breakeven).	?
5Y5Y_INFL_US	Forward inflation factor (United States), expressed in euro. Mark-to-market excess return on a long-inflation position in a 5y–5y forward zero-coupon inflation swap.	?
CURV_2S5S10S_EU	Yield-curve curvature factor. Excess return on a DV01-neutral butterfly in German government bond futures, long the 5-year belly (Bobl) and short the 2-year (Schatz) and 10-year (Bund) wings.	?
CURV_2S5S10S_US	Yield-curve curvature factor. Excess return on a DV01-neutral butterfly in U.S. Treasury futures, long the 5-year belly and short the 2- and 10-year wings.	?

Note: Factors are organized by category. Each series is reconstructed consistently with the original reference. Total: 75 factors.

Table A.XV (continued)

Factor ID	Description	Reference
FI_CARRY	Fixed-income carry factor. Self-financing long-short return on developed-market government bonds sorted on term-structure carry (yield plus roll-down); converted to euro.	?
FI_DEF	Fixed-income defensive factor. Self-financing long-short return on developed-market government bonds tilted toward low-risk exposures; converted to euro.	?
FI_MOM	Fixed-income momentum factor. Self-financing long-short return on developed-market government bonds sorted on their past 12-month return; converted to euro.	?
FI_VAL	Fixed-income value factor. Self-financing long-short return on a cross-section of developed-market government bonds sorted on a value signal; converted to euro.	?
GLOBAL_TERM	Excess return over cash on the Bloomberg Global Treasury Index (EUR-hedged).	?
GOVT_EU	Euro-area term-premium factor. Excess return over cash on the broad German government bond aggregate (all maturities), the euro-area analogue of the US Term factor.	?
INFL_LINK	Excess return on global inflation-linked bonds over cash or nominal Treasuries.	?
R10_EU	Excess return on the 10-year German Bund.	?
R2_EU	Excess return on the 2-year German government bond (Schatz).	?
R5_EU	Excess return on the 5-year German government bond (Bobl).	?
SLOPE_10S30S_EU	Yield-curve slope factor. Excess return on a DV01-matched steepener in German government bond futures, long the 10-year (Bund) and short the 30-year (Buxl).	?
SLOPE_2S10S_EU	Yield-curve slope factor. Excess return on a DV01-matched steepener in German government bond futures, long the 2-year (Schatz) and short the 10-year (Bund).	?
SLOPE_2S10S_US	Yield-curve slope factor. Excess return on a DV01-matched steepener in U.S. Treasury futures, long the 2-year and short the 10-year.	?
SS10Y	Swap-spread factor at the 10-year point. Excess return on a DV01-matched position long the German Bund future and pay-fixed on a 10-year constant-maturity EUR swap.	??
SS2Y	Swap-spread factor at the 2-year point. Excess return on a DV01-matched position long the German Schatz future and pay-fixed on a 2-year constant-maturity EUR swap.	??
SS5Y	Swap-spread factor at the 5-year point. Excess return on a DV01-matched position long the German Bobl future and pay-fixed on a 5-year constant-maturity EUR swap.	??
TERM_EU	Bond term structure factor. Excess total return on long-term German Bunds over the 1-month German government bill return.	?
TERM_US	Bond term structure factor. Excess total return on long-term U.S. government bonds over the 1-month U.S. Treasury bill return.	?

A.8 PCA Robustness

A.8.1 Robustness across the Number of Components

Table ?? reports α and $\bar{R}^2$ across the number of retained components K . Results are robust to the choice of K .

Table A.XVI
PCA Spanning Robustness across the Number of Components

Full-sample contemporaneous spanning $r_{i,t} = \alpha_i + \beta_i' \mathbf{PC}_t + \varepsilon_{i,t}$ for K principal components (baseline $K = 8$). Newey–West HAC t in parentheses. Expl. Var. is the cumulative variance captured by the K components. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

K	BTP Italia		CDS–Bond Basis		iTraxx Skew		Expl. Var.
	α (% p.a.)	$\bar{R}^2$	α (% p.a.)	$\bar{R}^2$	α (% p.a.)	$\bar{R}^2$	(%)
5	1.50*** (3.23)	0.089	4.79*** (6.22)	0.114	2.21*** (5.39)	0.050	55.0
8	1.44*** (3.30)	0.088	4.78*** (6.25)	0.109	2.21*** (5.95)	0.134	64.7
11	1.47*** (3.56)	0.086	4.77*** (6.13)	0.111	2.21*** (5.95)	0.136	71.8
13	1.75*** (4.05)	0.144	4.76*** (6.23)	0.105	2.21*** (6.12)	0.146	75.6

A.8.2 Subperiod and Regime Analysis

Half-sample splits and regime-dependent alpha confirm that PCA-based alpha is stable across subperiods and increases during stress.

Table A.XVII**PCA Subperiod Analysis (Contemporaneous)**

First/Second Half: temporal split at sample midpoint; stress regimes (LOW/MEDIUM/HIGH) are reported in Table ??.
 α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC, 4 lags). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	Full Sample	First Half	Second Half
BTP Italia			
α (% p.a.)	1.44*** (3.30)	1.27** (2.00)	1.38** (2.41)
β_{PC1}	-0.031** (-2.25)	-0.009 (-0.47)	-0.017 (-0.80)
β_{PC2}	-0.025** (-2.02)	-0.003 (-0.17)	-0.040** (-2.51)
β_{PC3}	0.030 (1.03)	-0.003 (-0.07)	0.099* (1.83)
β_{PC4}	0.049 (1.45)	0.114* (1.76)	0.039 (1.10)
β_{PC5}	0.002 (0.04)	0.045 (0.65)	-0.019 (-0.33)
β_{PC6}	-0.025 (-0.62)	-0.033 (-0.76)	0.009 (0.18)
β_{PC7}	-0.038 (-1.47)	-0.051 (-1.14)	0.021 (0.68)
β_{PC8}	-0.022 (-1.13)	-0.023 (-0.40)	-0.031 (-1.23)
$\bar{R}^2$	0.088	0.009	0.147
N	157	78	79
CDS–Bond			
α (% p.a.)	4.78*** (6.25)	6.06*** (4.14)	3.22*** (4.23)
β_{PC1}	0.036* (1.91)	0.005 (0.29)	0.121*** (4.62)
β_{PC2}	0.060*** (2.96)	0.075* (1.70)	0.035* (1.94)
β_{PC3}	-0.086*** (-2.80)	-0.057** (-1.97)	-0.107* (-1.82)
β_{PC4}	-0.014 (-0.59)	0.014 (0.26)	-0.003 (-0.08)
β_{PC5}	-0.017 (-0.55)	-0.011 (-0.28)	-0.048 (-0.81)
β_{PC6}	-0.007 (-0.24)	-0.056 (-0.82)	-0.044 (-1.58)
β_{PC7}	0.042 (1.33)	0.062 (1.34)	0.127*** (2.64)
β_{PC8}	0.029 (0.71)	0.049 (1.01)	0.086 (1.45)
$\bar{R}^2$	0.109	0.041	0.267
N	197	98	99

continued on next page

<i>(continued)</i>			
	Full Sample	First Half	Second Half
Index Skew			
α (% p.a.)	2.21*** (5.95)	2.47*** (4.86)	1.75*** (3.58)
β_{PC1}	-0.006 (-0.83)	0.001 (0.16)	0.013 (0.85)
β_{PC2}	0.012 (0.90)	0.001 (0.05)	0.000 (0.01)
β_{PC3}	0.037*** (3.03)	0.025*** (2.77)	0.120** (2.47)
β_{PC4}	0.018 (0.92)	-0.011 (-0.62)	0.074*** (2.62)
β_{PC5}	0.042** (2.46)	0.021 (1.43)	0.145*** (2.72)
β_{PC6}	-0.075*** (-2.95)	-0.038* (-1.68)	-0.043 (-1.59)
β_{PC7}	0.020 (1.07)	0.002 (0.06)	0.019 (0.80)
β_{PC8}	0.042* (1.66)	-0.007 (-0.28)	0.097*** (2.62)
$\bar{R}^2$	0.134	0.005	0.253
N	209	104	105

A.8.3 Conditional Alpha: Christopherson–Ferson–Glassman (PCA)

Table ?? reports the Christopherson–Ferson–Glassman conditional alpha using PCA factors: because the model uses the full continuous variation in the stress proxy rather than a discrete split, it establishes that the regime effect of Table ?? does not hinge on the choice of cutoffs; the cutoffs themselves are the 60 and 100 basis-point thresholds that govern the transaction-cost tiers (Appendix ??).

Table A.XVIII**Conditional Alpha (PCA, Contemporaneous): Christopherson, Ferson, and Glassman (1998) Model**

Christopherson, Ferson, and Glassman (1998) full conditional model, $r_t = \alpha_0 + \alpha_1 z_{t-1} + (\beta + \delta z_{t-1})' \mathbf{PC}_t + \varepsilon_t$, with the $K = 8$ full-sample PCA factors: both the alpha and the factor loadings are conditioned on z_{t-1} , the one-month-lagged standardized iTraxx Main 5Y spread (a predetermined conditioning variable). Timing: Contemporaneous. Significance of α_1 from a moving-block bootstrap ($B = 9,999$, $H_0: \alpha_1 = 0$ imposed, block length 9); Newey–West HAC t -statistics in parentheses; the HAC p -value is a robustness check. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	BTP Italia	CDS–Bond Basis	iTraxx Skew
α_0 (ann. %)	+1.32***	+4.80***	+2.18***
α_1 (% p.a. per 1σ)	+1.33***	+2.24***	+1.39***
t -stat	(3.17)	(3.50)	(3.97)
p (bootstrap)	0.0021	0.0089	0.0015
p (HAC, robustness)	0.0015	0.0005	0.0001
$\alpha(z = +1\sigma)$ ann.	+2.73	+6.73	+3.67
$\alpha(z = -1\sigma)$ ann.	+0.06	+2.25	+0.88
T	156	196	208

A.9 Factor-Selection Robustness*A.9.1 Sensitivity to the Subset Cardinality*

Table ?? re-estimates the best-subset selection as the cardinality k varies around the baseline. The ℓ_0 support is not nested by construction, and the data show both cases: for BTP Italia and the iTraxx skew the smaller supports are contained in the larger ones, whereas for the CDS–Bond basis correlated candidates substitute for one another as k grows. The alpha is flat in k in all three cases, so the conclusions do not hinge on the choice of k .

Table A.XIX
Best-Subset Support as the Cardinality k Varies

Each row reports the best-subset solution of exactly k factors (Bertsimas, King and Mazumder, 2016) on the full factor panel; the primary specification fixes $k = 8$. α is annualized (% p.a.), with Newey–West HAC t -statistics; $\bar{R}^2$ is the in-sample adjusted R^2 on the standardized design. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

k	α (% p.a.)	t -stat	$\bar{R}^2$	Selected factors
<i>BTP Italia</i>				
5	+2.24 ***	5.84	0.27	MKT_EU, SMB_EU, FI_VAL, FI_DEF, SS5Y
8	+2.19 ***	5.59	0.31	DEF_EU, LIQ_V, MKT_EU, SMB_EU, FI_VAL, FI_DEF, 5Y5Y_INFL_US, SS5Y
10	+1.85 ***	4.07	0.32	DEF_EU, LIQ_V, MKT_EU, SMB_EU, UMD_EU, FI_VAL, FI_DEF, PTFSBD, 5Y5Y_INFL_US, SS5Y
<i>CDS–Bond Basis</i>				
5	+4.67 ***	5.76	0.18	PB_CDS_1Y_EU, DEF_US, EMERG_DEBT, RI_EU, PTFSFX
8	+4.80 ***	5.87	0.21	PB_CDS_5Y_EU, DEF_US, HY_CORP, EMERG_DEBT, RI_EU, FX_MOM, PTFSFX, CURV_2S5S10S_EU
10	+3.87 ***	5.32	0.22	EMERG_DEBT, RI_EU, SMB_EU, COM_MOM, PTFSFX, SVS_RET_1M, IV_BUND, INFL_LINK, CURV_2S5S10S_EU, SS5Y
<i>iTraxx Skew</i>				
5	+3.02 ***	4.99	0.19	HKM_IC, BAB_EU, FI_MOM, SVS_RET_1M, ATM_IV_ITRX
8	+2.93 ***	5.30	0.25	CDX_IG, HKM_IC, HML_EU, BAB_EU, GMOM, FI_MOM, SVS_RET_1M, ATM_IV_ITRX
10	+2.72 ***	5.54	0.28	CDX_IG, DEF_US, HKM_IC, HML_EU, BAB_EU, GMOM, FI_MOM, SVS_RET_1M, ATM_IV_ITRX, SLOPE_2S10S_US

A.9.2 Factor Selection Across Methods

Table ?? reports the factor set selected for each strategy by every method of Table ?. The overlap across selectors is partial, as expected when highly correlated candidates substitute for one another, but the economically central exposures recur in every method, and the double-selection control set is a strict subset of exactly these consensus exposures—the fixed-income defensive factor for BTP Italia, emerging-market debt and the industrial bond return for the CDS–Bond basis—so the identified risks are not an artifact of the cardinality constraint or of any single penalization scheme.

Table A.XX
Selected Factors Across Selection Methods

The factor set selected for each strategy by every method of Table ??; the methods themselves are described in Section ?. $\emptyset$ marks a method that retains no factor: the relaxed lasso, post-double-selection and the knockoffs let the data choose the cardinality, and for the iTraxx skew none of the 69 candidates clears their threshold. The identity of the selected factors varies across methods, as expected when correlated candidates substitute for one another, while the alpha of Table ?? does not.

Method	Selected factors
<i>BTP Italia</i>	
Best subset	DEF_EU, LIQ_V, MKT_EU, SMB_EU, FI_VAL, FI_DEF, 5Y5Y_INFL_US, SS5Y
One per risk category	CRED_SPR_EU, LIQ_V, QMJ_EU, FI_DEF, PTFSCOM
Adaptive elastic net	HY_CORP, UMD_EU, FI_DEF, R2_EU
Relaxed lasso	XOVER_MAIN, LIQ_V, MKT_EU, SMB_EU, UMD_EU, QMJ_EU, FI_CARRY, FI_DEF, PTFSCOM, 5Y5Y_INFL_US, INFL_LINK, SS5Y
Random forest	PB_CDS_5Y_US, XOVER_MAIN, DEF_US, SMB_EU, UMD_EU, FI_DEF, 5Y5Y_INFL_US, R2_EU
Model-X knockoffs	$\emptyset$
Double-selection	FI_DEF
<i>CDS–Bond Basis</i>	
Best subset	PB_CDS_5Y_EU, DEF_US, HY_CORP, EMERG_DEBT, RI_EU, FX_MOM, PTFSEFX, CURV_2S5S10S_EU
One per risk category	RI_EU, HKM_IC, RB_EU, IV_BUND, R10_EU
Adaptive elastic net	CRED_SPR_EU, EMERG_DEBT, RI_EU, SS5Y
Relaxed lasso	EMERG_DEBT, RI_EU
Random forest	PB_CDS_5Y_EU, PB_CDS_1Y_EU, DEF_EU, EMERG_DEBT, RI_EU, COM_MOM, SVS_RET_1M, IV_BUND
Model-X knockoffs	$\emptyset$
Double-selection	EMERG_DEBT, RI_EU
<i>iTraxx Skew</i>	
Best subset	CDX_IG, HKM_IC, HML_EU, BAB_EU, GMOM, FI_MOM, SVS_RET_1M, ATM_IV_ITRX
One per risk category	CDX_IG, HKM_IC, BAB_EU, FI_MOM, SVS_RET_1M
Adaptive elastic net	PB_CDS_5Y_US, SLOPE_3S5S_MAIN, CRED_SPR_EU, DEF_EU, EMERG_DEBT, HKM_IC, HML_EU, BAB_EU, GMOM, FI_MOM, FI_DEF, FX_CARRY, PTFSEFX, SVS_RET_1M, ATM_IV_ITRX, IV_TSY, TERM_US, SLOPE_2S10S_EU, SLOPE_2S10S_US, SS5Y, SS10Y
Relaxed lasso	$\emptyset$
Random forest	CDX_IG, DEF_EU, EMERG_DEBT, SMB_EU, HML_EU, FI_MOM, SVS_RET_1M, ATM_IV_ITRX
Model-X knockoffs	$\emptyset$
Double-selection	$\emptyset$

A.9.3 Conditional Alpha: Detailed Results (Best-Subset)

Table ?? reports the full Christopherson–Ferson–Glassman conditional alpha model with best-subset-selected factors; as in the PCA case, the continuous conditioning variable makes the regime effect independent of the choice of discrete cutoffs.

Table A.XXI
Conditional Alpha: Christopherson, Ferson, and Glassman (1998) Model

Christopherson, Ferson, and Glassman (1998) full conditional performance model, $r_t = \alpha_0 + \alpha_1 z_{t-1} + (\beta + \delta z_{t-1})' \mathbf{f}_{\hat{S},t} + \varepsilon_t$, with the best-subset-selected factors $\hat{S}$: both the alpha and the factor loadings are conditioned on z_{t-1} , the one-month-lagged standardized iTraxx Main 5Y spread (mean zero, unit variance, a predetermined conditioning variable); $\alpha_1 > 0$ indicates alpha rises with credit stress. Significance of α_1 is from a moving-block bootstrap ($B = 9,999$, $H_0: \alpha_1 = 0$ imposed, block length 9); t -statistics in parentheses are Newey–West HAC (lag 4). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	BTP Italia	CDS–Bond Basis	iTraxx Skew
<i>Panel A: Conditional alpha</i>			
α_0 (ann. %)	+2.09***	+4.83***	+2.92***
α_1 (% p.a. per 1σ)	+0.48	+2.55***	+0.07
t -stat	(1.18)	(4.31)	(0.16)
p (bootstrap)	0.2290	0.0063	0.8892
p (HAC, robustness)	0.2384	0.0000	0.8734
$\bar{R}^2$ (unconditional)	0.2980	0.2116	0.2495
$\bar{R}^2$ (conditional)	0.3619	0.2570	0.2548
$\alpha(z=+1\sigma)$ (% p.a.)	+2.39	+7.46	+2.97
$\alpha(z=-1\sigma)$ (% p.a.)	+1.43	+2.37	+2.83
T	156	196	210
<i>Panel B: Conditional betas (δ_j)</i>			
5Y5YINFL_US $\times z_t$	−0.146*	—	—
ATM _{IV} _ITRX $\times z_t$	—	—	−0.269*
BAB _{EU} $\times z_t$	—	—	+0.002
CDX _{IG} $\times z_t$	—	—	+0.018
CURV _{2S5S10S_EU} $\times z_t$	—	+0.016***	—
DEF _{EU} $\times z_t$	+0.028	—	—
DEF _{US} $\times z_t$	—	+0.024	—
EMERG _{DEBT} $\times z_t$	—	+0.041*	—
FI _{DEF} $\times z_t$	−0.016	—	—
FI _{MOM} $\times z_t$	—	—	+0.019*
FI _{VAL} $\times z_t$	−0.056*	—	—
FX _{MOM} $\times z_t$	—	−0.001	—
GMOM $\times z_t$	—	—	−0.019
HKM _{IC} $\times z_t$	—	—	+0.000
HML _{EU} $\times z_t$	—	—	−0.011
HY _{CORP} $\times z_t$	—	−0.063*	—
LIQ _V $\times z_t$	−0.020*	—	—
MKT _{EU} $\times z_t$	+0.001	—	—
PB _{CDS_5Y_EU} $\times z_t$	—	+0.094	—
PTFSFX $\times z_t$	—	+0.003	—
RI _{EU} $\times z_t$	—	−0.044	—
SMB _{EU} $\times z_t$	−0.044***	—	—
SS5Y $\times z_t$	+0.008**	—	—
SVS _{RET_1M} $\times z_t$	—	—	+0.078

A.9.4 Rolling Alpha

Figure ?? plots rolling 36-month alpha with regime shading for the best-subset-selected factors. Alpha spikes during stress episodes.



Figure A.1

Rolling Alpha with Regime Shading (Best-Subset Factors). Grey bands: 95% confidence intervals. Red shading: iTraxx Main 5Y > 100 bps. Rolling window: 36 months.

A.10 Cross-Strategy Interdependencies: Additional Tests

This appendix collects robustness checks and supplementary analyses for Section ?. All tests use the same mispricing magnitudes $m_{i,t}$ and regime definitions as the main text.

A.10.1 DCC-GARCH

We estimate bivariate DCC-GARCH(1,1) models on each pair of Δm series. Figure ?? plots the time-varying conditional correlations. Spikes in conditional correlation coincide with stress episodes (COVID-19, ECB hiking cycle), corroborating the regime-dependent results in Section ??.




Figure A.2

Dynamic Conditional Correlations (DCC-GARCH). Bivariate DCC(1,1) estimates for each pair of Δm . Shaded areas: HIGH stress regime (iTraxx Main > 100 bps).

A.10.2 Granger Causality: Lag Robustness

Table ?? reports Granger causality tests at lag lengths $p = 1, 2, 3$. The main result — iTraxx Skew Granger-causes both the CDS–Bond basis and BTP Italia — is robust across all specifications.

Table A.XXII
Granger Causality: Robustness to Lag Length

Each entry is the F -statistic for the null that the row variable does not Granger-cause the column variable, from a $\text{VAR}(p)$ on the Δm series estimated separately at $p = 1, 2, 3$ lags. Lag selection: HQC selects $p = 1$; AIC/FPE select $p = 2$; BIC selects $p = 0$ (forced to 1). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Cause	Effect	$p = 1$		$p = 2$		$p = 3$	
		F	p -val	F	p -val	F	p -val
BTP Italia	CDS–Bond Basis	0.05	0.8176	0.55	0.5790	1.61	0.1858
iTraxx Skew	CDS–Bond Basis	6.52**	0.0110	3.97**	0.0195	2.14*	0.0946
CDS–Bond Basis	BTP Italia	3.08*	0.0802	1.40	0.2465	1.53	0.2071
iTraxx Skew	BTP Italia	7.39***	0.0068	4.16**	0.0162	3.04**	0.0289
CDS–Bond Basis	iTraxx Skew	0.01	0.9043	0.16	0.8537	0.13	0.9445
BTP Italia	iTraxx Skew	0.48	0.4897	1.22	0.2970	0.87	0.4589

Tables

Table I
Summary Statistics for Fixed Income Arbitrage Strategies

This table reports summary statistics for the signal-weighted arbitrage strategies, net of transaction costs. Panel A reports monthly basis levels (in basis points), with P_{10} and P_{90} the 10th and 90th percentiles; Panel B reports monthly percentage excess returns. The basis for each strategy is constructed as described in Section ?? . In Panel B, the mean and standard deviation are annualized and the Sharpe ratio is their ratio; kurtosis is excess kurtosis; the t -statistic for the mean is corrected for serial correlation (Newey–West, 4 lags); and %Neg is the proportion of negative monthly excess returns. All series end in May 2025; the samples begin in April 2012 (BTP Italia), September 2005 (iTraxx Skew), and September 2008 (CDS–Bond Basis).

Panel A: Basis Levels (bps)										
Strategy	Mean	Median	Std Dev	Min	Max	Skew	Kurt	AC(1)	P_{10}	P_{90}
<i>BTP Italia</i>	44.7	46.3	28.6	-66.5	101.0	-1.43	3.29	0.81	14.4	73.8
<i>CDS-Bond Basis</i>	6.9	9.8	18.5	-73.9	39.5	-0.97	1.69	0.83	-18.7	30.5
<i>iTraxx Skew</i>	0.2	1.1	6.7	-37.5	21.9	-1.92	7.76	0.66	-6.1	6.6

Panel B: Strategy Returns (%)											
Strategy	n	Mean (% p.a.)	t -Stat	Std Dev (% p.a.)	Min (%)	Max (%)	Skew	Kurt	%Neg	AC(1)	Sharpe
<i>BTP Italia</i>	163	1.59	3.52	1.62	-1.25	1.81	0.81	1.77	36.8	0.08	0.98
<i>iTraxx Skew</i>	243	1.91	4.92	1.53	-2.06	2.50	1.01	7.37	38.3	0.01	1.25

Table II
Manipulation-Proof Performance Measure by Risk Aversion

MPPM is the manipulation-proof performance measure of Goetzmann et al. (2007), computed from monthly returns; ρ is the coefficient of relative risk aversion and the risk-free rate is 1-month Euribor. Returns are net of transaction costs.

Strategy	$\rho = 2$ (% p.a.)	$\rho = 3$ (% p.a.)	$\rho = 4$ (% p.a.)
BTP Italia	1.50	1.48	1.47
CDS–Bond Basis	4.43	4.39	4.35
iTraxx Skew	1.88	1.87	1.86

Table III
Duarte Factor Model Regressions

This table reports estimates of the regression $r_{i,t} = \alpha_i + \beta_1 Mkt_t + \beta_2 SMB_t + \beta_3 HML_t + \beta_4 UMD_t + \beta_5 R_{S,t} + \beta_6 R_{I,t} + \beta_7 R_{B,t} + \beta_8 R_{2,t} + \beta_9 R_{5,t} + \beta_{10} R_{10,t} + \varepsilon_{i,t}$, where $r_{i,t}$ is the monthly excess return of strategy i and the regressors are the ten Duarte et al. (2007) euro-area equity and bond risk factors. α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC). Mkt : European value-weighted market excess return over the 1-month German bill; SMB , HML , UMD : the Fama–French European size, value, and momentum factors. R_S : excess return on the EuroStoxx Banks index. R_I , R_B : excess returns on European A-rated industrial and European bank bond indices. R_2 , R_5 , R_{10} : excess returns on European 2-, 5-, and 10-year government bond portfolios. $\bar{R}^2$ is the adjusted R^2 . Sample: BTP Italia 163 months, CDS–Bond 207 months, iTraxx 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	α (% p.a.)	β_{Mkt}	β_{SMB}	β_{HML}	β_{UMD}	β_{R_S}	β_{R_I}	β_{R_B}	β_{R_2}	β_{R_5}	$\beta_{R_{10}}$	$\bar{R}^2$
<i>BTP Italia</i>	1.62*** (3.29)	−0.04* (−1.79)	−0.06*** (−2.86)	−0.03 (−1.42)	0.02 (1.13)	0.01 (1.01)	−0.03 (−0.27)	0.04 (0.49)	0.03 (0.06)	−0.04 (−0.14)	−0.02 (−0.16)	0.11
<i>CDS–Bond</i>	4.24*** (5.27)	0.00 (0.20)	0.04 (0.90)	0.03 (1.11)	−0.02 (−0.88)	−0.02* (−1.74)	0.45*** (3.94)	−0.06 (−1.57)	−0.27 (−0.45)	−0.04 (−0.10)	−0.10 (−0.71)	0.15
<i>iTraxx</i>	2.19*** (4.90)	−0.02 (−1.25)	−0.03** (−2.04)	−0.07** (−2.26)	−0.02* (−1.95)	0.01 (1.53)	0.05 (0.75)	−0.02 (−0.78)	0.06 (0.16)	0.05 (0.21)	−0.06 (−0.78)	0.06

Table IV
Fung & Hsieh Factor Model Regressions

This table reports estimates of the regression $r_{i,t} = \alpha_i + \beta_1 S\&P_t + \beta_2 SC-LC_t + \beta_3 BdOpt_t + \beta_4 FXOpt_t + \beta_5 ComOpt_t + \beta_6 10Y_t + \beta_7 CredSpr_t + \varepsilon_{i,t}$, where $r_{i,t}$ is the monthly excess return of strategy i and the regressors are the seven Fung & Hsieh (2004) euro-area hedge fund risk factors. α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC). S&P: European value-weighted market excess return. SC–LC: the Fama–French European size (small-minus-big) factor. BdOpt, FXOpt, ComOpt: trend-following returns from lookback straddles on government-bond, currency, and commodity futures. 10Y: excess return on a 10-year German government bond total-return index. CredSpr: return spread between euro-area BBB corporate bonds and 10-year German Bunds. $\bar{R}^2$ is the adjusted R^2 . Sample: BTP Italia 163 months, CDS–Bond 207 months, iTraxx 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	α (% p.a.)	$\beta_{S\&P}$	β_{SC-LC}	β_{BdOpt}	β_{FXOpt}	β_{ComOpt}	β_{10Y}	$\beta_{CredSpr}$	$\bar{R}^2$
<i>BTP Italia</i>	1.81*** (4.08)	−0.03** (−2.53)	−0.06*** (−2.70)	0.00 (1.25)	0.00 (0.03)	0.00 (1.09)	−0.01 (−0.27)	0.02 (0.30)	0.10
<i>CDS–Bond</i>	3.98*** (5.05)	−0.02 (−1.18)	0.05 (1.24)	0.00 (1.01)	−0.01* (−1.95)	−0.00 (−0.42)	0.17*** (3.43)	0.19*** (2.87)	0.11
<i>iTraxx</i> <i>Skew</i>	1.84*** (4.49)	0.00 (0.41)	−0.02 (−1.50)	0.00 (0.60)	0.00 (1.01)	0.00 (0.57)	0.03 (0.83)	0.02 (0.62)	0.00

Table V
Active Fixed Income Factor Model Regressions

This table reports estimates of the regression $r_{i,t} = \alpha_i + \beta_1 \text{Term}_t + \beta_2 \text{Global Term}_t + \beta_3 \text{Global Agg}_t + \beta_4 \text{Infl. Linkers}_t + \beta_5 \text{Corp Credit}_t + \beta_6 \text{Emerg Debt}_t + \beta_7 \text{Emerg Curr}_t + \beta_8 \text{Rates Vol}_t + \varepsilon_{i,t}$, where $r_{i,t}$ is the monthly excess return of strategy i and the regressors are the eight Active Fixed Income risk premia of Brooks et al. (2020). α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC). Term: euro-area government bond excess return over cash. Global Term: Bloomberg Global Treasury Index (hedged), in excess of cash. Global Agg: Bloomberg Global Aggregate Bond Index (hedged), in excess of cash. Infl. Linkers: global inflation-linked bonds (hedged), in excess of cash. Corp Credit: pan-European high-yield corporate bond index, in excess of cash. Emerg Debt: excess return on hard-currency emerging-market debt (hedged). Emerg Curr: excess return on a basket of emerging-market currencies against the U.S. dollar. Rates Vol: excess return on a one-month variance swap on the ten-year Bund future, in percentage points of annualised variance. $\bar{R}^2$ is the adjusted R^2 . Sample: BTP Italia 163 months, CDS–Bond 207 months, iTraxx 243 months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	α (% p.a.)	β_{Term}	β_{GlobTerm}	β_{GlobAgg}	$\beta_{\text{InflLinkers}}$	$\beta_{\text{CorpCredit}}$	$\beta_{\text{EmergDebt}}$	$\beta_{\text{EmergCurr}}$	$\beta_{\text{Rates Vol}}$	$\bar{R}^2$
<i>BTP Italia</i>	2.12*** (3.84)	-0.01 (-0.47)	-0.05 (-0.16)	0.13 (0.49)	-0.06 (-1.60)	-0.12** (-2.23)	0.05 (0.97)	-0.01 (-0.30)	-0.16 (-1.36)	0.08
<i>CDS–Bond</i>	4.20*** (5.20)	0.08* (1.89)	-0.43 (-0.73)	0.30 (0.56)	-0.07 (-1.05)	-0.03 (-0.53)	0.14*** (3.15)	0.01 (0.16)	0.32 (1.15)	0.12
<i>iTraxx</i> <i>Skew</i>	1.75*** (4.24)	-0.04* (-1.92)	-0.09 (-0.42)	0.33 (1.42)	-0.00 (-0.01)	0.02 (1.61)	-0.08** (-2.18)	0.01 (0.55)	0.19 (1.53)	0.04

Table VI
Alpha Estimates Across Benchmark Factor Models

Note: Panel A reports annualized alpha (% p.a.) from monthly OLS regressions with Newey–West HAC standard errors (4 lags) using EUR factors. Panel B reports the intercept α_0 and the stress-loading coefficient α_1 from the full conditional model $r_t = \alpha_0 + \alpha_1 z_{t-1} + (\beta + \delta z_{t-1})' X_t + \varepsilon_t$, in which both the alpha and the factor loadings are conditioned on z_{t-1} , the one-month-lagged standardized iTraxx Main 5Y spread (a predetermined conditioning variable); $\bar{R}^2$ is the adjusted R^2 . Significance of α_1 is assessed by a moving-block bootstrap ($B = 9,999$, $H_0: \alpha_1 = 0$ imposed, block length 9). Panel C reports out-of-sample alpha in the recursive design of β : loadings estimated by OLS on the expanding past-only window $[1, t]$ (60-month burn-in), and the realized abnormal return $e_{i,t} = r_{i,t} - \hat{\beta}_{i,t-1}' \mathbf{f}_t$ tested with Newey–West HAC; IR is the annualized information ratio of the hedged return, and N the number of out-of-sample months, common to the three benchmarks within each strategy (rolling-window counterpart: Appendix ??). t -statistics in parentheses are Newey–West HAC throughout. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Panel A: Full-sample alpha (% p.a.)

Strategy	Duarte et al. (2007)			Brooks et al. (2020)			Fung & Hsieh (2004)		
	α	t	N	α	t	N	α	t	N
<i>BTP Italia</i>	+1.62***	3.29	163	+2.12***	3.84	163	+1.81***	4.08	163
<i>iTraxx Skew</i>	+2.19***	4.90	243	+1.75***	4.24	243	+1.84***	4.49	243
<i>CDS–Bond Basis</i>	+4.24***	5.27	207	+4.20***	5.20	207	+3.98***	5.05	207

Panel B: Christopherson, Ferson, and Glassman (1998) conditional alpha (% p.a.)

	Duarte et al. (2007)			Brooks et al. (2020)			Fung & Hsieh (2004)		
	α_0	α_1	$\bar{R}^2$	α_0	α_1	$\bar{R}^2$	α_0	α_1	$\bar{R}^2$
<i>BTP Italia</i>	+1.58*** (3.78)	+1.20** (2.25)	0.23	+1.91*** (3.03)	+1.02 (1.14)	0.13	+1.70*** (5.24)	+0.50 (1.22)	0.19
<i>iTraxx Skew</i>	+1.94*** (4.75)	+0.93** (2.55)	0.11	+1.82*** (4.43)	+1.43*** (3.40)	0.14	+1.89*** (4.85)	+1.51*** (4.66)	0.05
<i>CDS–Bond Basis</i>	+4.67*** (5.99)	+2.10** (2.74)	0.25	+4.27*** (6.24)	+3.18*** (4.02)	0.18	+3.96*** (5.45)	+1.73** (2.47)	0.17

Panel C: Out-of-sample alpha (% p.a.)

	N	Duarte et al. (2007)			Brooks et al. (2020)			Fung & Hsieh (2004)		
		α	t	IR	α	t	IR	α	t	IR
<i>BTP Italia</i>	103	+1.29**	2.24	+0.78	+2.10***	3.83	+1.25	+1.69***	2.86	+1.05
<i>iTraxx Skew</i>	183	+2.27***	4.71	+1.37	+1.77***	3.48	+1.07	+2.02***	4.05	+1.20
<i>CDS–Bond Basis</i>	147	+3.48***	4.29	+1.25	+3.86***	4.30	+1.31	+3.45***	3.87	+1.19

Table VII
PCA Spanning Regressions

This table reports estimates of the spanning regression $r_{i,t} = \alpha_i + \beta_1 PC_{1,t} + \beta_2 PC_{2,t} + \dots + \beta_8 PC_{8,t} + \varepsilon_{i,t}$, where $r_{i,t}$ is the monthly excess return of strategy i and the regressors are the first 8 principal components of the standardized candidate factor panel, extracted over the full estimation sample (the number of components is selected by the Bai–Ng IC_{p2} criterion). α is annualized (% p.a.). t -statistics in parentheses (Newey–West HAC, $\lfloor 4(T/100)^{2/9} \rfloor = 4$ lags). $\bar{R}^2$ is the adjusted R^2 ; N is the number of monthly observations. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	α (% p.a.)	β_{PC1}	β_{PC2}	β_{PC3}	β_{PC4}	β_{PC5}	β_{PC6}	β_{PC7}	β_{PC8}	$\bar{R}^2$	N
Contemporaneous ($PC_t \rightarrow R_t$)											
<i>BTP Italia</i>	1.44*** (3.30)	-0.03** (-2.25)	-0.02** (-2.02)	0.03 (1.03)	0.05 (1.45)	0.00 (0.04)	-0.02 (-0.62)	-0.04 (-1.47)	-0.02 (-1.13)	0.088	157
<i>CDS–Bond</i>	4.78*** (6.25)	0.04* (1.91)	0.06*** (2.96)	-0.09*** (-2.80)	-0.01 (-0.59)	-0.02 (-0.55)	-0.01 (-0.24)	0.04 (1.33)	0.03 (0.71)	0.109	197
<i>Index Skew</i>	2.21*** (5.95)	-0.01 (-0.83)	0.01 (0.90)	0.04*** (3.03)	0.02 (0.92)	0.04** (2.46)	-0.07*** (-2.95)	0.02 (1.07)	0.04* (1.66)	0.134	209

Table VIII
Top Factor Loadings by Principal Component

For each of the first 3 principal components, the 6 candidate factors with the largest absolute loading, with the corresponding eigenvector entry. Loadings are the full-sample eigenvectors of the standardized candidate factor panel; the sign indicates the direction of the relationship.

PC1		PC2		PC3	
Factor	Loading	Factor	Loading	Factor	Loading
CRED_SPR_EU	+0.20	GLOBAL_AGG	+0.30	SNRFIN_MAIN	+0.30
DEF_US	+0.20	GLOBAL_TERM	+0.28	BTP_BUND_2Y	+0.21
DEF_EU	+0.19	R10_EU	+0.28	PB_CDS_5Y_EU	+0.20
CRED_SPR_US	+0.19	GOVT_EU	+0.28	CREDIT_US	-0.20
ITRX_MAIN	+0.19	TERM_EU	+0.28	BTP_BUND	+0.19
HY_CORP	+0.19	RI_EU	+0.27	FI_VAL	+0.19

Table IX
Out-of-Sample Alpha: PCA and Best-Subset

Out-of-sample alpha of the PCA (Panel A) and best-subset (Panel B) hedges, in the recursive design of ?. At each month t the hedge is estimated on a past-only window—expanding with a 60-month burn-in (headline) or rolling 60-month, which relaxes the constant-loadings assumption—and the realized abnormal return is $e_{i,t} = r_{i,t} - \hat{\beta}'_{i,t-1} \mathbf{f}_t$; the in-window intercept (the alpha under measurement) is excluded from the hedge. In Panel A the principal components are re-estimated within each window, so the design is invariant to the sign and rotation indeterminacy of PCA; in Panel B the best-subset factor set is frozen and only the loadings are re-estimated. α is the annualized mean of $e_{i,t}$ (% p.a.), tested with Newey–West HAC standard errors; IR is the annualized information ratio $\bar{e}_i / \sigma(e_i) \times \sqrt{12}$; N is the number of out-of-sample months. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Strategy	N	Expanding (60-month burn-in)			Rolling 60-month		
		α	t	IR	α	t	IR
Panel A: PCA							
<i>BTP Italia</i>	97	+1.61***	2.86	+0.97	+1.65***	3.00	+1.02
<i>CDS–Bond Basis</i>	137	+4.55***	4.98	+1.55	+3.88***	3.75	+1.25
<i>iTraxx Skew</i>	149	+1.85***	3.75	+1.11	+1.96***	3.97	+1.13
Panel B: Best-Subset							
<i>BTP Italia</i>	97	+2.05***	4.06	+1.35	+2.13***	3.95	+1.32
<i>CDS–Bond Basis</i>	137	+3.90***	4.43	+1.42	+3.74***	3.90	+1.32
<i>iTraxx Skew</i>	149	+2.50***	5.35	+1.60	+1.90***	3.52	+1.01

Table X
Post-Selection OLS on the Best-Subset Factors

The table reports OLS estimates of $r_{i,t} = \alpha_i + \sum_j \beta_{i,j} F_{j,t} + \varepsilon_{i,t}$, where $r_{i,t}$ is the monthly excess return of arbitrage strategy i and $F_{j,t}$ are the factors selected for that strategy by best-subset selection (Bertsimas, King and Mazumder, 2016). Factors enter in original (un-standardized) units; α is annualized (% p.a.). The 95% confidence interval for α is the percentile interval from a stationary bootstrap (9999 replications). t -statistics use Newey–West HAC standard errors (4 lags). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. Factor definitions are in Table ??.

	Coefficient	t -stat
<i>Panel A: BTP Italia</i>		
α (% p.a.)	2.2***	5.59
95% CI (bootstrap)	[1.3, 3.0]	
DEF_EU	0.058**	2.43
LIQ_V	−0.016*	−1.88
MKT_EU	−0.038***	−3.99
SMB_EU	−0.056***	−2.79
FI_VAL	−0.094**	−2.28
FI_DEF	−0.207***	−4.94
5Y5Y_INFL_US	−0.232***	−2.70
SS5Y	−0.010**	−2.26
$T = 157, k = 8, \bar{R}^2 = 0.31$		
<i>Panel B: CDS–Bond Basis</i>		
α (% p.a.)	4.8***	5.87
95% CI (bootstrap)	[2.9, 6.6]	
PB_CDS_5Y_EU	−0.223**	−2.28
DEF_US	0.112***	3.21
HY_CORP	−0.109**	−2.20
EMERG_DEBT	0.074**	2.17
RI_EU	0.217***	3.06
FX_MOM	−0.041	−1.33
PTFSFX	−0.006**	−2.26
CURV_2S5S10S_EU	−0.015*	−1.82
$T = 197, k = 8, \bar{R}^2 = 0.21$		
<i>Panel C: iTraxx Skew</i>		
α (% p.a.)	2.9***	5.30
95% CI (bootstrap)	[1.5, 3.8]	
CDX_IG	0.167**	2.40
HKM_IC	−0.025***	−4.14
HML_EU	−0.026*	−1.91
BAB_EU	−0.044***	−4.26
GMOM	−0.052***	−3.05
FI_MOM	−0.041***	−3.87
SVS_RET_1M	−0.857**	−2.52
ATM_IV_ITRX	0.304	1.32
$T = 209, k = 8, \bar{R}^2 = 0.25$		

Table XI
Alpha by Stress Regime: PCA and Best-Subset

Annualized alpha (% p.a.) per stress regime, from a single regression with three intercept dummies and betas common across regimes (no constant), $r_t = \sum_g \alpha_g D_{g,t} + \sum_j \beta_j X_{jt} + \varepsilon_t$, Newey–West HAC standard errors (4 lags). PCA: spanning regression on the $K = 8$ full-sample principal components. Best-Subset: post-selection OLS on the selected factors (Table ??). Regimes on the 5-year iTraxx Europe Main spread: LOW < 60, MEDIUM $\in [60, 100)$, HIGH ≥ 100 bps. t -statistics in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Strategy	PCA			Best-Subset		
	LOW	MEDIUM	HIGH	LOW	MEDIUM	HIGH
<i>BTP Italia</i>	+0.37 (0.65)	+1.68*** (3.17)	+3.27* (1.82)	+1.28** (2.29)	+2.40*** (5.08)	+3.77*** (2.95)
<i>CDS–Bond Basis</i>	+2.87*** (3.12)	+4.44*** (3.87)	+7.55*** (4.84)	+2.69*** (2.99)	+4.69*** (3.95)	+7.59*** (4.93)
<i>iTraxx Skew</i>	+1.03* (1.95)	+2.23*** (4.62)	+3.45*** (4.32)	+2.48*** (3.95)	+2.94*** (4.38)	+3.39*** (4.14)

Table XII
Alpha Across Empirical Layers: Factor Benchmarks, PCA, and Best-Subset

Each cell is the annualized alpha (% p.a.) and adjusted R^2 from a spanning regression of the strategy excess return on the factor set of the corresponding layer: the three established hedge-fund factor benchmarks of Section ??, the principal-component factors, and the best-subset selection (Bertsimas, King and Mazumder, 2016). A residual alpha that survives all three layers is evidence that the premium is not compensation for the identified systematic exposures. Alphas are annualized; stars denote significance from each layer's own inference. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Method	<i>BTP Italia</i>		<i>CDS–Bond</i>		<i>iTraxx</i>	
	α	$\bar{R}^2$	α	$\bar{R}^2$	α	$\bar{R}^2$
Duarte ($k = 10$)	+1.6***	0.11	+4.2***	0.15	+2.2***	0.06
Brooks ($k = 8$)	+2.1***	0.08	+4.2***	0.12	+1.7***	0.04
Fung–Hsieh ($k = 7$)	+1.8***	0.10	+4.0***	0.11	+1.8***	0.00
PCA ($k = 8$)	+1.4***	0.09	+4.8***	0.11	+2.2***	0.13
Best-subset ($k = 8$)	+2.2***	0.31	+4.8***	0.21	+2.9***	0.25

Table XIII
Alpha Across Factor-Selection Methods

For each selection method, n is the number of selected factors and two post-selection alphas of the monthly net-of-cost strategy excess return are reported. α^{split} : honest split-sample alpha — factors are selected on the first half of the sample only, and the alpha is the intercept of an OLS on the second half (valid post-selection inference by sample splitting); on each half the adaptive elastic net is re-fitted, so the split-sample alpha is available for every selector. α^{OOS} : frozen-set out-of-sample alpha — with the selected set held fixed, loadings are re-estimated each month on an expanding window (60-month burn-in) and the alpha is the mean abnormal return $y_t - \mathbf{x}_t' \hat{\beta}_{t-1}$, with the estimated intercept excluded from the fitted value. Alphas are annualized (% p.a.); t -statistics in parentheses use Newey–West HAC standard errors (4 lags). The ℓ_0 problem is solved by the splicing algorithm (abess) polished against greedy forward search; risk categories are the panels of Table ???. Selection with zero factors reports the raw strategy alpha. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Method	BTP Italia			CDS–Bond Basis			iTraxx Skew		
	n	α^{split}	α^{OOS}	n	α^{split}	α^{OOS}	n	α^{split}	α^{OOS}
Best subset ($\ell_0, k=8$)	8	+1.69*** (2.67)	+2.05*** (3.96)	8	+2.92*** (4.31)	+3.90*** (4.43)	8	+2.49*** (3.59)	+2.50*** (5.35)
Best subset, one factor per risk category	5	+1.91*** (2.83)	+1.23*** (2.59)	5	+2.52*** (3.47)	+3.21*** (3.58)	5	+2.18*** (3.19)	+2.18*** (4.47)
Adaptive elastic net	4	+1.96*** (2.94)	+1.43*** (2.91)	4	+2.85*** (4.40)	+3.29*** (3.70)	21	+2.10*** (4.30)	+2.74*** (6.50)
Relaxed lasso	12	+1.96*** (2.94)	+1.60*** (3.02)	2	+3.11*** (3.76)	+3.54*** (3.85)	0	+2.20*** (3.00)	+1.79*** (3.36)
Random forest screening (top-8)	8	+1.54*** (3.44)	+1.30*** (2.73)	8	+3.41*** (4.03)	+3.80*** (4.29)	8	+2.19*** (3.76)	+1.87*** (3.66)
Model-X knockoffs (FDR 10%)	0	+1.96*** (2.94)	+1.79*** (3.19)	0	+3.11*** (3.76)	+3.77*** (4.04)	0	+2.06*** (2.81)	+1.79*** (3.36)

Table XIV
Unconditional and Regime-Dependent Correlations

$m_{i,t}$ is the market-level mispricing magnitude of strategy i and $\Delta m_{i,t}$ its monthly change. Panel A reports correlations of monthly strategy returns; Panel B, of the Δm series. Unconditional: Pearson correlations with Newey–West HAC t -statistics (4 lags). Regime columns: within-regime Pearson correlations, with regimes on the iTraxx Main 5Y level (LOW < 60, MEDIUM [60, 100), HIGH $\geq$ 100 bps); stars from a two-sided t -test with $n - 2$ degrees of freedom. HIGH^{FR} applies the ? adjustment for the higher return variance in the HIGH regime, so that the rise in correlation is not a heteroskedasticity artifact. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Panel A: Strategy Returns

Pair	Unconditional		By Regime			
	Pearson	t -stat	LOW	MED	HIGH	HIGH ^{FR}
CDS–Bond Basis vs BTP Italia	−0.189	−1.54	−0.123	−0.285**	−0.134	−0.066
CDS–Bond Basis vs iTraxx Skew	−0.020	−0.19	−0.079	−0.019	−0.038	−0.017
BTP Italia vs iTraxx Skew	0.217***	2.70	0.151	0.251**	0.154	0.082

Panel B: Mispricing Changes (Δm)

Pair	Unconditional		By Regime			
	Pearson	t -stat	LOW	MED	HIGH	HIGH ^{FR}
CDS–Bond Basis vs BTP Italia	−0.298***	−3.04	−0.134	−0.324***	−0.465**	−0.326
CDS–Bond Basis vs iTraxx Skew	0.191*	1.73	−0.208*	0.201*	0.316**	0.135*
BTP Italia vs iTraxx Skew	−0.100	−1.07	0.150	−0.147	−0.144	−0.069

Table XV
Cross-Strategy Spanning Regressions (Δm)

Pairwise regressions $\Delta m_{i,t} = \alpha + \beta \Delta m_{j,t} + \varepsilon_t$, where $\Delta m_{i,t}$ is the monthly change in the mispricing magnitude of strategy i . Newey–West HAC standard errors (4 lags). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Dependent	Regressor	β	t -stat	p	$\bar{R}^2$
CDS–Bond Basis	BTP Italia_L0	−0.148**	-2.20	0.0282	0.0913
	BTP Italia_L1	−0.010	-0.22	0.8288	
	BTP Italia_L2	−0.039	-1.27	0.2034	
	iTraxx Skew_L0	+0.370	0.91	0.3637	
	iTraxx Skew_L1	+0.710*	1.93	0.0530	
	iTraxx Skew_L2	+0.030	0.08	0.9367	
BTP Italia	iTraxx Skew_L0	−0.687	-1.41	0.1581	0.1000
	iTraxx Skew_L1	−0.950	-1.36	0.1728	
	iTraxx Skew_L2	+0.391	0.46	0.6437	
	CDS–Bond Basis_L0	−0.424**	-2.43	0.0150	
	CDS–Bond Basis_L1	+0.215	1.57	0.1161	
	CDS–Bond Basis_L2	+0.072	0.62	0.5340	
iTraxx Skew	BTP Italia_L0	−0.010	-0.76	0.4482	0.0103
	BTP Italia_L1	+0.016	0.80	0.4210	
	BTP Italia_L2	+0.028*	1.94	0.0529	
	CDS–Bond Basis_L0	+0.017	0.52	0.5997	
	CDS–Bond Basis_L1	+0.019	0.75	0.4520	
	CDS–Bond Basis_L2	+0.008	0.39	0.6945	

Table XVI
Stress Amplification: Interaction Regressions

$\Delta m_{i,t} = \alpha + \beta_0 \Delta m_{j,t} + \beta_1 (\Delta m_{j,t} \times z_t) + \gamma z_t + \varepsilon_t$, where $\Delta m_{i,t}$ is the monthly change in the mispricing magnitude of strategy i and z_t is the contemporaneous standardized iTraxx Main 5Y level. Newey–West HAC standard errors (4 lags). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Dependent	Variable	Coefficient	t -stat	p	$\bar{R}^2$
CDS–Bond Basis	Δm_j (BTP Italia)	−0.144**	-2.02	0.043	0.1221
	$\Delta m_j \times z_t$	−0.086	-1.50	0.133	
	z_t (stress)	+1.591***	2.63	0.009	
BTP Italia	Δm_j (CDS–Bond Basis)	−0.464***	-3.22	0.001	0.1286
	$\Delta m_j \times z_t$	−0.637***	-2.79	0.005	
	z_t (stress)	+1.502	0.95	0.341	
CDS–Bond Basis	Δm_j (iTraxx Skew)	+0.482	1.05	0.291	0.0531
	$\Delta m_j \times z_t$	+0.016	0.07	0.940	
	z_t (stress)	+0.934**	2.29	0.022	
iTraxx Skew	Δm_j (CDS–Bond Basis)	+0.030	1.13	0.261	0.0731
	$\Delta m_j \times z_t$	+0.079**	2.01	0.044	
	z_t (stress)	−0.076	-0.31	0.757	
BTP Italia	Δm_j (iTraxx Skew)	−0.389	-0.66	0.510	0.0154
	$\Delta m_j \times z_t$	−0.618	-0.93	0.353	
	z_t (stress)	−0.128	-0.07	0.941	
iTraxx Skew	Δm_j (BTP Italia)	−0.010	-0.82	0.412	0.0201
	$\Delta m_j \times z_t$	−0.020	-1.01	0.313	
	z_t (stress)	−0.045	-0.19	0.849	

Table XVII
AR(1) Persistence of Mispricing Levels

AR(1) regression of the mispricing level $m_{i,t} = c_i + \phi_i m_{i,t-1} + u_{i,t}$, where $m_{i,t}$ is the market-level mispricing magnitude of strategy i . The half-life is $h_i^* = \ln(0.5)/\ln(\phi_i)$ months. Newey–West HAC standard errors (4 lags). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Strategy	ϕ	HAC t -stat	Half-life (months)	N
CDS–Bond Basis	0.7913***	14.43	3.0	148
BTP Italia	0.7272***	13.54	2.2	148
iTraxx Skew	0.7258***	9.06	2.2	148

Table XVIII
Common Factor (PC1) and Funding Stress

OLS regression of the first principal component of the three Δm series on intermediary stress proxies and on a macro placebo set. Panel A: one proxy per Duffie channel (intermediary capital, funding cost, dealer health, market illiquidity). Panel B: all available funding and liquidity proxies. Panel C: macroeconomic placebo set (macro and real-activity uncertainty, long-term inflation expectations, economic policy uncertainty). Significant coefficients in Panels A and B indicate that the common factor driving correlated mispricings is linked to intermediary balance-sheet conditions (???). The non-significant coefficients in Panel C are consistent with the common factor not being driven by macroeconomic fundamentals. Newey–West HAC standard errors (4 lags). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Factor	Coefficient	<i>t</i> -stat	<i>p</i> -value
Panel A: Core Intermediary Proxies			
$T = 149, \bar{R}^2 = 0.233$			
HKM	−4.278**	−2.28	0.023
LIBOR_OIS	+3.349***	3.28	0.001
PB_CDS_5Y_EU	+0.349**	2.01	0.044
ILLIQ	+1.766***	2.83	0.005
Panel C: Macroeconomic Placebo			
$T = 151, \bar{R}^2 = 0.106$			
$\Delta U M$	+0.363	0.03	0.976
$\Delta U R$	+10.313	1.23	0.219
5Y5Y_INFL	−0.627	−0.66	0.510
EPU_EU	+0.003	1.32	0.185




Figure 1

The TIPS–Treasury Basis, 2005–2025. The series is the average TIPS–Treasury mispricing across matched bond pairs, in basis points, following ?. The vertical dashed line marks the end of their sample (19 November 2009). The basis exceeded 140 bps during the 2008–09 crisis; that extreme dislocation has not recurred, and the basis has since compressed, turning slightly negative in 2021–22.



Figure 2

Market-Level Basis Time Series for Fixed Income Arbitrage Strategies. For BTP Italia, the basis is the cross-sectional mean of all outstanding issues with at least six months to maturity. For iTraxx Skew, it is the mean of the four on-the-run sub-index skews. For CDS-Bond Basis, it is the cross-sectional median of the single-name bases in the iTraxx Main universe. Dashed lines indicate entry thresholds.




Figure 3

Cumulative Return Indices and Drawdowns of Fixed Income Arbitrage Strategies (Signal-Weighted). Each row shows one strategy: the upper sub-panel reports the cumulative return index from inception, and the lower sub-panel reports the running drawdown (peak-to-trough). Sample periods differ across strategies: BTP Italia from April 2012, iTraxx Skew from September 2005, CDS–Bond Basis from September 2008. All series end in May 2025 and are net of transaction costs.




Figure 4

Rolling 36-Month Alpha Across Benchmark Factor Models. Each panel shows one strategy. Lines represent rolling alpha under Duarte et al. (2007, blue), Brooks et al. (2020, orange), and Fung & Hsieh (2004, green). Red shading: HIGH stress regime (iTraxx Main 5Y $\geq$ 100 bps).




Figure 5

Scree Plot: Cumulative Variance Explained by Principal Components. Bars show the variance explained by each component; the line shows the cumulative share. The vertical dashed line marks the baseline $K = 8$, selected by the Bai–Ng IC_{p2} criterion; the first eight components capture 64.7% of total variance.




Figure 6

Pairwise Correlations by Stress Regime. Bars show Pearson correlations of Δm in LOW, MEDIUM, and HIGH stress regimes defined by iTraxx Main 5Y thresholds at 60 and 100 bps, as in Table ??.




Figure 7
Generalized Impulse Response Functions. Responses of Δm to one-standard-deviation shocks. Shaded areas: 95% bootstrap confidence intervals (1000 replications).